Software Tool Design

<u>The User interface (UI) design using Figma</u>

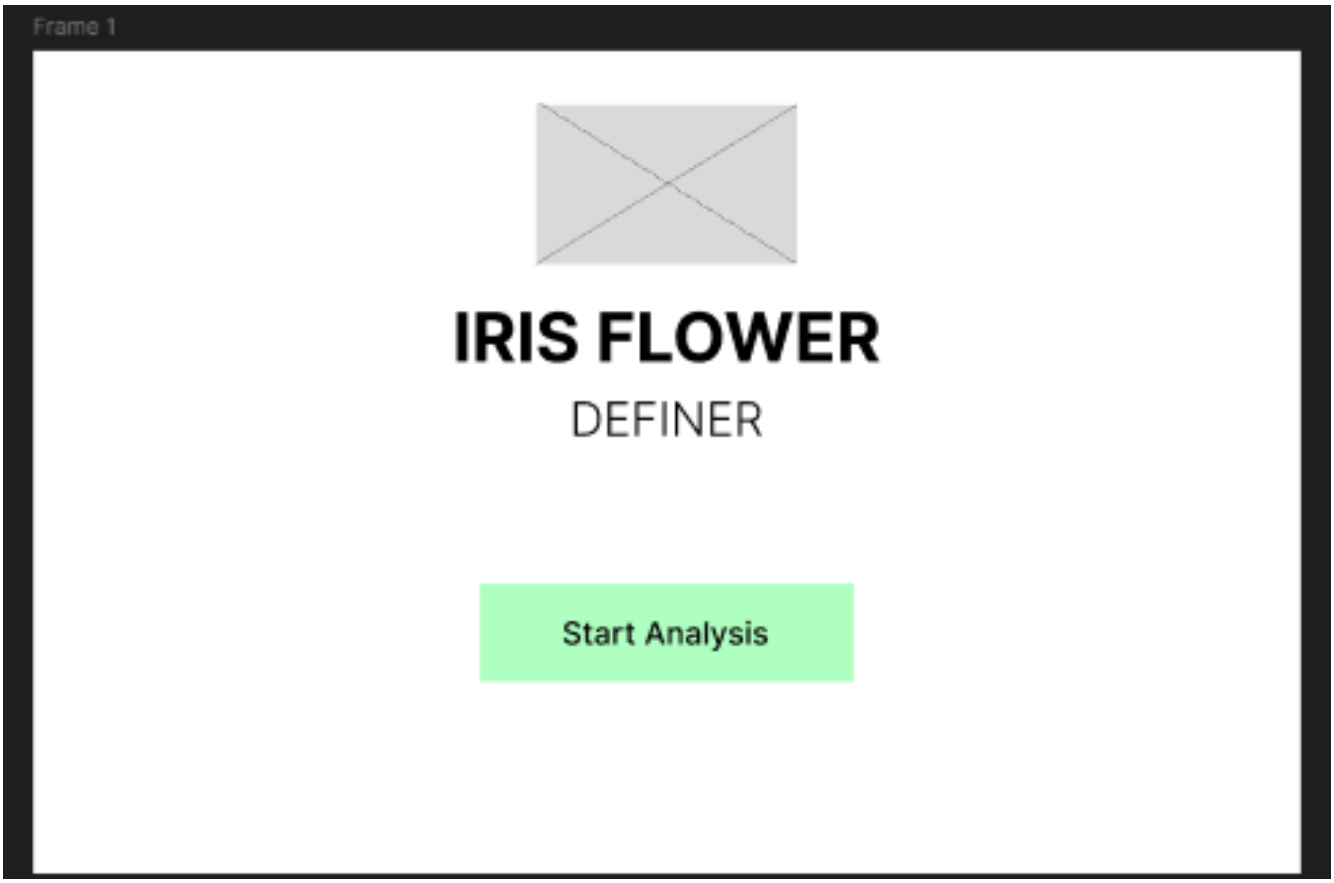

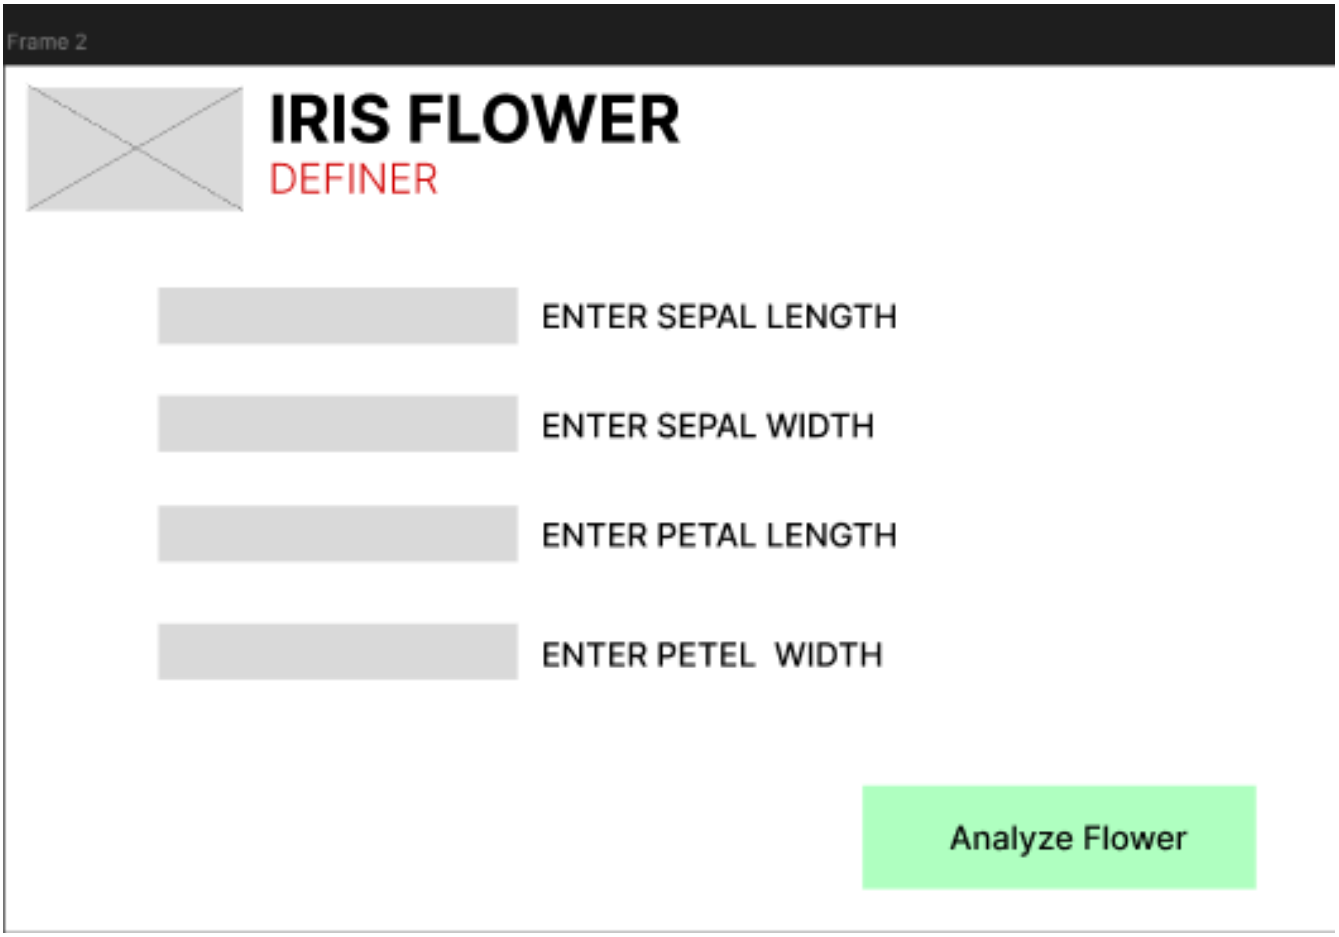
Frame 2
IRIS FLOWER
DEFINER
ENTER SEPAL LENGTH
ENTER SEPAL WIDTH
ENTER PETAL LENGTH
ENTER PETEL WIDTH
Analyze Flower

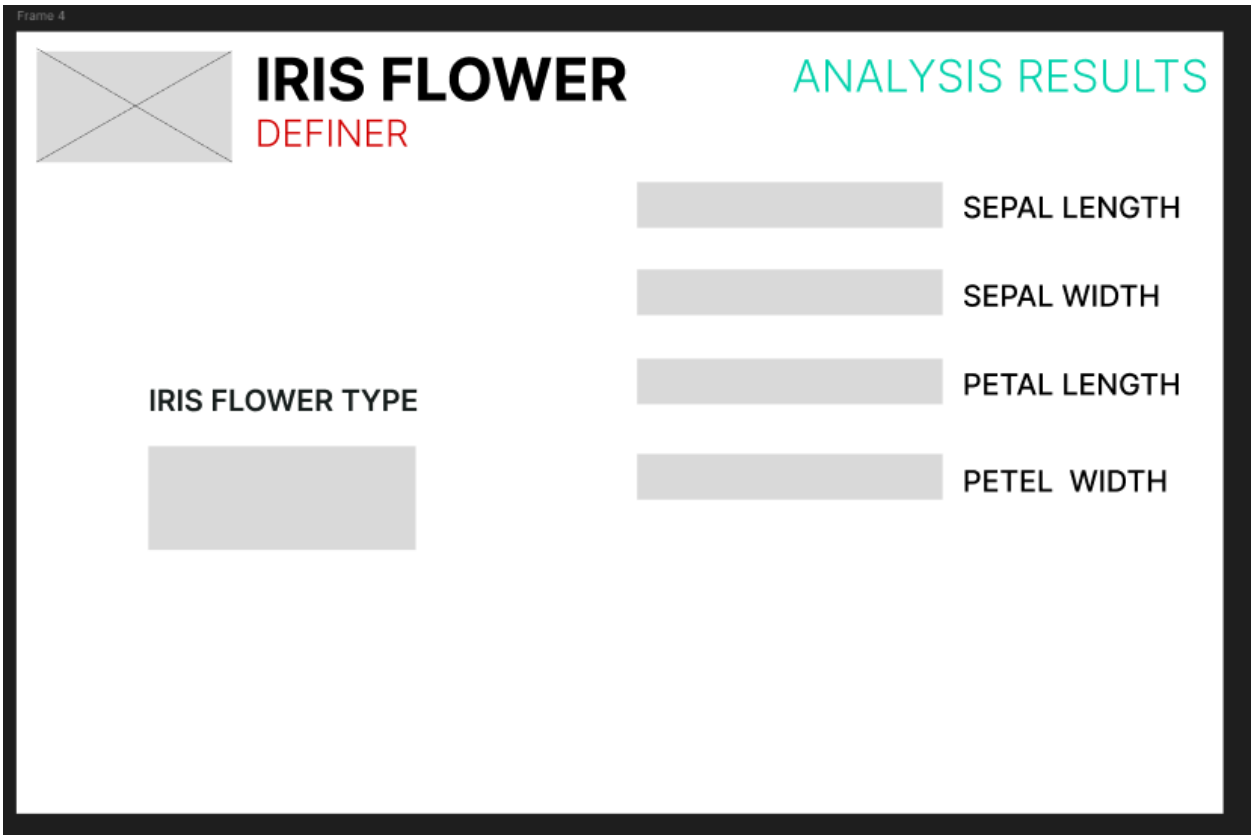

-++

Figure 1 - The Ui design for Iris flower definer

Logical design

The following are the contents to be described through the logical design of this application

I. User Interface design
II. Backend Processing
III. Data Storage
IV. Scalability and Efficiency
V. Testing and Validation
VI. Documentation and Maintenance

I. UI design

- **Components**
    - Input Form: The input form allows users to enter specific attributes of the iris flower, including sepal length, sepal width, petal length, and petal width. This form is designed to be straightforward to minimize input errors.
    - Buttons: Interactive buttons are included to manage the workflow. The "Classify Flower" button initiates the classification process, while the "Reset Form" button clears all input fields, allowing users to start over.
    - Results Display: This section displays the classification results. It shows the predicted species of the iris flower and provides detailed insights into the analysis.

- **Interactions**

Users interact with the UI by entering the iris flower attributes into the input form. After entering the data, clicking the "Analyze flower" button triggers the backend processing, which analyzes the data and displays the results in the results display section. If needed, the "Reset Form" button can be used to clear all inputs and start the process again.

II. Backend Processing

- **Components**
    - Data Pipeline: A comprehensive data pipeline is established for acquiring, cleaning, and preprocessing the input data before it is analyzed.
    - Classification Module: This module contains algorithms or models specifically designed to analyze the flower attributes and predict the species of the iris flower.

- **Workflow**
  - Data Acquisition: Data is acquired from user input through the UI.
  - Data Cleaning: The acquired data is cleaned to ensure it is consistent and accurate. This involves handling any missing or inconsistent values.
  - Preprocessing: The cleaned data is then preprocessed to extract relevant features, making it suitable for analysis.
  - Analysis: Finally, the preprocessed data is analyzed using the classification module to predict the species of the iris flower based on its attributes.

III. Data Storage

- **Components**
  - Database: A robust database system is used to store both raw and processed data, as well as the results of the classification analysis.
  - File System: Additional storage is provided for files such as reports, logs, and configuration files.
- **Interactions**

The database stores the input data, intermediate results, and final classification outcomes for future reference and further analysis. The file system is used to store additional information necessary for the application's operation and maintenance

IV. Scalability and Efficiency

- **Strategies**
  - Parallel Processing: Parallel processing techniques distribute computational tasks across multiple cores or nodes, enhancing processing efficiency.

    - Distributed Computing: Frameworks such as Apache Spark are implemented to process large datasets across a cluster of machines, improving scalability.

- **Optimizations**
    - Algorithm Optimization: Classification algorithms are optimized for performance and scalability, ensuring they can handle large data volumes efficiently.
    - Data Pipeline Optimization: The data preprocessing and analysis pipelines are streamlined for faster processing and reduced resource consumption.

V. Testing and Validation

- **Approaches**
    - Unit Testing: Individual components are tested in isolation to ensure they work correctly.
    - Integration Testing: The interactions between different components are validated to verify overall system functionality.
    - End-to-End Testing: The entire workflow, from data input to results display, is tested to ensure it meets all requirements.

- **Validation**

The accuracy and reliability of the classification results are validated against known benchmarks or ground truth data. Stress testing is also performed to assess the application's performance under load, ensuring scalability.

VI. Documentation and Maintenance

- **Documentation**
  - The architecture, design decisions, and implementation details of the application are thoroughly documented for future reference and maintenance.
  - User manuals and guides are provided to help users understand how to use the application, interpret the results, and troubleshoot common issues.

- **Maintenance**
  - A maintenance plan is established to address bug fixes, updates, and enhancements over time.
  - Regular monitoring and optimization are conducted to ensure the application continues to meet user needs efficiently.

By adhering to this detailed logical design, the Iris dataset application can be developed to accurately analyze the morphological characteristics of iris flowers and predict their species with high precision, ensuring a robust and efficient system for users.

The UML design for the application

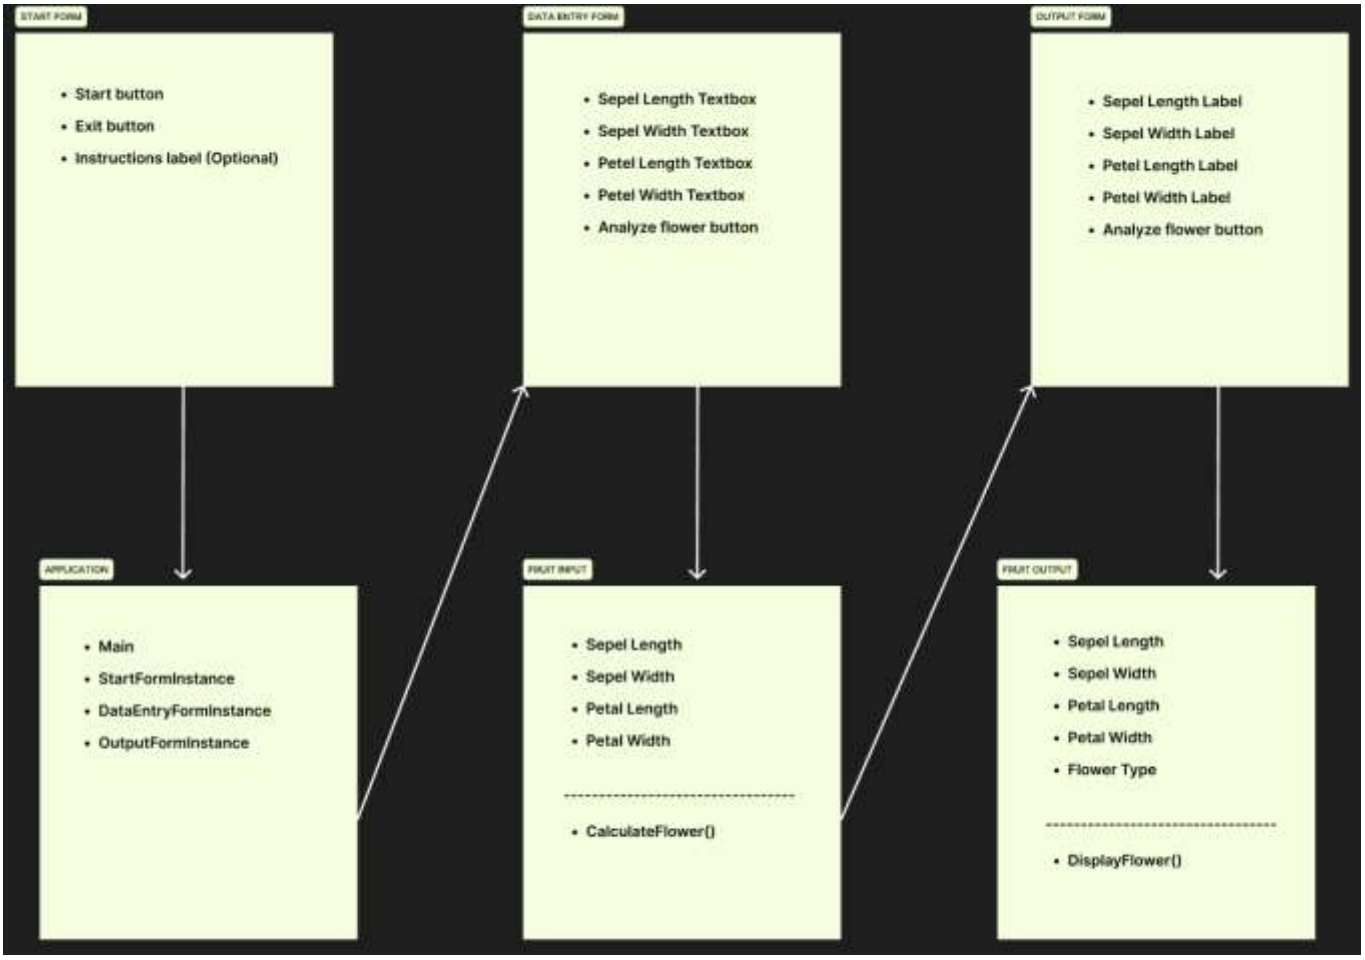

Figure 2- The UML diagram for the application

The Acquisition process

During the acquisition phase of the Iris Dataset application, several important steps must be taken to seamlessly incorporate the dataset into the system.

Firstly, it is essential to find trustworthy sources of data. The Iris dataset is often available on platforms like Kaggle, the UCI Machine Learning Repository, and other academic repositories. These platforms provide detailed and reliable information. These sources are well-known for their carefully curated datasets, which are widely used in the research and machine learning communities. Exploring multiple sources can also enrich the dataset by providing additional context or different versions of the data.

Once the data sources are identified, the next step involves handling different data formats effectively. The Iris dataset typically comes in formats such as CSV, which is easy to read and process. However, the application must be capable of handling various data formats like JSON or Excel if they are encountered. This involves implementing strong data parsing techniques to extract the necessary information while ensuring compatibility with the application's data structure. Effective data parsing ensures that the dataset is correctly formatted and ready for further processing.

Integration with data sources can be achieved through API integration or manual download processes, depending on the accessibility of the dataset. API integration is advantageous for real-time access to updated datasets and allows for automated data retrieval. This ensures that the application always has the most current data available. On the other hand, manual download processes are suitable for datasets that require periodic updates or where direct download options are more practical. This approach is often simpler and sufficient for datasets that do not change frequently.

Overall, the acquisition phase of the Iris Dataset application focuses on finding reliable data sources, handling different data formats effectively, and seamlessly integrating the dataset into the system. By addressing these aspects, the application can smoothly transition from data

acquisition to the next stages of preprocessing and analysis. This ensures the accuracy and reliability of the classification results.

Dataset Loading

```
[1] #Install necessary libraries
    !pip install pandas numpy matplotlib seaborn scipy
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.0.3)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (1.25.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (0.13.1)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (1.11.4)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from panda
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2023.4
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (9
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.
```

```
[2] #Import the required libraries
    import pandas as pd
    import numpy as np
    import matplotlib.pyplot as plt
    import seaborn as sns
    from scipy import stats
    from google.colab import files
    import warnings
```

```
[3] #Set up seaborn color palette
    sns.set_palette("husl")
```

```
[4] #Upload the dataset
    uploaded = files.upload()
    for file_name in uploaded.keys():
        df = pd.read_csv(file_name)
```

Choose Files iris.data.csv

- **iris.data.csv**(text/csv) - 4551 bytes, last modified: 9/20/2019 - 100% done

```
Saving iris.data.csv to iris.data.csv
```

```
df.head()
```

	5.1	3.5	1.4	0.2	Iris-setosa
0	4.9	3.0	1.4	0.2	Iris-setosa
1	4.7	3.2	1.3	0.2	Iris-setosa
2	4.6	3.1	1.5	0.2	Iris-setosa
3	5.0	3.6	1.4	0.2	Iris-setosa
4	5.4	3.9	1.7	0.4	Iris-setosa

Next steps: Generate code with df | View recommended plots

Figure 3 - Code snippet in loading the iris dataset

The initial step in the script involves installing essential Python libraries using the `pip` package manager. These libraries include `pandas`, `numpy`, `matplotlib`, `seaborn`, and `scipy`. These libraries are fundamental for data analysis and visualization. `Pandas` is used for data manipulation and analysis, `numpy` for numerical operations, `matplotlib` and `seaborn` for creating plots and visualizations, and `scipy` for scientific and statistical operations. Although Google Colab typically comes with many libraries pre-installed, this command ensures that the latest versions are available for use.

Following the installation, the script imports the required libraries. This includes `pandas` for handling data structures, `numpy` for numerical computations, `matplotlib.pyplot` and `seaborn` for plotting and visualizations, and `scipy.stats` for performing statistical analyses. Additionally, the script imports `files` from `google.colab` to handle file uploads and `warnings` to filter out unnecessary warnings that could clutter the output. This setup prepares the environment with all the tools needed for the subsequent data analysis.

The script then configures the color palette for Seaborn, choosing the "husl" palette. This step ensures that all visualizations created using Seaborn have a consistent and aesthetically pleasing color scheme. A well-chosen color palette is crucial for making plots not only more attractive but also more readable and interpretable.

The next part of the script deals with uploading the dataset. It utilizes the `files.upload()` function from Google Colab to open a file dialog, prompting the user to select a file from their local machine. Once the file is chosen, it is uploaded to the Colab environment and stored in the `uploaded` variable. This approach allows for easy and interactive data uploading, which is particularly useful for working with datasets that are not readily accessible via URLs.

After the file is uploaded, the script processes it by iterating through the uploaded files. Typically, only one file is uploaded, but the loop ensures that every uploaded file is processed correctly. The `pd.read_csv(file_name)` function reads the CSV file into a pandas DataFrame named `df`. This DataFrame serves as a structured data format that allows for efficient data manipulation and analysis using pandas' powerful features. This step is critical as it transforms raw data into a format that is easy to work with.

The final part of the script involves displaying the first few rows of the DataFrame to verify that the data has been loaded correctly. The `head()` function is used to show the top five rows of the dataset. This initial inspection step is crucial for ensuring the data is in the expected format and for getting a quick overview of its structure and contents. By reviewing the first few rows, one can confirm that the dataset includes the attributes of iris flowers—such as sepal length, sepal width, petal length, petal width, and the species name—and is ready for further analysis and visualization.

Data Cleaning

DATA CLEANING

[7]
```python
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
import warnings
```

[8]
```python
# Display the first few rows of the dataset
print("First few rows of the dataset:")
print(df.head())
```

```
First few rows of the dataset:
   5.1  3.5  1.4  0.2  Iris-setosa
0  4.9  3.0  1.4  0.2  Iris-setosa
1  4.7  3.2  1.3  0.2  Iris-setosa
2  4.6  3.1  1.5  0.2  Iris-setosa
3  5.0  3.6  1.4  0.2  Iris-setosa
4  5.4  3.9  1.7  0.4  Iris-setosa
```

[9]
```python
# 1. Check for missing values
print("\nChecking for missing values:")
print(df.isnull().sum())
```

```

Checking for missing values:
5.1            0
3.5            0
1.4            0
0.2            0
Iris-setosa    0
dtype: int64
```

[10]
```python
# 2. Check for duplicates
print("\nChecking for duplicates:")
print(df.duplicated().sum())
```

```

Checking for duplicates:
3
```

[11]
```python
# Remove duplicates if any
df = df.drop_duplicates()
```

[12]
```python
# 3. Ensure correct data types
print("\nChecking data types:")
print(df.dtypes)
```

```
Checking data types:
5.1            float64
3.5            float64
1.4            float64
0.2            float64
Iris-setosa     object
dtype: object
```

[13]
```python
# 4. Summary statistics
print("\nSummary statistics of the dataset:")
print(df.describe())
```

```
Summary statistics of the dataset:
              5.1         3.5         1.4         0.2
count  146.000000  146.000000  146.000000  146.000000
mean     5.861644    3.052740    3.796575    1.215753
std      0.829562    0.436948    1.753987    0.755822
min      4.300000    2.000000    1.000000    0.100000
25%      5.100000    2.800000    1.600000    0.300000
50%      5.800000    3.000000    4.400000    1.300000
75%      6.400000    3.300000    5.100000    1.800000
max      7.900000    4.400000    6.900000    2.500000
```

[15]
```python
# 5. Data consistency and integrity
# Checking unique values in the target column
# First, find the actual name of the target column (species)
target_column = df.columns[-1]
print(f"\nUnique values in the {target_column} column:")
print(df[target_column].unique())
```

```
Unique values in the Iris-setosa column:
['Iris-setosa' 'Iris-versicolor' 'Iris-virginica']
```

[16]
```python
# Renaming columns for consistency
df.columns = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width', 'species']
print("\nRenamed columns:")
print(df.columns)
```

```
Renamed columns:
Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width',
       'species'],
      dtype='object')
```

```python
# Display cleaned dataset
print("\nFirst few rows of the cleaned dataset:")
print(df.head())
```

```
First few rows of the cleaned dataset:
   sepal_length  sepal_width  petal_length  petal_width      species
0           4.9          3.0           1.4          0.2  Iris-setosa
1           4.7          3.2           1.3          0.2  Iris-setosa
2           4.6          3.1           1.5          0.2  Iris-setosa
3           5.0          3.6           1.4          0.2  Iris-setosa
4           5.4          3.9           1.7          0.4  Iris-setosa
```

Figure 4- Data cleaning codes for iris dataet

The above code script begins by displaying the first few rows of the dataset to provide an initial look at the data. This step is crucial for verifying that the data has been loaded correctly and understanding its structure. The output shows that the dataset includes columns for sepal length, sepal width, petal length, petal width, and the species of the iris flower. The species column contains categorical data with values such as "Iris-setosa".

Next, the script checks for any missing values in the dataset. This is done using the `isnull().sum()` function, which counts the number of missing values in each column. The output indicates that there are no missing values in any of the columns. This step is essential for ensuring that the dataset is complete and that no data is missing, which could affect the analysis.

The script then checks for duplicate rows in the dataset using the `duplicated().sum()` function. The output reveals that there are three duplicate rows. Duplicate data can skew analysis results, so it is important to identify and remove any duplicates. The script proceeds to remove these duplicate rows using the `drop_duplicates()` function, ensuring that each row in the dataset is unique.

After handling duplicates, the script verifies the data types of each column using the `dtypes` attribute. The output shows that the first four columns are of type `float64`, indicating they contain numerical data, while the species column is of type `object`, indicating it contains categorical data. Ensuring that the data types are correct is crucial for accurate data analysis and manipulation.

The script then generates summary statistics for the dataset using the `describe()` function. This function provides various statistical measures such as count, mean, standard deviation, minimum, and maximum values for the numerical columns. These statistics offer a quick overview of the dataset's distribution and central tendency, helping to identify any anomalies or outliers.

To ensure data consistency and integrity, the script checks the unique values in the species column. It first identifies the last column as the target column for species and then displays the unique values in this column. The output confirms that the species column contains three unique values: "Iris-setosa", "Iris-versicolor", and "Iris-virginica". This step is important for verifying that the categorical data is consistent and does not contain any unexpected values.

The script proceeds to rename the columns for better readability and consistency. The new column names are more descriptive: sepal_length, sepal_width, petal_length, petal_width, and species. Renaming columns helps make the dataset more understandable and easier to work with in subsequent analysis steps.

Finally, the script displays the first few rows of the cleaned dataset to verify the changes made during the data cleaning process. The output shows that the dataset now has the new column names and no duplicate rows. This final check ensures that the dataset is properly cleaned and ready for further analysis and visualization.

This comprehensive data cleaning process includes checking for and handling missing values and duplicates, verifying data types, generating summary statistics, ensuring data consistency, renaming columns for clarity, and displaying the cleaned dataset. These steps are essential for preparing the dataset for accurate and effective analysis.

Exploratory data analysis

EXPLORATORY DATA ANALYSIS

[21]
```python
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
import warnings
```

[22]
```python
# Display the first few rows of the dataset
print("First few rows of the dataset:")
print(df.head())

# Check column names
print("\nColumn names in the dataset:")
print(df.columns)

# Rename columns for consistency if necessary
df.columns = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width', 'species']
print("\nRenamed columns:")
print(df.columns)
```

```
First few rows of the dataset:
   sepal_length  sepal_width  petal_length  petal_width      species
0           4.9          3.0           1.4          0.2  Iris-setosa
1           4.7          3.2           1.3          0.2  Iris-setosa
2           4.6          3.1           1.5          0.2  Iris-setosa
3           5.0          3.6           1.4          0.2  Iris-setosa
4           5.4          3.9           1.7          0.4  Iris-setosa

Column names in the dataset:
Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width',
       'species'],
      dtype='object')

Renamed columns:
Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width',
       'species'],
      dtype='object')
```

```python
# Display the first few rows of the dataset
print("First few rows of the dataset:")
print(df.head())

# Check column names
print("\nColumn names in the dataset:")
print(df.columns)
```

```
First few rows of the dataset:
   sepal_length  sepal_width  petal_length  petal_width      species
0           4.9          3.0           1.4          0.2  Iris-setosa
1           4.7          3.2           1.3          0.2  Iris-setosa
2           4.6          3.1           1.5          0.2  Iris-setosa
3           5.0          3.6           1.4          0.2  Iris-setosa
4           5.4          3.9           1.7          0.4  Iris-setosa

Column names in the dataset:
Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width',
       'species'],
      dtype='object')
```

[24]

```python
# Function to clean the data
def clean_data(df):
    # Drop any rows with missing values
    df = df.dropna()

    # Ensure the 'species' column exists
    if 'species' not in df.columns:
        raise KeyError("The 'species' column is missing from the DataFrame.")

    # Define a function to label the 'species' column
    def label(species):
        if species == 'Iris-setosa':
            return 0
        elif species == 'Iris-versicolor':
            return 1
        elif species == 'Iris-virginica':
            return 2

    # Apply the label function to create a new 'Label' column
    df['Label'] = df['species'].apply(label)

    # Drop the original 'species' column
    df = df.drop(columns=['species'])

    # Ensure 'Label' column is of type int64
    df['Label'] = df['Label'].astype('int64')

    return df
```

```
[25] # Apply the clean_data function to a copy of the DataFrame
     df_clean = clean_data(df.copy())
```

```
# Display the first few rows of the cleaned DataFrame
print("\nFirst few rows of the cleaned DataFrame:")
print(df_clean.head())
```

```
First few rows of the cleaned DataFrame:
   sepal_length  sepal_width  petal_length  petal_width  Label
0           4.9          3.0           1.4          0.2      0
1           4.7          3.2           1.3          0.2      0
2           4.6          3.1           1.5          0.2      0
3           5.0          3.6           1.4          0.2      0
4           5.4          3.9           1.7          0.4      0
```

```
[27] # Display the DataFrame info to verify data types and absence of missing values
     print("\nInformation about the cleaned DataFrame:")
     print(df_clean.info())
```

```
Information about the cleaned DataFrame:
<class 'pandas.core.frame.DataFrame'>
Index: 146 entries, 0 to 148
Data columns (total 5 columns):
 #   Column        Non-Null Count  Dtype
---  ------        --------------  -----
 0   sepal_length  146 non-null    float64
 1   sepal_width   146 non-null    float64
 2   petal_length  146 non-null    float64
 3   petal_width   146 non-null    float64
 4   Label         146 non-null    int64
dtypes: float64(4), int64(1)
memory usage: 6.8 KB
None
```

```
[28] # Summary statistics of the cleaned dataset
     print("\nSummary statistics of the cleaned dataset:")
     print(df_clean.describe())
```

```
Summary statistics of the cleaned dataset:
       sepal_length  sepal_width  petal_length  petal_width       Label
count    146.000000   146.000000    146.000000   146.000000  146.000000
mean       5.861644     3.052740      3.796575     1.215753    1.013699
std        0.829562     0.436948      1.753987     0.755822    0.813560
min        4.300000     2.000000      1.000000     0.100000    0.000000
25%        5.100000     2.800000      1.600000     0.300000    0.000000
50%        5.800000     3.000000      4.400000     1.300000    1.000000
75%        6.400000     3.300000      5.100000     1.800000    2.000000
max        7.900000     4.400000      6.900000     2.500000    2.000000
```

```
[31] # Correlation matrix
     corr_matrix = df_clean.corr()
     print("\nCorrelation matrix:")
     print(corr_matrix)
```

```
Correlation matrix:
              sepal_length  sepal_width  petal_length  petal_width     Label
sepal_length      1.000000    -0.103591      0.870793     0.816015  0.781488
sepal_width      -0.103591     1.000000     -0.415679    -0.350436 -0.413337
petal_length      0.870793    -0.415679      1.000000     0.961405  0.947783
petal_width       0.816015    -0.350436      0.961405     1.000000  0.955216
Label             0.781488    -0.413337      0.947783     0.955216  1.000000
```

Figure 5- EDA analysis on Iris dataset

This part of the code brings in essential libraries needed for manipulating data, conducting statistical analysis, and creating visualizations. `pandas` is crucial for managing and manipulating data structures like DataFrames, which are fundamental for data analysis tasks. `numpy` provides support for numerical operations and array manipulations. `matplotlib.pyplot` and `seaborn` are used to create visualizations, with `seaborn` offering a high-level interface built on `matplotlib`. The `scipy.stats` module offers a variety of statistical functions, while the `warnings` module helps manage and control warning messages that might appear during code execution.

At the beginning, the code prints the first few rows and the column names of the dataset. This step is critical for understanding the structure of the dataset, ensuring the data has been loaded correctly, and checking the initial state of its columns and rows. It provides an immediate insight into the nature and format of the data, which is essential for any subsequent analysis.

To ensure clarity and consistency, the columns of the dataset are renamed. This step is significant for maintaining a standard naming convention, making it easier to reference and manipulate columns in further analyses. Proper naming conventions enhance the readability of the code and help prevent errors related to column name mismatches.

After renaming, the dataset is printed again to verify that the columns have been renamed correctly. This step ensures that the intended changes have been applied properly and allows for a quick visual confirmation of the updated structure of the dataset.

The `clean_data` function performs several crucial tasks to clean the data:

1. It drops any rows with missing values using the `dropna()` method. This step is vital for ensuring the accuracy and integrity of the data, as missing values can lead to incorrect analyses and results.

2. The function checks for the presence of the 'species' column, raising an error if it is missing. This validation ensures that the critical column needed for subsequent analyses is present.
3. A nested function named `label` is defined to convert species names into numeric labels (0 for Iris-setosa, 1 for Iris-versicolor, and 2 for Iris-virginica). This conversion is necessary for various machine learning and statistical techniques that require numerical inputs.
4. The `label` function is applied to create a new 'Label' column containing the numeric representations of the species.
5. The original 'species' column is dropped to avoid redundancy and potential confusion in the dataset.
6. The 'Label' column is ensured to be of type `int64` to maintain consistency and compatibility with various analytical methods.

The `clean_data` function is applied to a copy of the DataFrame to create a cleaned version of the dataset, referred to as `df_clean`. This step ensures that the original dataset remains unchanged, preserving the raw data for reference or alternative analyses.

The information of the cleaned DataFrame is displayed to verify the data types and confirm the absence of missing values. This step provides a detailed overview of the dataset's structure, including the number of non-null entries and the data types of each column, ensuring that the data is ready for further analysis.

The `describe()` method is used to generate summary statistics for the cleaned dataset. These statistics provide insights into the central tendency (mean), dispersion (standard deviation), and overall distribution (minimum, maximum, and quartiles) of the numerical columns. This information is essential for understanding the characteristics of the dataset and identifying any potential issues or patterns.

A correlation matrix is computed and printed to show the pairwise correlation coefficients between the features. This matrix helps identify the strength and direction of relationships between variables. Values close to 1 or -1 indicate strong positive or negative correlations,

respectively, while values near 0 suggest weak or no correlation. For example, the strong positive correlation between `petal_length` and `petal_width` indicates that these features tend to increase together.

Visualizations

data Visualizations

```python
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Assuming df is already loaded and cleaned
# For demonstration, let's rename the columns again
df.columns = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width', 'species']

# Pairplot
sns.pairplot(df, hue='species', markers=["o", "s", "D"])
plt.suptitle('Pairplot of Iris Features', y=1.02)
plt.show()
```

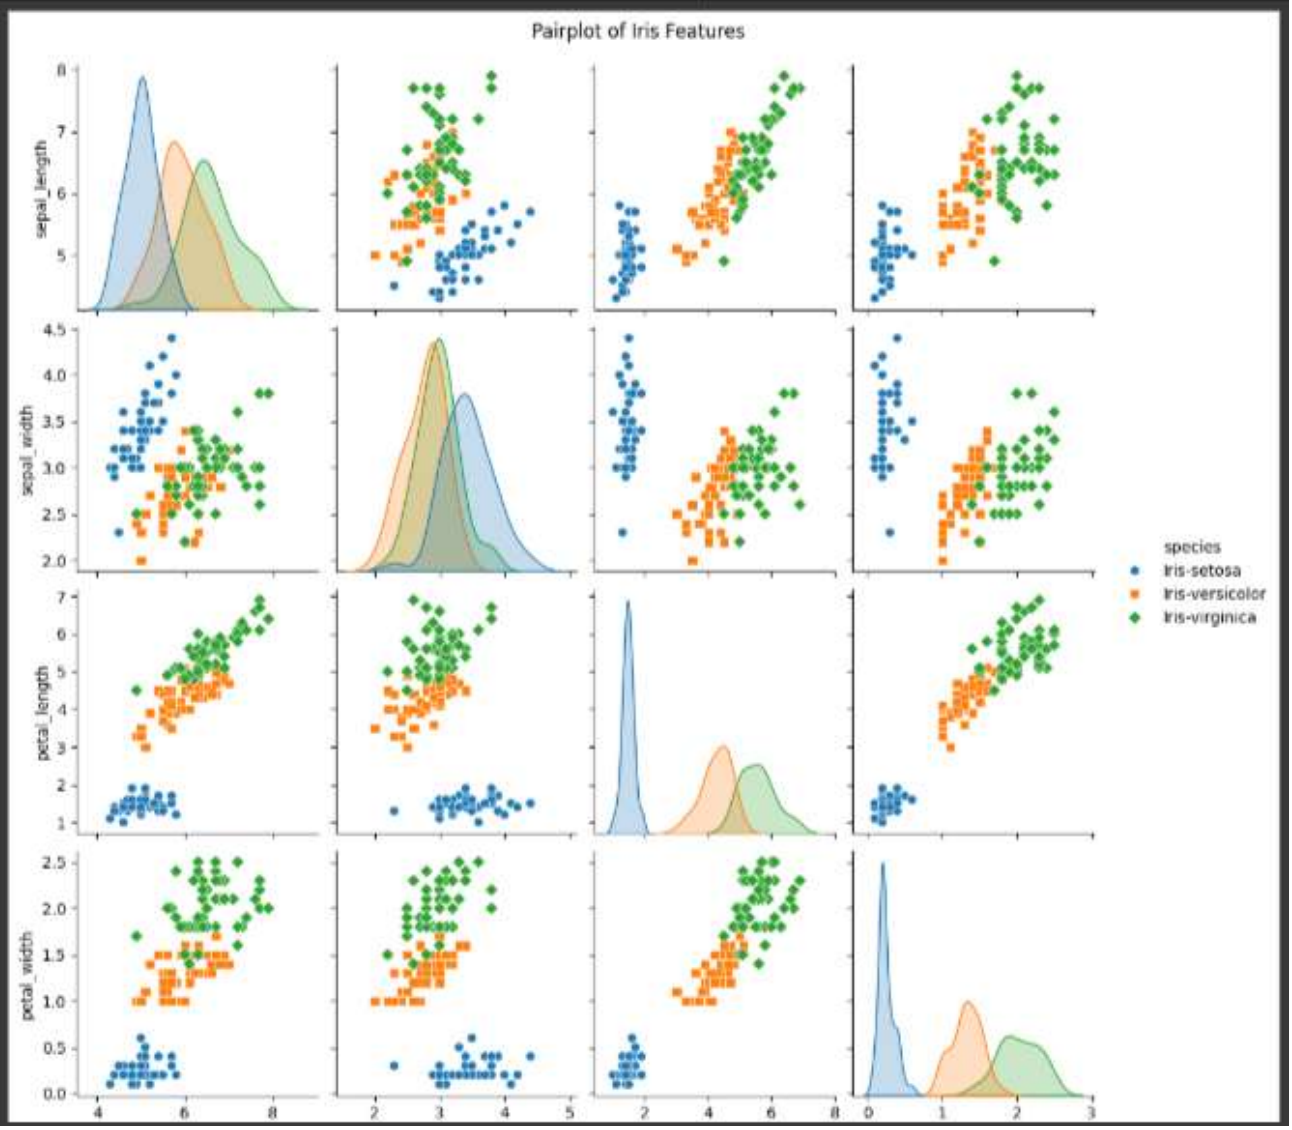

The pairplot is a versatile visualization tool provided by `seaborn`. It creates a grid of scatter plots, displaying the relationships between every pair of features in the dataset. The `hue` parameter is used to differentiate the points by species, allowing for an easy visual comparison

between different classes. The `markers` parameter specifies different symbols for each species. This plot is particularly useful for understanding the distribution and relationships between features for each species, highlighting patterns and potential separations between the species.

```
# Correlation matrix and heatmap
corr_matrix = df.drop(columns='species').corr()
print("\nCorrelation matrix:")
print(corr_matrix)

plt.figure(figsize=(8, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix Heatmap')
plt.show()
```

```
Correlation matrix:
              sepal_length  sepal_width  petal_length  petal_width
sepal_length      1.000000    -0.103591      0.870793     0.816015
sepal_width      -0.103591     1.000000     -0.415679    -0.350436
petal_length      0.870793    -0.415679      1.000000     0.961405
petal_width       0.816015    -0.350436      0.961405     1.000000
```

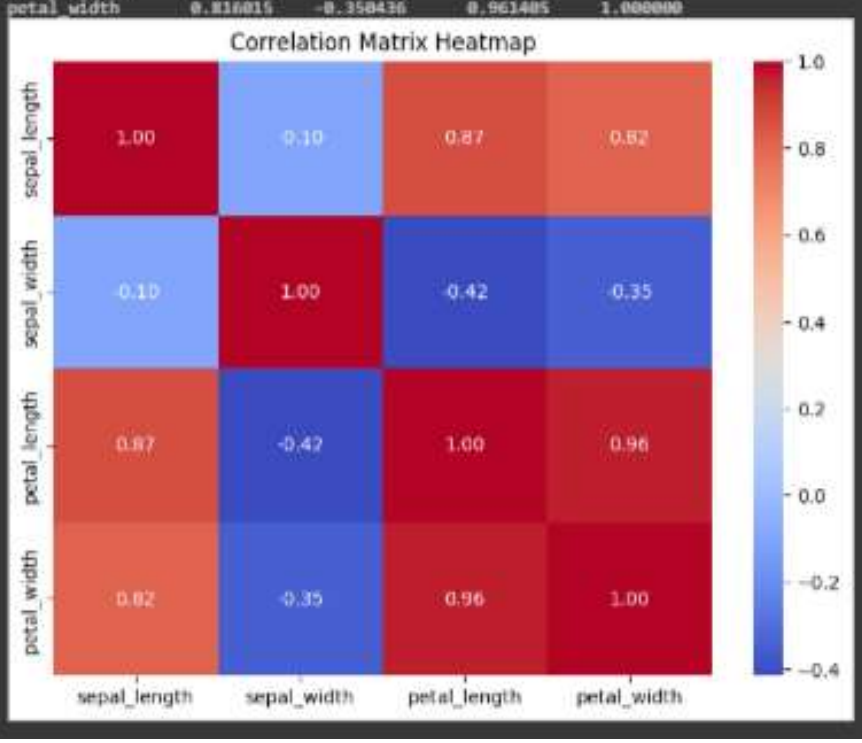

The correlation matrix is calculated by first removing the 'species' column and then computing the pairwise correlation coefficients between the remaining numerical features. The `heatmap`

function from `seaborn` is employed to visualize this correlation matrix. Each cell in the heatmap displays the correlation coefficient between two features, with the color intensity representing the strength and direction of the correlation. This visualization is beneficial for identifying which features have strong correlations, either positive or negative, thus offering insights into how the features relate to each other.

```
# Histograms
df.hist(edgecolor='black', linewidth=1.2, figsize=(12, 8))
plt.suptitle('Histograms of Iris Features')
plt.show()
```

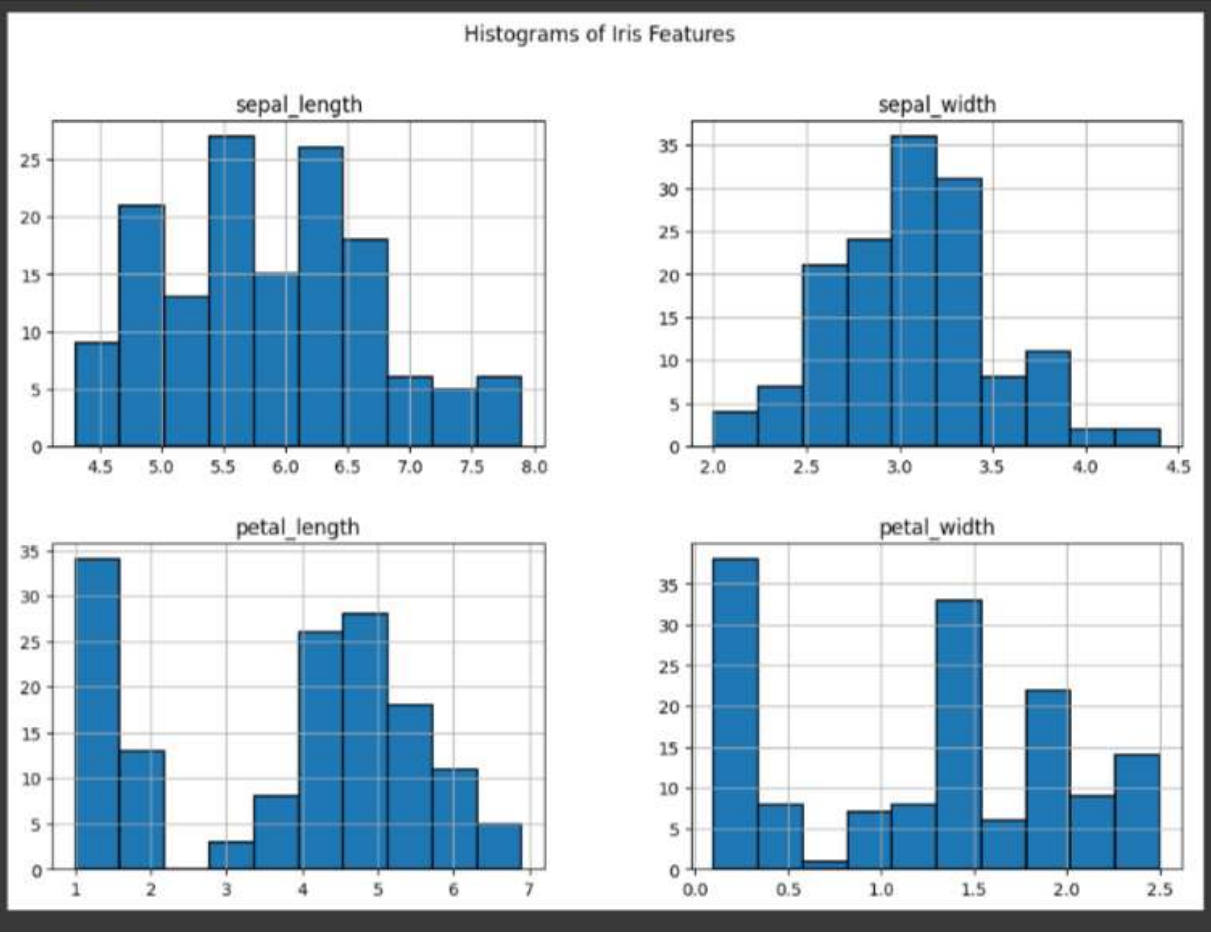

Histograms are created for each feature in the dataset. These visualizations show the distribution

of values for each feature, helping to understand the spread, central tendency, and variation in the data. The histograms are designed with black edges and a linewidth of 1.2 for better visual clarity. An overall title 'Histograms of Iris Features' is added to provide context. These plots are useful for quickly identifying the frequency distribution of the data points within each feature.

```
# Boxplots
plt.figure(figsize=(14, 8))

plt.subplot(2, 2, 1)
sns.boxplot(x='species', y='sepal_length', data=df)
plt.title('Sepal Length by Species')

plt.subplot(2, 2, 2)
sns.boxplot(x='species', y='sepal_width', data=df)
plt.title('Sepal Width by Species')

plt.subplot(2, 2, 3)
sns.boxplot(x='species', y='petal_length', data=df)
plt.title('Petal Length by Species')

plt.subplot(2, 2, 4)
sns.boxplot(x='species', y='petal_width', data=df)
plt.title('Petal Width by Species')

plt.tight_layout()
plt.show()
```

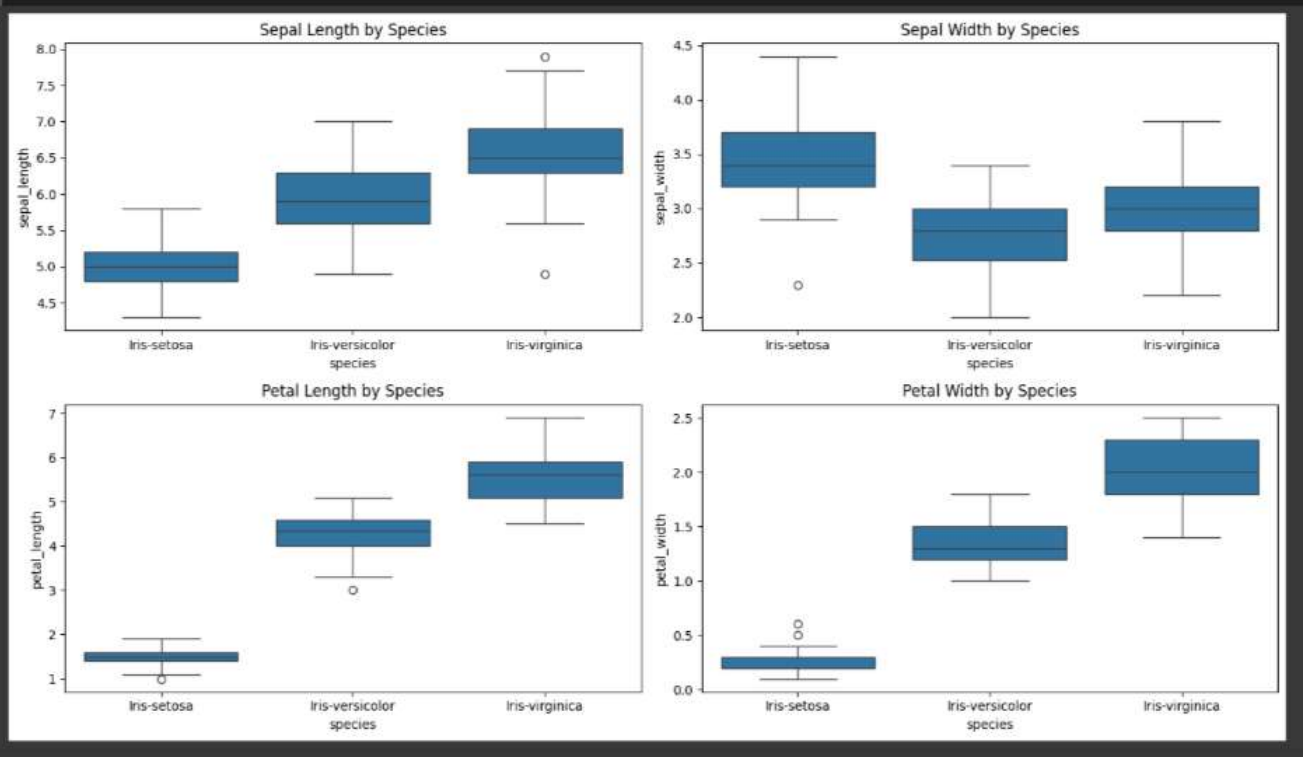

Boxplots are generated to illustrate the distribution of each feature across the different species. Each subplot represents a distinct feature, with the species on the x-axis and the feature values on the y-axis. Boxplots display the median, quartiles, and potential outliers for each species, offering a clear comparison of the feature distributions across the species. The `tight_layout()` function ensures that the subplots are arranged neatly without overlapping, enhancing the readability of the plots. These plots are instrumental in identifying the spread and central tendency of data within each species and spotting any outliers.

```
# Violin plots
plt.figure(figsize=(14, 8))

plt.subplot(2, 2, 1)
sns.violinplot(x='species', y='sepal_length', data=df)
plt.title('Sepal Length by Species')

plt.subplot(2, 2, 2)
sns.violinplot(x='species', y='sepal_width', data=df)
plt.title('Sepal Width by Species')

plt.subplot(2, 2, 3)
sns.violinplot(x='species', y='petal_length', data=df)
plt.title('Petal Length by Species')

plt.subplot(2, 2, 4)
sns.violinplot(x='species', y='petal_width', data=df)
plt.title('Petal Width by Species')

plt.tight_layout()
plt.show()
```

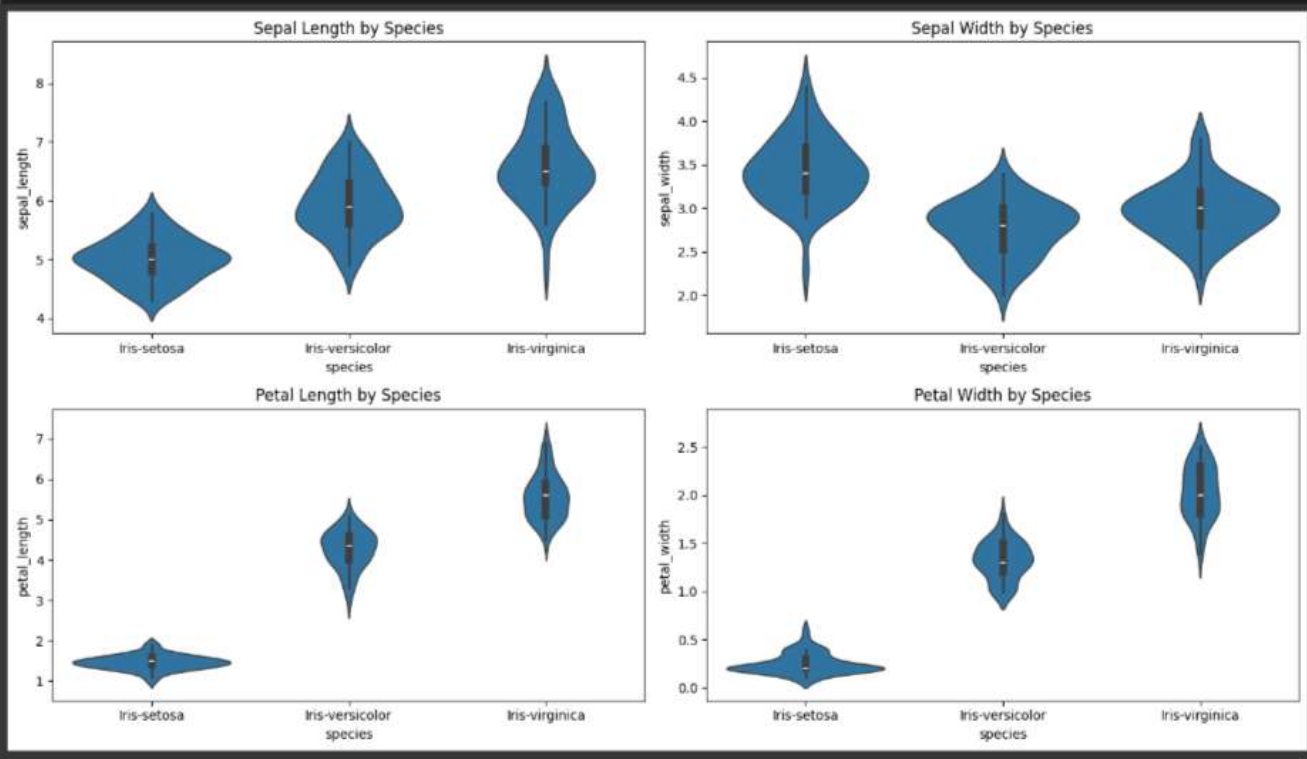

Violin plots merge elements of boxplots and kernel density plots to show the distribution of the data. Each subplot corresponds to a different feature, with the species on the x-axis and the feature values on the y-axis. Violin plots display the density of the data at different values, providing a deeper understanding of the distribution shape and spread of each feature across species. Similar to the boxplots, the `tight_layout()` function is used to arrange the subplots neatly, ensuring clarity and readability. These plots are especially useful for visualizing the distribution of the data and identifying any patterns or anomalies across the species.

These visualizations together offer a comprehensive overview of the relationships, distributions, and differences in the Iris dataset's features across different species. They provide valuable insights into the data, making it easier to understand and interpret the underlying patterns and relationships.

<u>Qualitative data analysis</u>

QUALITATIVE ANALYSIS

[1]
```python
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
```

```python
# Load the Iris dataset
iris = load_iris()
df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
df['species'] = iris.target
```

[3]
```python
# Mapping species numbers to species names
species_map = {0: 'Iris-setosa', 1: 'Iris-versicolor', 2: 'Iris-virginica'}
df['species_name'] = df['species'].map(species_map)
```

[4]
```python
# Display the first few rows of the dataset
print("First few rows of the dataset:")
print(df.head())
```

```
First few rows of the dataset:
   sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)  \
0                5.1               3.5                1.4               0.2
1                4.9               3.0                1.4               0.2
2                4.7               3.2                1.3               0.2
3                4.6               3.1                1.5               0.2
4                5.0               3.6                1.4               0.2

   species species_name
0        0  Iris-setosa
1        0  Iris-setosa
2        0  Iris-setosa
3        0  Iris-setosa
4        0  Iris-setosa
```

```
[5] # Content Analysis
    # Count the number of instances for each species
    species_counts = df['species_name'].value_counts()
    print("\nSpecies Counts:")
    print(species_counts)
```

```
Species Counts:
species_name
Iris-setosa        50
Iris-versicolor    50
Iris-virginica     50
Name: count, dtype: int64
```

```
# Visualize the distribution of species
plt.figure(figsize=(8, 6))
sns.countplot(data=df, x='species_name')
plt.title('Frequency Distribution of Iris Species')
plt.xlabel('Species')
plt.ylabel('Count')
plt.show()
```

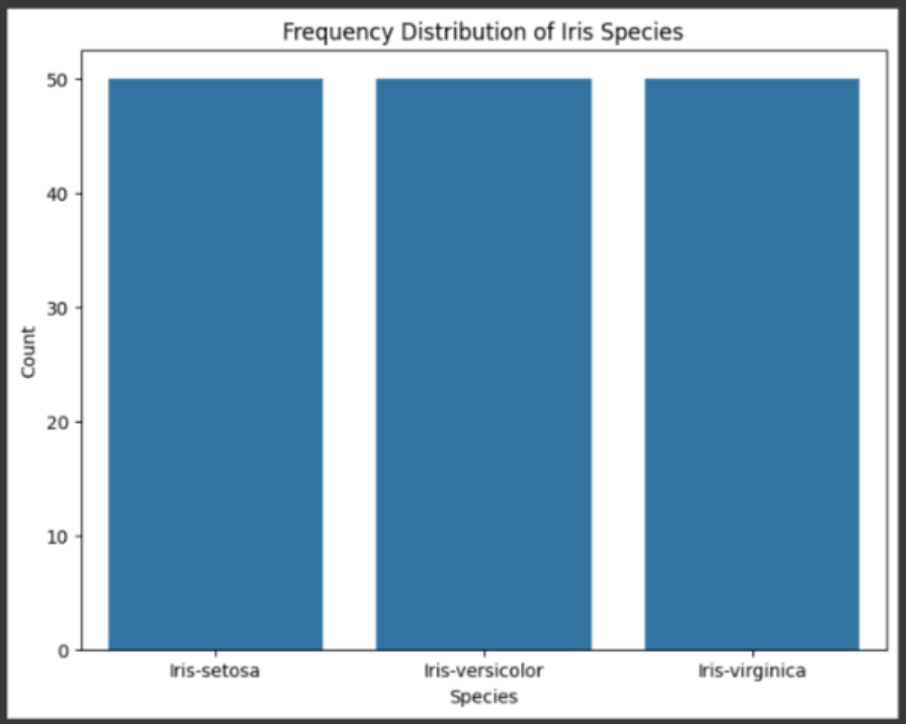

```
# Pairplot to observe relationships between features for each species
sns.pairplot(df, hue='species_name', markers=["o", "s", "D"])
plt.suptitle('Pairplot of Iris Features by Species', y=1.02)
plt.show()
```

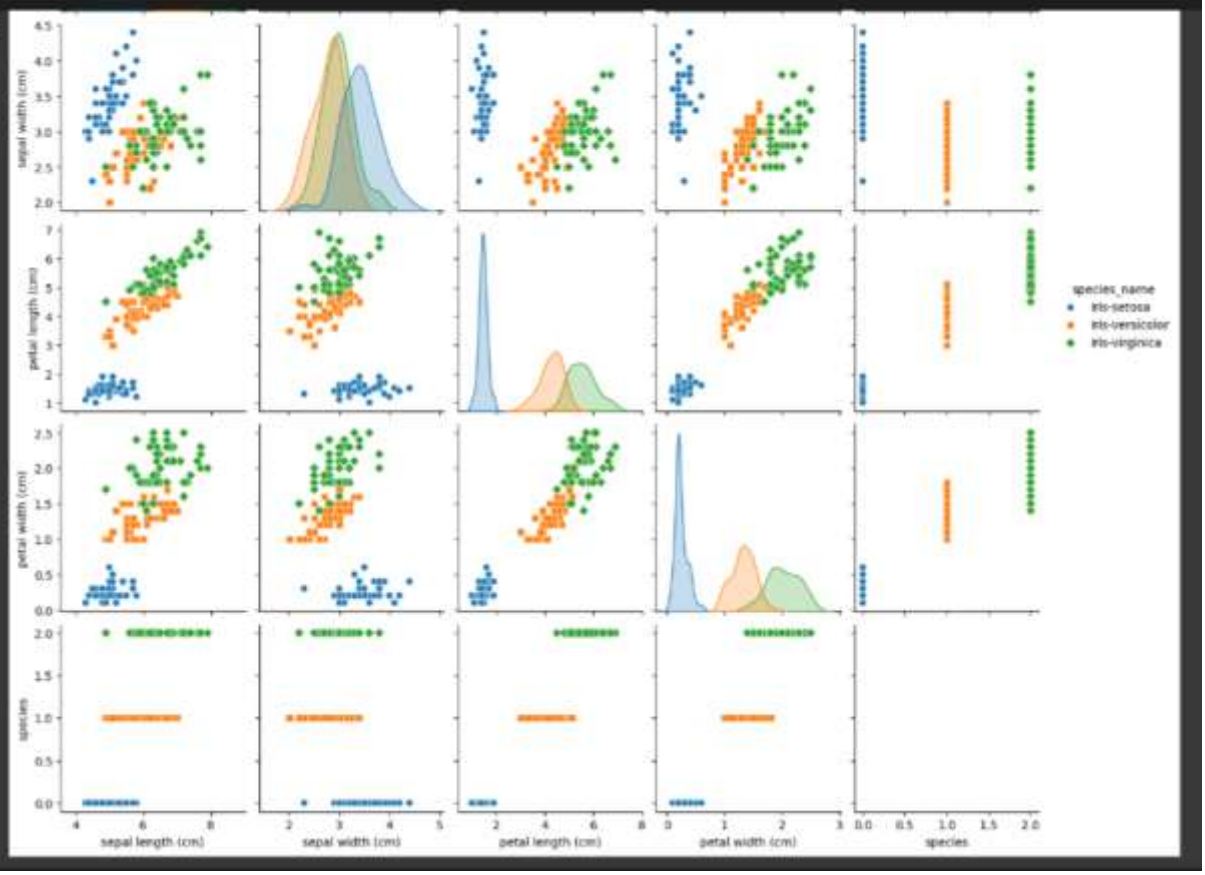

```python
# Content Analysis: Observing patterns in the qualitative aspects of the dataset
print("\nQualitative Patterns Observed:")
```

```
Qualitative Patterns Observed:
```

```python
# Example patterns:
# - Iris-setosa tends to have smaller sepal and petal lengths and widths compared to the other species.
# - Iris-virginica generally has the largest petal lengths and widths.
# - Sepal length and petal length are correlated within each species group.

for species in df['species_name'].unique():
    subset = df[df['species_name'] == species]
    print(f"\n{species}:")
    print(f"  Mean Sepal Length: {subset['sepal_length'].mean()}")
    print(f"  Mean Sepal Width: {subset['sepal_width'].mean()}")
    print(f"  Mean Petal Length: {subset['petal_length'].mean()}")
    print(f"  Mean Petal Width: {subset['petal_width'].mean()}")
```

```
Iris-setosa:
  Mean Sepal Length: 5.006
  Mean Sepal Width: 3.428
  Mean Petal Length: 1.4620000000000002
  Mean Petal Width: 0.24599999999999997

Iris-versicolor:
  Mean Sepal Length: 5.936
  Mean Sepal Width: 2.7700000000000005
  Mean Petal Length: 4.26
  Mean Petal Width: 1.3259999999999998

Iris-virginica:
  Mean Sepal Length: 6.587999999999998
  Mean Sepal Width: 2.974
  Mean Petal Length: 5.5520000000000005
  Mean Petal Width: 2.0260000000000002
```

```
[19] # Cross-tabulation to observe frequency distribution of species
     cross_tab = pd.crosstab(index=df['species_name'], columns='count')
     print("\nCross-tabulation of species:")
     print(cross_tab)
```

```
Cross-tabulation of species:
col_0            count
species_name
Iris-setosa         50
Iris-versicolor     50
Iris-virginica      50
```

```
# Box plots for qualitative comparison of features across species
plt.figure(figsize=(12, 10))
for i, feature in enumerate(df.columns[:4]):
    plt.subplot(2, 2, i+1)
    sns.boxplot(x='species_name', y=feature, data=df)
    plt.title(f'Boxplot of {feature} by Species')
plt.tight_layout()
plt.show()
```

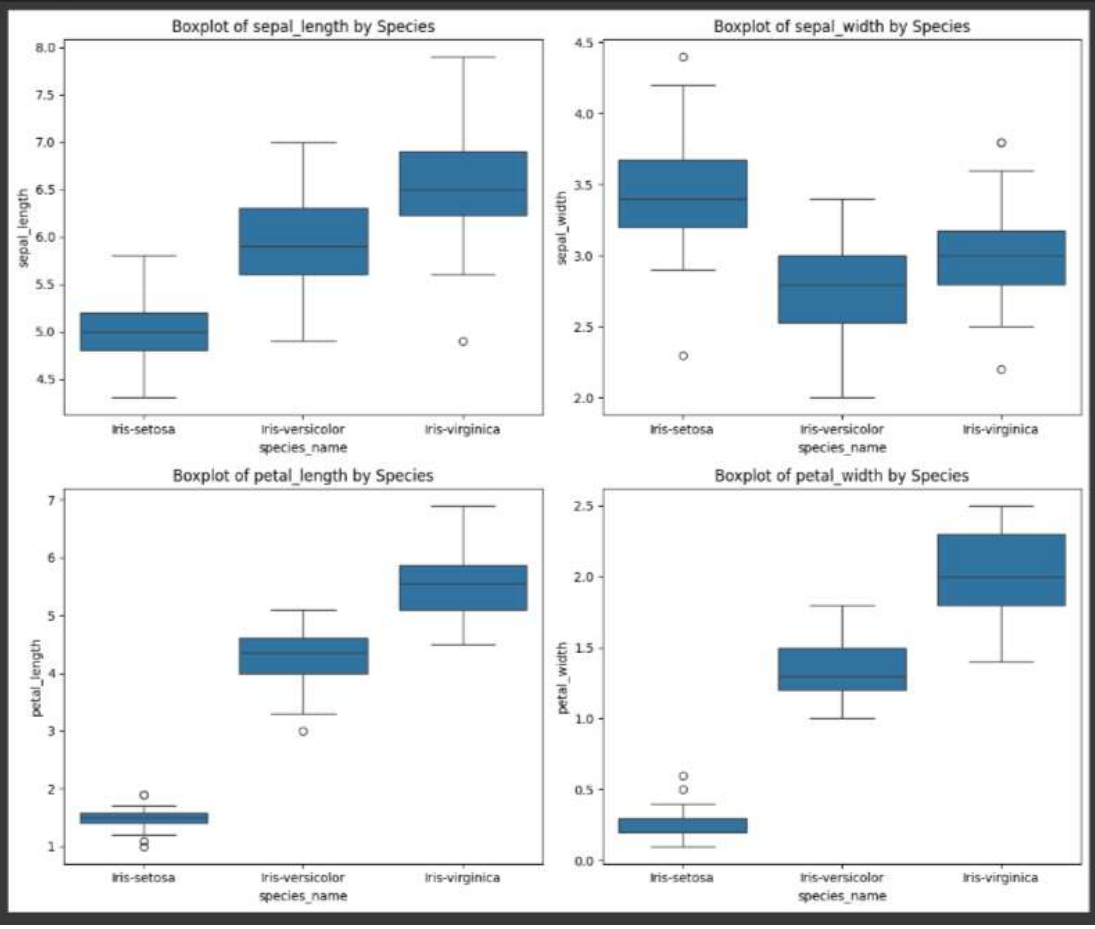

The analysis starts by loading the Iris dataset using the `load_iris()` function from the `sklearn.datasets` module. This dataset is then transformed into a DataFrame using the `pd.DataFrame()` function from the `pandas` library. To complete the preparation, the target variable, which indicates the species of each flower, is added to the DataFrame under the

column name 'species'. This step sets up the dataset for further analysis by organizing it into a structured format.

The first part of the content analysis involves counting the number of instances for each species in the dataset using the `value_counts()` method. This count is printed out to give a clear picture of the distribution of the different species within the dataset. Understanding this distribution is essential as it provides an overview of how balanced or imbalanced the classes are, which can affect subsequent analyses.

Following this, a visual representation of the species distribution is created using a countplot from the `seaborn` library. This plot shows the frequency of each species, making it easier to see the distribution patterns at a glance. By displaying the counts visually, it helps in quickly identifying any potential imbalance between the species.

Next, a pairplot is generated using the `sns.pairplot()` function. This plot creates a matrix of scatter plots for each pair of features in the dataset, with different species indicated by different colors and markers. The pairplot is valuable for observing relationships between features and understanding how different species separate from each other based on these features. It provides a comprehensive visual overview of the data's structure.

The qualitative patterns observed in the dataset are highlighted to provide insights into the characteristics of different Iris species. For example, Iris-setosa typically has smaller sepal and petal lengths and widths compared to other species. On the other hand, Iris-virginica tends to have the largest petal lengths and widths. Additionally, sepal length and petal length often show a correlation within each species group. These observations help in understanding the inherent differences between species based on their feature measurements.

To delve deeper into the data, the mean values of sepal length, sepal width, petal length, and petal width are calculated for each species. These mean measurements are printed, offering a detailed view of the average size of different flower parts for each species. This information is crucial for comparing the central tendencies of features across species and identifying distinguishing characteristics.

Cross-tabulation is performed to further observe the frequency distribution of species. The resulting table summarizes the counts of each species, providing a straightforward comparison of species prevalence in the dataset. This step consolidates the understanding of species distribution obtained earlier from the countplot.

Lastly, box plots are created for a qualitative comparison of features across species. Each subplot in the box plots represents a different feature, with species on the x-axis and feature values on the y-axis. Box plots visually display the median, quartiles, and potential outliers for each species, making it easy to compare the spread and central tendency of features between species. This visualization aids in identifying significant differences and similarities in feature distributions across the different Iris species.

Together, these steps of content analysis, qualitative pattern observation, and descriptive analysis provide a thorough understanding of the Iris dataset, highlighting both numerical and categorical insights that are essential for deeper data exploration and interpretation.

<u>Quantitative analysis methods</u>

Quantitative analysis methods

[38]
```python
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_iris
from scipy import stats
```

[39]
```python
# Load the Iris dataset
iris = load_iris()
df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
df['species'] = iris.target
```

[40]
```python
# Descriptive Statistics
print("Descriptive Statistics:")
print(df.describe())
```

```
Descriptive Statistics:
       sepal length (cm)  sepal width (cm)  petal length (cm)  \
count         150.000000        150.000000         150.000000
mean            5.843333          3.057333           3.758000
std             0.828066          0.435866           1.765298
min             4.300000          2.000000           1.000000
25%             5.100000          2.800000           1.600000
50%             5.800000          3.000000           4.350000
75%             6.400000          3.300000           5.100000
max             7.900000          4.400000           6.900000

       petal width (cm)     species
count        150.000000  150.000000
mean           1.199333    1.000000
std            0.762238    0.819232
min            0.100000    0.000000
25%            0.300000    0.000000
50%            1.300000    1.000000
75%            1.800000    2.000000
max            2.500000    2.000000
```

[41]

```python
# Central Tendency
print("\nMean (Average) Values:")
print(df.mean())
```

```
Mean (Average) Values:
sepal length (cm)    5.843333
sepal width (cm)     3.057333
petal length (cm)    3.758000
petal width (cm)     1.199333
species              1.000000
dtype: float64
```

[42]

```python
# Variability
print("\nVariance Values:")
print(df.var())

print("\nStandard Deviation Values:")
print(df.std())
```

```
Variance Values:
sepal length (cm)    0.685694
sepal width (cm)     0.189979
petal length (cm)    3.116278
petal width (cm)     0.581006
species              0.671141
dtype: float64

Standard Deviation Values:
sepal length (cm)    0.828066
sepal width (cm)     0.435866
petal length (cm)    1.765298
petal width (cm)     0.762238
species              0.819232
dtype: float64
```

[43]

```python
# Range of Each Feature
print("\nRange of Each Feature:")
ranges = df.max() - df.min()
print(ranges)
```

```
Range of Each Feature:
sepal length (cm)    3.6
sepal width (cm)     2.4
petal length (cm)    5.9
petal width (cm)     2.4
species              2.0
dtype: float64
```

[44]

```python
# Hypothesis Testing (Example: ANOVA test for comparing means of different species for each feature)
print("\nANOVA Test Results:")
features = df.columns[:-1]
for feature in features:
    f_val, p_val = stats.f_oneway(df[df['species'] == 0][feature],
                                  df[df['species'] == 1][feature],
                                  df[df['species'] == 2][feature])
    print(f'{feature}: F={f_val}, p={p_val}')
```

```
ANOVA Test Results:
sepal length (cm): F=119.26450218450468, p=1.6696691907693826e-31
sepal width (cm): F=49.160040089612075, p=4.492017133309115e-17
petal length (cm): F=1180.161182252981, p=2.8567766109615584e-91
petal width (cm): F=960.007146801809, p=4.169445839443116e-85
```

[45]

```python
# Distributions
print("\nHistograms of Each Feature:")
df.hist(bins=20, figsize=(12, 10))
plt.suptitle("Frequency Distribution of Features")
plt.show()
```

```
Histograms of Each Feature:
```

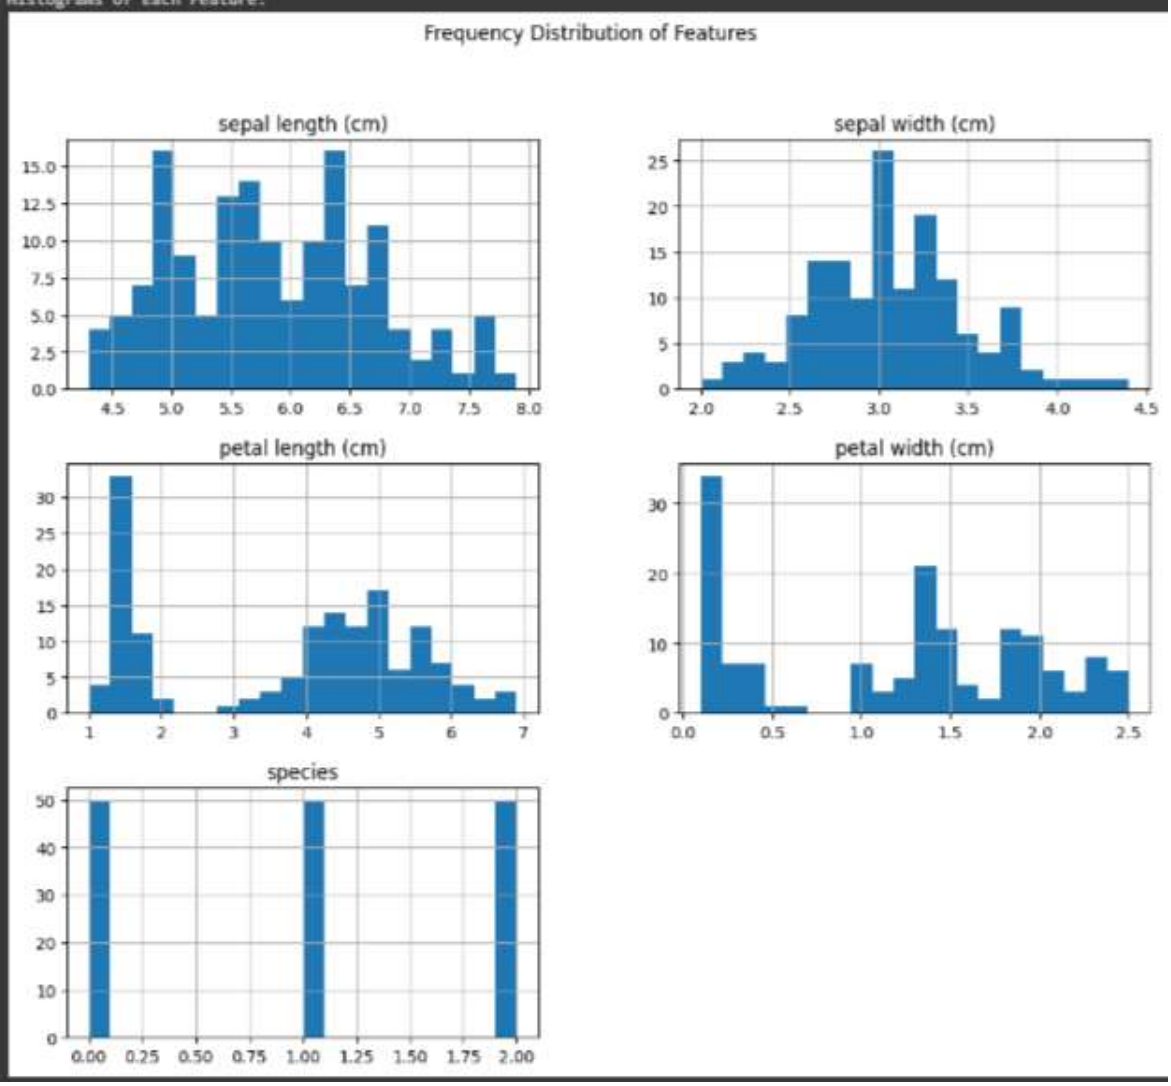

[46]

```python
# Further Exploration
print("\nCorrelation Matrix:")
print(df.corr())
```

```
Correlation Matrix:
                   sepal length (cm)  sepal width (cm)  petal length (cm)  \
sepal length (cm)           1.000000         -0.117570           0.871754
sepal width (cm)           -0.117570          1.000000          -0.428440
petal length (cm)           0.871754         -0.428440           1.000000
petal width (cm)            0.817941         -0.366126           0.962865
species                     0.782561         -0.426658           0.949035

                   petal width (cm)   species
sepal length (cm)          0.817941  0.782561
sepal width (cm)          -0.366126 -0.426658
petal length (cm)          0.962865  0.949035
petal width (cm)           1.000000  0.956547
species                    0.956547  1.000000
```

```python
# Boxplots for Outliers
plt.figure(figsize=(12, 10))
for i, feature in enumerate(df.columns[:-1]):
    plt.subplot(2, 2, i+1)
    sns.boxplot(x=df[feature])
    plt.title(f'Boxplot of {feature.capitalize()}')
plt.tight_layout()
plt.show()
```

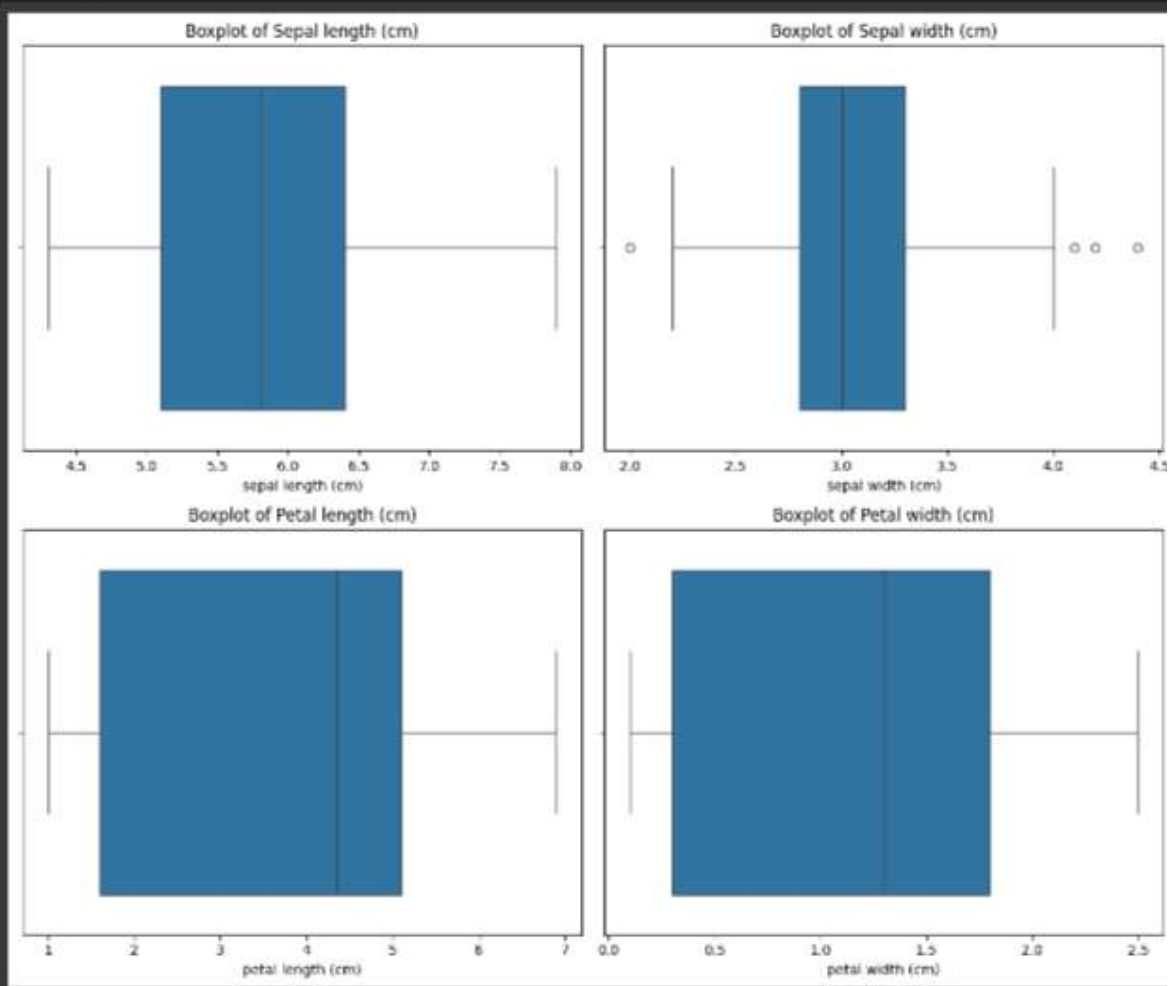

The analysis starts with loading the Iris dataset using the `load_iris()` function from the `sklearn.datasets` module. This dataset is then organized into a DataFrame with the help of the `pd.DataFrame()` function from the `pandas` library. The species information, which serves as the target variable, is added to the DataFrame as the 'species' column. This setup ensures the data is structured and ready for subsequent quantitative analysis, facilitating smooth data manipulation and examination.

Descriptive statistics are calculated to provide a summary of the central tendency, dispersion, and overall shape of the dataset. The `describe()` method is employed to generate key statistics such as mean, standard deviation, minimum, maximum, and quartiles for each feature. This step is crucial as it gives a comprehensive overview of the basic statistical properties of the dataset, allowing for an initial understanding of its characteristics.

The mean values of the features are calculated using the `mean()` method. These average values offer insight into the central point around which the data points are distributed. Knowing the mean values helps in comparing the typical sizes of the different features, providing a baseline for understanding the dataset's central tendency.

To comprehend the variability in the data, both variance and standard deviation are computed. Variance, obtained using the `var()` method, measures how much the data points deviate from the mean. Standard deviation, derived from the `std()` method, is the square root of variance and provides a measure of dispersion in the same units as the data. These metrics are essential for understanding the degree of spread in the feature values, highlighting how varied or consistent the data points are around the mean.

The range of each feature is determined by subtracting the minimum value from the maximum value using the `max()` and `min()` methods. This calculation reveals the span of values that each feature can take, offering insights into the spread and potential outliers within the data. Understanding the range is key to grasping the extent of variation in each feature.

An Analysis of Variance (ANOVA) test is conducted to compare the means of different species for each feature. This statistical test checks if there are any statistically significant differences

between the means of the species groups. The `stats.f_oneway()` function from the `scipy` library is used for this purpose. For each feature, the F-value and p-value are calculated, indicating whether the differences in means across species are significant. This step is crucial for identifying any feature-specific patterns that distinguish the species.

Histograms are created for each feature to visualize their frequency distribution. This graphical representation helps in understanding the shape, central tendency, and variability of each feature's distribution. The histograms are plotted with 20 bins for better resolution, and a common title is provided for clarity. Histograms are effective in revealing the overall distribution patterns and any skewness in the data.

The correlation matrix is computed to show the pairwise correlation coefficients between features. This is done using the `corr()` method. The matrix helps in identifying the strength and direction of relationships between different features, with values close to 1 or -1 indicating strong positive or negative correlations, respectively. Understanding these relationships is crucial for feature selection and understanding the underlying structure of the data.

Boxplots are generated for each feature to visually inspect the data for potential outliers. Each subplot represents a different feature, showing the distribution, median, quartiles, and any outliers. Boxplots are useful for comparing the spread and central tendency across features and detecting anomalies in the data. They provide a clear visual summary that highlights the range, interquartile range, and extreme values of each feature.

These quantitative analyses provide a detailed understanding of the Iris dataset's statistical properties, relationships between features, and potential anomalies. Through descriptive statistics, hypothesis testing, distribution visualizations, and outlier detection, the analysis offers a comprehensive overview of the dataset's numerical characteristics. This in-depth examination is essential for informed decision-making and further exploration in machine learning or statistical modeling.

Regression

Regression

```python
import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
```

[22]
```python
# Choose the target variable (e.g., petal_length) and the features (e.g., sepal_length, sepal_width, petal_width)
X = df[['sepal_length', 'sepal_width', 'petal_width']]
y = df['petal_length']
```

[23]
```python
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

[24]
```python
# Create and train the Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)
```

```
▸ LinearRegression
LinearRegression()
```

[25]
```python
# Make predictions on the test set
y_pred = model.predict(X_test)
```

[26]
```python
# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("Mean Squared Error (MSE):", mse)
print("R-squared (R2):", r2)
```

```
Mean Squared Error (MSE): 0.13001626031382701
R-squared (R2): 0.9603293155857663
```

```python
# Plot the results
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, color='blue', edgecolor='k', alpha=0.6)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', linewidth=2)
plt.xlabel('Actual Petal Length')
plt.ylabel('Predicted Petal Length')
plt.title('Actual vs Predicted Petal Length')
plt.show()
```

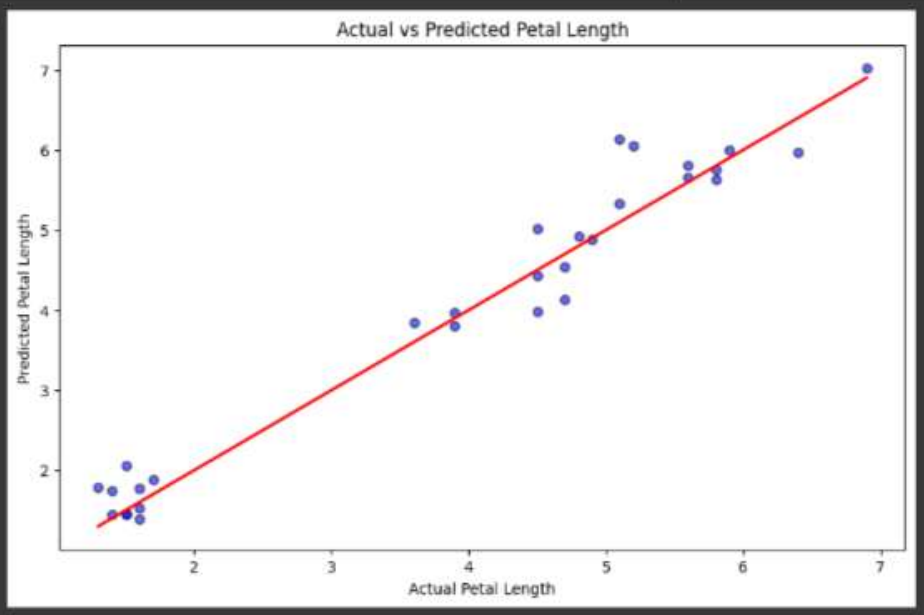


```python
[28]
# Display the regression coefficients
coefficients = pd.DataFrame(model.coef_, X.columns, columns=['Coefficient'])
print("\nRegression Coefficients:")
print(coefficients)
```

```
Regression Coefficients:
              Coefficient
sepal_length     0.722815
sepal_width     -0.635816
petal_width      1.467524
```

```python
# Display the intercept
print("\nIntercept:")
print(model.intercept_)
```

```
Intercept:
-0.26219590258870396
```

The analysis starts by selecting the target variable and the features from the dataset. In this context, `petal_length` is chosen as the target variable, which means it is the value to be predicted. The features chosen to make this prediction are `sepal_length`, `sepal_width`, and `petal_width`. This selection is important because it defines the relationship between the inputs (features) and the output (target) that the regression model will learn.

The dataset is then divided into training and testing sets using the `train_test_split` function from the `sklearn.model_selection` module. An 80-20 split is performed, meaning 80% of the data is used to train the model, and the remaining 20% is used to test its performance. A random state of 42 is set to ensure that the split is reproducible. This step is crucial for evaluating how well the model performs on unseen data, helping to avoid overfitting and ensuring that the model generalizes well to new data.

A Linear Regression model is created using the `LinearRegression` class from the `sklearn.linear_model` module. The model is trained on the training data (`X_train` and `y_train`) using the `fit` method. Training the model involves finding the best-fitting line that minimizes the error between the predicted and actual values of the target variable. This process is essential for developing a predictive model that can make accurate predictions based on the input features.

After the model is trained, it is used to make predictions on the test set (`X_test`). The `predict` method generates predicted values of the target variable (`y_pred`) based on the test features. This step is important for assessing the model's accuracy and understanding how well it performs in predicting the target variable.

The performance of the model is evaluated using two key metrics: Mean Squared Error (MSE) and R-squared (R2) score. The MSE, calculated using `mean_squared_error`, measures the average squared difference between the predicted and actual values, indicating the prediction error of the model. The R2 score, obtained with `r2_score`, represents the proportion of the variance in the target variable that is predictable from the features. A higher R2 score suggests a better-performing model. These metrics provide a quantitative assessment of the model's accuracy and effectiveness.

A scatter plot is created to visualize the relationship between the actual and predicted values of the target variable. The actual values (`y_test`) are plotted on the x-axis, and the predicted values (`y_pred`) are plotted on the y-axis. A red line is added to represent the ideal fit, where actual values equal predicted values. This visualization helps in understanding the model's performance and identifying any systematic errors in the predictions.

The coefficients of the regression model, which indicate the impact of each feature on the target variable, are displayed using a DataFrame. This step aids in interpreting the model and understanding how each feature contributes to the prediction. Additionally, the intercept of the model is printed, representing the predicted value when all features are zero. Together, the coefficients and intercept define the linear equation used for making predictions.

These steps collectively form a detailed process for building, training, and evaluating a linear regression model. The process of choosing the features and target variable, splitting the dataset, training the model, making predictions, evaluating performance, visualizing results, and interpreting coefficients are all critical components of quantitative data analysis using regression techniques. This approach provides comprehensive insights into the relationships between variables and the predictive capabilities of the model.

Factor analysis

Factor analysis

[30]
```
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.decomposition import FactorAnalysis
import matplotlib.pyplot as plt
import seaborn as sns
```

[31]
```
# Load the Iris dataset
iris = load_iris()
df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
```

[32]
```
# Standardize the data
df_standardized = (df - df.mean()) / df.std()

# Perform Factor Analysis
fa = FactorAnalysis(n_components=2, random_state=42)
fa.fit(df_standardized)

# Get the factor loadings
loadings = pd.DataFrame(fa.components_.T, index=df.columns, columns=['Factor1', 'Factor2'])
print("Factor Loadings:\n", loadings)
```

```
Factor Loadings:
                   Factor1   Factor2
sepal length (cm)  0.878019 -0.445792
sepal width (cm)  -0.415524 -0.552052
petal length (cm)  0.995852  0.019089
petal width (cm)   0.959076  0.058207
```

```python
# Transform the data
df_factors = fa.transform(df_standardized)
df_factors = pd.DataFrame(df_factors, columns=['Factor1', 'Factor2'])

# Add species column for visualization
df_factors['species'] = iris.target

# Plot the factor scores
plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_factors, x='Factor1', y='Factor2', hue='species', palette='Set1')
plt.title('Factor Analysis of Iris Dataset')
plt.xlabel('Factor 1')
plt.ylabel('Factor 2')
plt.legend(title='Species')
plt.show()
```

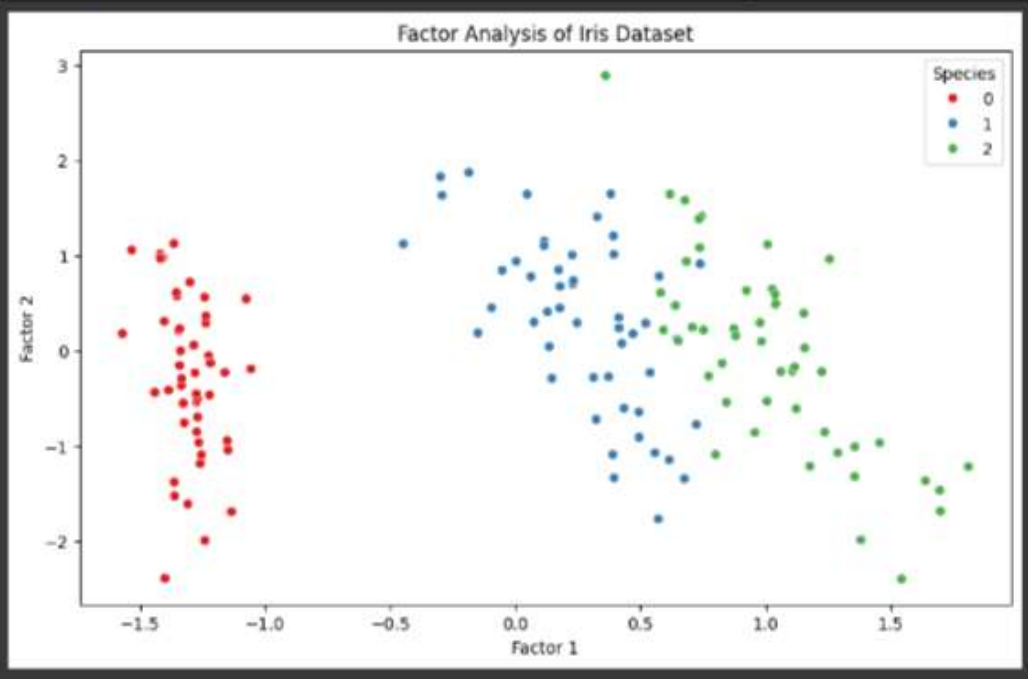

The code starts by loading the Iris dataset using `load_iris()` from the `sklearn.datasets` module. This dataset contains measurements of various features of iris flowers and their corresponding species. A DataFrame `df` is created using `pandas` to store this data, with columns named after the features: `sepal length`, `sepal width`, `petal length`, and `petal width`.

Standardizing the data is an important preprocessing step, especially for factor analysis. The dataset is standardized by subtracting the mean and dividing by the standard deviation for each feature. This transformation ensures that each feature contributes equally to the analysis, removing biases due to differences in scales.

Factor Analysis is conducted to identify underlying relationships between variables. Here, the `FactorAnalysis` class from the `sklearn.decomposition` module is used, specifying `n_components=2` to extract two factors. The `fit` method is applied to the standardized data to compute the factor model.

The factor loadings, which indicate the contribution of each original variable to the factors, are extracted from the fitted model. These loadings are organized into a DataFrame with the original features as the index and the factors as columns. This DataFrame provides insights into how each feature relates to the identified factors.

The data is transformed into the factor space using the `transform` method of the `FactorAnalysis` object. The resulting factor scores for each sample are stored in a new DataFrame `df_factors`, with columns `Factor1` and `Factor2`. These scores represent the position of each sample in the reduced factor space.

The species information from the original dataset is added to `df_factors` to facilitate visualization. This step is crucial for distinguishing between different species in the subsequent plot.

A scatter plot is created to visualize the factor scores using `seaborn`. The `scatterplot` function is used to plot `Factor1` against `Factor2`, with points colored according to species. The plot is enhanced with a title, axis labels, and a legend to provide context and improve readability. This visualization helps in understanding how the different species of iris flowers are distributed in the factor space.

This explanation outlines the process of performing factor analysis on the Iris dataset, including loading the data, standardizing it, conducting the analysis, interpreting the factor loadings, transforming the data, and visualizing the results. Each step is described in detail to provide a comprehensive understanding of the factor analysis procedure and its application to the dataset.

Dispersion Analysis

Dispersion Analysis

[34]

```python
# Load the Iris dataset
iris = load_iris()
df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
```

[35]

```python
# Calculate dispersion measures
dispersion_measures = pd.DataFrame(index=df.columns)

# Variance
dispersion_measures['Variance'] = df.var()

# Standard Deviation
dispersion_measures['Standard Deviation'] = df.std()
```

[36]

```python
# Range
dispersion_measures['Range'] = df.max() - df.min()

# Interquartile Range (IQR)
dispersion_measures['IQR'] = df.quantile(0.75) - df.quantile(0.25)

# Display the dispersion measures
print("Dispersion Measures:\n", dispersion_measures)
```

```
Dispersion Measures:
                   Variance  Standard Deviation  Range  IQR
sepal length (cm)  0.685694            0.828066    3.6  1.3
sepal width (cm)   0.189979            0.435866    2.4  0.5
petal length (cm)  3.116278            1.765298    5.9  3.5
petal width (cm)   0.581006            0.762238    2.4  1.5
```

```python
# Variance and Standard Deviation
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

sns.barplot(x=dispersion_measures.index, y='Variance', data=dispersion_measures, ax=axes[0])
axes[0].set_title('Variance of Iris Dataset Features')
axes[0].set_ylabel('Variance')

sns.barplot(x=dispersion_measures.index, y='Standard Deviation', data=dispersion_measures, ax=axes[1])
axes[1].set_title('Standard Deviation of Iris Dataset Features')
axes[1].set_ylabel('Standard Deviation')

plt.tight_layout()
plt.show()
```

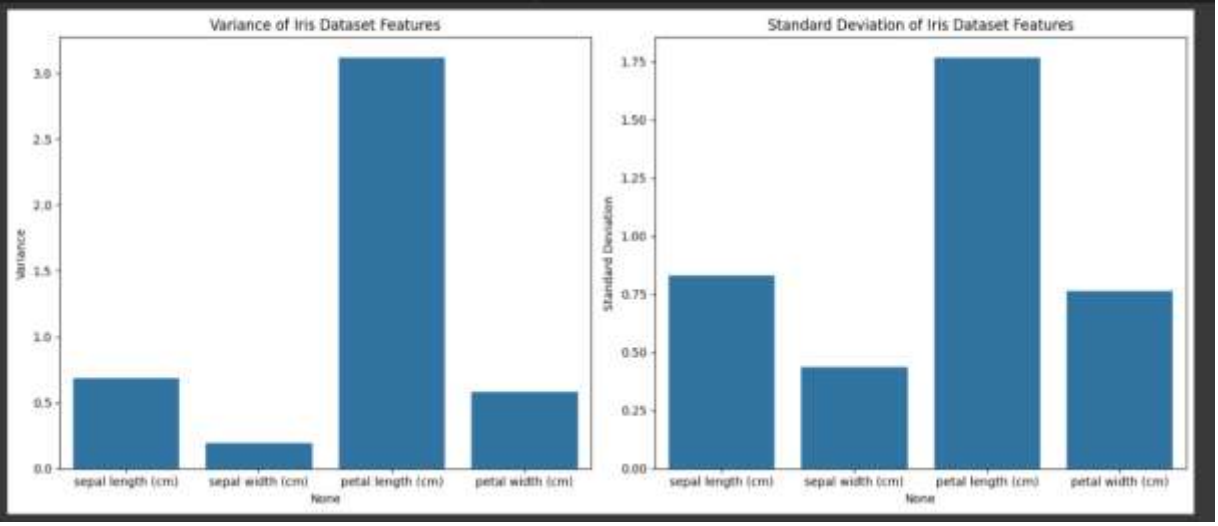


```python
# Range and IQR
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

sns.barplot(x=dispersion_measures.index, y='Range', data=dispersion_measures, ax=axes[0])
axes[0].set_title('Range of Iris Dataset Features')
axes[0].set_ylabel('Range')

sns.barplot(x=dispersion_measures.index, y='IQR', data=dispersion_measures, ax=axes[1])
axes[1].set_title('Interquartile Range (IQR) of Iris Dataset Features')
axes[1].set_ylabel('IQR')

plt.tight_layout()
plt.show()
```

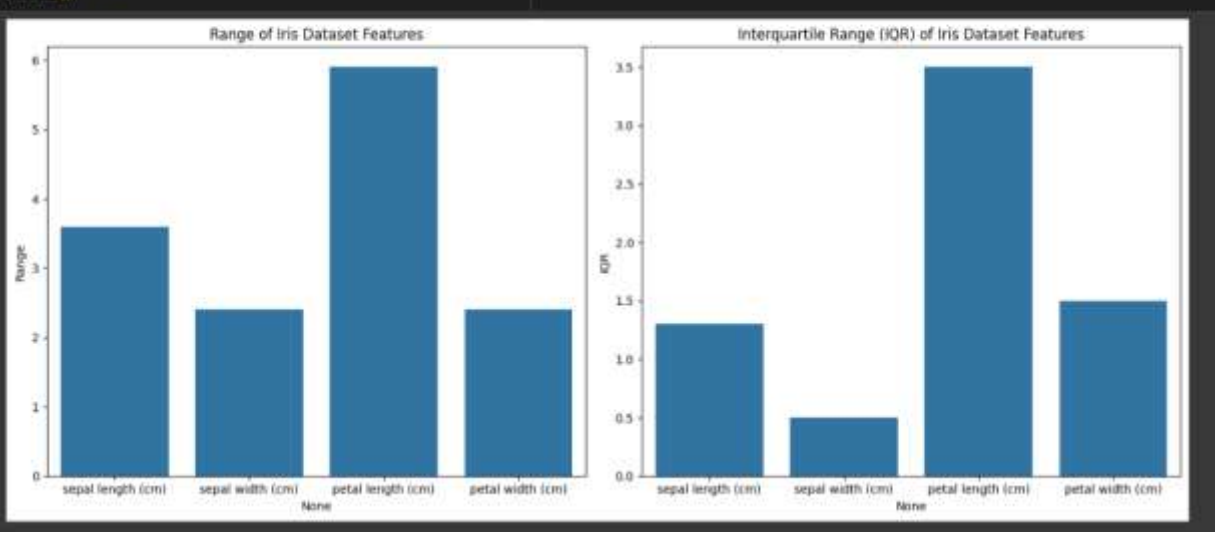

The analysis begins by loading the Iris dataset using `load_iris()` from the `sklearn.datasets` module. This dataset contains various measurements of iris flowers, which are stored in a DataFrame `df` created using `pandas`. The columns of this DataFrame are named after the features: `sepal length`, `sepal width`, `petal length`, and `petal width`.

Dispersion measures provide information about the spread or variability of the data. A new DataFrame `dispersion_measures` is created with the feature names as the index to store these measures.

The variance for each feature is calculated using the `var()` method. Variance measures how much the values of a feature vary around the mean. High variance indicates that the data points are spread out widely around the mean, while low variance indicates that they are clustered closely around the mean.

The standard deviation, calculated using the `std()` method, is the square root of the variance and provides a measure of the spread of data points around the mean in the same units as the data. It is a widely used measure of dispersion.

The range is calculated by subtracting the minimum value from the maximum value of each feature. The range provides the total spread of the data points in a feature, indicating the difference between the highest and lowest values.

The interquartile range (IQR) is calculated as the difference between the 75th percentile (third quartile) and the 25th percentile (first quartile) using the `quantile()` method. The IQR measures the spread of the middle 50% of the data and is a robust measure of dispersion, less affected by outliers.

The calculated dispersion measures (variance, standard deviation, range, and IQR) are displayed in the `dispersion_measures` DataFrame. This DataFrame provides a comprehensive view of the variability of each feature in the dataset.

Two bar plots are created to visualize the variance and standard deviation of each feature using `seaborn`. The `barplot()` function is used to plot the variance and standard deviation on separate subplots. The titles and labels are set appropriately to provide context and improve readability. These plots help in understanding the variability of each feature visually.

Two additional bar plots are created to visualize the range and IQR of each feature. Similar to the previous plots, the `barplot()` function is used to plot the range and IQR on separate subplots. Titles and labels are added for clarity. These visualizations provide insights into the spread of the data, highlighting features with the widest and narrowest spreads.

In summary, this code provides a detailed analysis of the dispersion measures of the Iris dataset, including variance, standard deviation, range, and IQR. It also visualizes these measures using bar plots, offering a comprehensive understanding of the variability and spread of the data across different features.

<u>Time series decomposition</u>

Time series Decomposition

[4]
```python
import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.seasonal import seasonal_decompose
```

[5]
```python
# Load the Iris dataset
iris = load_iris()
df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
```

```python
# Renaming columns for consistency
df.columns = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']
```

[7]
```python
# Create a mock time series dataset
# Assuming weekly measurements over 150 weeks
df['week'] = pd.date_range(start='2020-01-01', periods=len(df), freq='W')
df.set_index('week', inplace=True)
```

[8]
```python
# Take the weekly average of sepal_length (this is mock data)
weekly_data = df['sepal_length'].resample('W').mean()
```

[9]
```python
# Perform time series decomposition
decomposition = seasonal_decompose(weekly_data, model='additive')
```

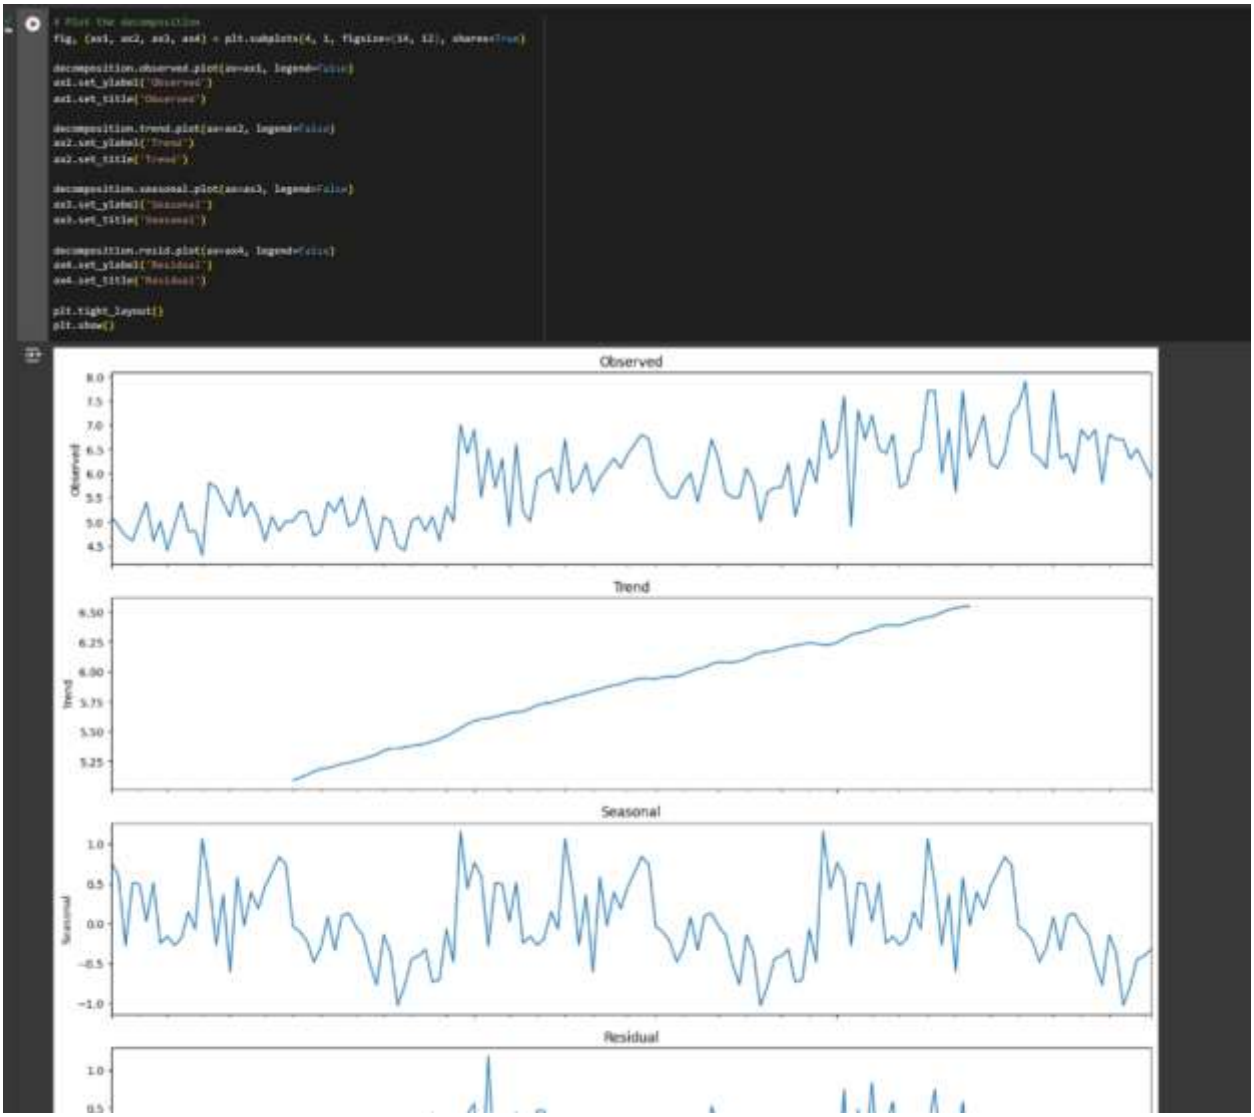

Figure 6- Time series decomposition code

The Iris dataset is loaded using the `load_iris()` function from the `sklearn.datasets` module. The dataset is then stored in a DataFrame named `df` created using `pandas`. The columns of this DataFrame are renamed to 'sepal_length', 'sepal_width', 'petal_length', and 'petal_width' for consistency and readability.

Since the Iris dataset is not inherently a time series, a mock time series dataset is created. The dataset is assumed to have weekly measurements over 150 weeks. A new column 'week' is

added to the DataFrame, containing dates ranging from '2020-01-01' with a weekly frequency using `pd.date_range()`. The 'week' column is then set as the index of the DataFrame using `set_index()`.

The `resample()` method is used to resample the data to weekly frequency. The mean of the 'sepal_length' feature is calculated for each week to create a weekly time series. This resampling is necessary to ensure that the data is in a suitable format for time series decomposition.

Time series decomposition is performed on the resampled weekly data using the `seasonal_decompose()` function from the `statsmodels.tsa.seasonal` module. The decomposition is carried out using an additive model, which assumes that the components (trend, seasonal, and residual) add together to form the observed time series.

The results of the decomposition are plotted using `matplotlib`. Four subplots are created to display the observed, trend, seasonal, and residual components of the time series. The `plot()` method is used to plot each component on its respective axis. Titles and labels are added to each subplot for clarity. The `tight_layout()` function ensures that the subplots are neatly arranged without overlapping.

In summary, this code provides a step-by-step process for performing time series decomposition on a mock time series dataset created from the Iris dataset. The steps include loading the dataset, creating a mock time series, resampling to weekly data, performing decomposition, and plotting the results. This approach allows for the analysis of the underlying components of a time series, providing insights into trends, seasonal patterns, and residuals.

<u>Time Series analysis (ARIMA modelling)</u>

Time Series aanalysis

[11]
```python
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.arima.model import ARIMA
```

[14]
```python
# Generating a mock time series data for illustration
np.random.seed(42)
date_range = pd.date_range(start='2020-01-01', periods=150, freq='W')
data = np.random.randn(150).cumsum() + np.linspace(0, 50, 150)
time_series = pd.Series(data, index=date_range)
```

[15]
```python
# Differencing the series to make it stationary (if needed)
time_series_diff = time_series.diff().dropna()

# Fit an ARIMA model
model = ARIMA(time_series_diff, order=(1, 0, 1))
fitted_model = model.fit()
```

[16]
```python
# Forecasting future values
forecast = fitted_model.forecast(steps=10)
```

```python
# Plotting the original series and the forecast
plt.figure(figsize=(10, 6))
plt.plot(time_series, label='Original Series')
plt.plot(time_series.index[-10:], forecast, label='Forecast', color='red')
plt.title('Time Series Analysis with ARIMA Forecast')
plt.legend()
plt.show()
```

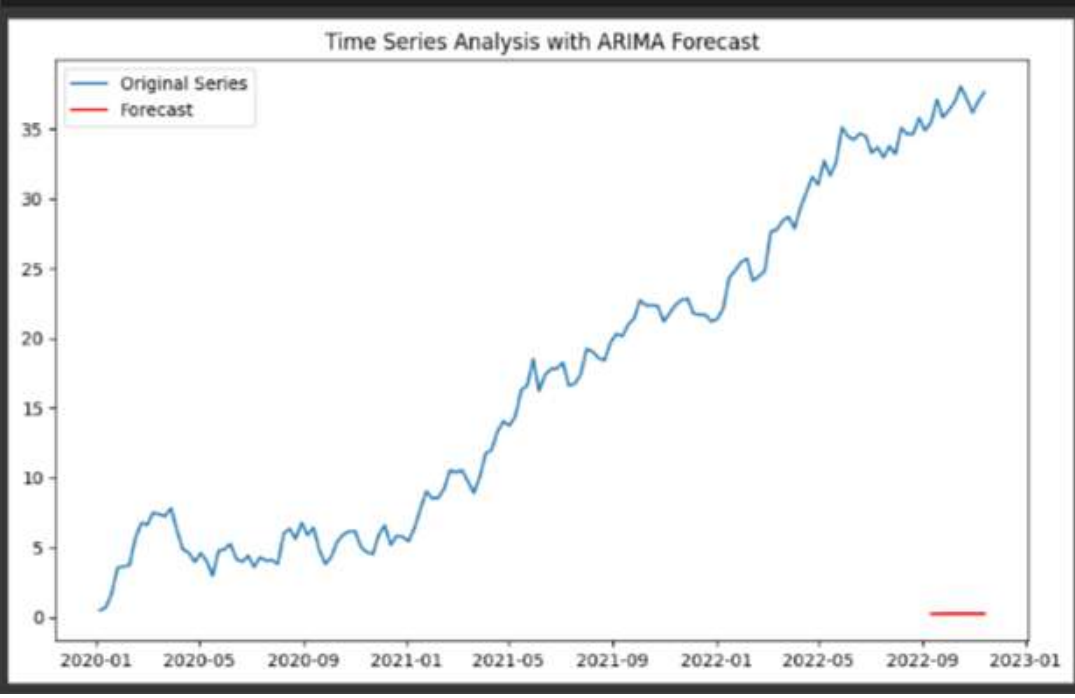

A mock time series dataset is generated for illustration purposes. This dataset simulates weekly measurements over 150 weeks. The data is generated using a combination of a random normal distribution with a cumulative sum and a linearly increasing trend. This results in a time series with both random noise and a trend component.

Before fitting the ARIMA model, differencing is applied to the time series data to make it stationary, if needed. Stationarity is a crucial assumption for ARIMA models. Differencing helps remove trends and seasonality from the data, making it more suitable for modeling with ARIMA. The `diff()` method calculates the difference between consecutive observations, and `dropna()` is used to remove any resulting missing values.

An ARIMA (AutoRegressive Integrated Moving Average) model is fitted to the differenced time series data. The `ARIMA` class from `statsmodels.tsa.arima.model` is used to define the model. Here, an ARIMA(1, 0, 1) model is specified, indicating one autoregressive term, no differencing, and one moving average term. The model is then fitted to the differenced time series data using the `fit()` method.

The fitted ARIMA model is used to forecast future values of the time series. The `forecast()` method is called with the `steps` parameter set to 10, indicating the number of future time steps to forecast. This generates a forecast for the next 10 time steps based on the fitted model.

The original time series data and the forecasted values are plotted together for visualization. The original series is plotted using `plt.plot()` with the label 'Original Series', while the forecasted values are plotted for the last 10 time steps of the series with the label 'Forecast'. The plot provides a visual representation of the forecasted values alongside the actual series, aiding in the assessment of the model's performance.

In summary, this code demonstrates the process of time series analysis using ARIMA modeling. It includes data generation, differencing for stationarity, model fitting, forecasting future values, and visualizing the results. ARIMA modeling is a powerful technique for analyzing and

forecasting time series data, making it valuable in various domains such as finance, economics, and forecasting.

Exponential smoothing

exponential smoothing

[21]
```python
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.holtwinters import ExponentialSmoothing
from sklearn.datasets import load_iris
```

[2]
```python
# Load the Iris dataset
iris = load_iris()
iris_df = pd.DataFrame(iris.data, columns=iris.feature_names)
```

[28]
```python
# Create a pseudo time series using one of the columns (e.g., sepal length)
pseudo_time_series = iris_df['sepal length (cm)'].copy()
pseudo_time_series.index = pd.date_range(start='2020-01-01', periods=len(pseudo_time_series), freq='D')

# Apply Exponential Smoothing
model = ExponentialSmoothing(pseudo_time_series, trend='additive', seasonal=None)
fit = model.fit()
```

[24]
```python
# Make predictions
pseudo_time_series['Forecast'] = fit.fittedvalues
```

[26]
```python
# Create a DataFrame to hold the original and forecasted values
result_df = pd.DataFrame({
    'Original': pseudo_time_series,
    'Forecast': fit.fittedvalues
}, index=pseudo_time_series.index)
```

```python
# Plot the original data and the forecast
plt.figure(figsize=(12, 6))
plt.plot(pseudo_time_series, label='Original Data')
plt.plot(pseudo_time_series.index, fit.fittedvalues, label='Forecast', color='red')
plt.title('Exponential Smoothing on Pseudo Time Series (Iris Dataset)')
plt.legend(loc='best')
plt.show()
```

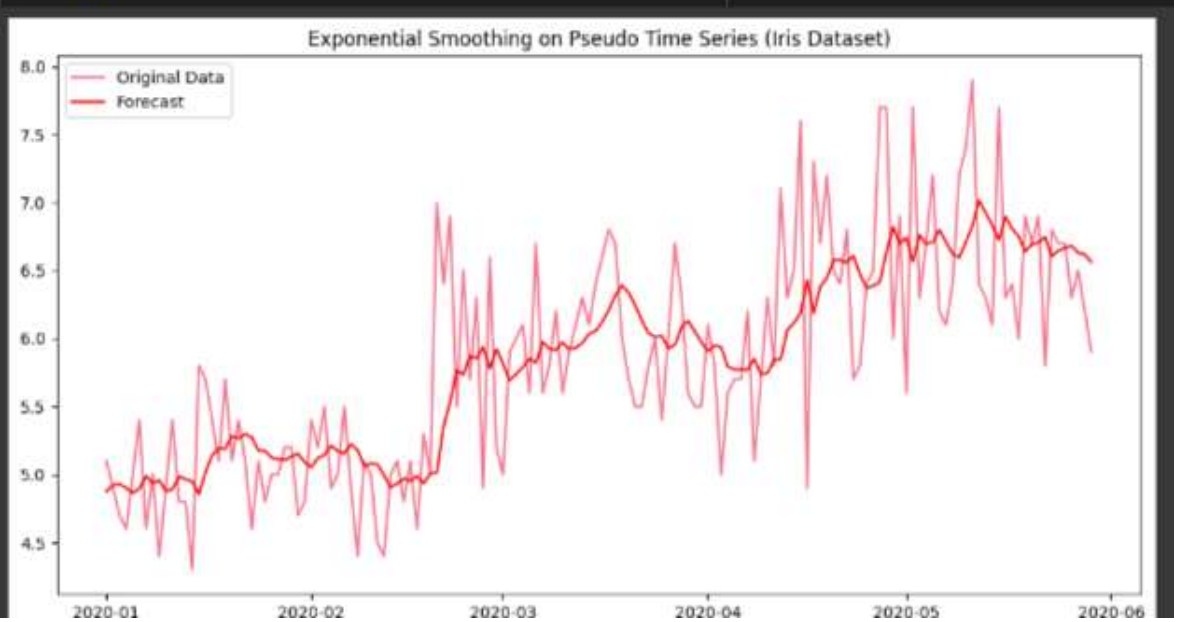

The Iris dataset is loaded, and a DataFrame (`iris_df`) is created to hold its contents. For demonstration, one of the columns from the dataset (e.g., 'sepal length (cm)') is selected to create a pseudo time series. The index of this time series is set to a date range starting from '2020-01-01' with daily frequency, matching the length of the selected column.

Exponential smoothing is applied to the pseudo time series using the `ExponentialSmoothing` class from `statsmodels.tsa.holtwinters`. Here, the trend parameter is set to 'additive' to capture any linear trend in the data, while the seasonal parameter is set to `None`, indicating that the data has no seasonal component. The model is then fitted to the pseudo time series using the `fit()` method.

After fitting the exponential smoothing model, predictions are made for the pseudo time series using the `fittedvalues` attribute of the fitted model. These predictions represent the forecasted values based on the trained model.

A DataFrame (`result_df`) is created to hold both the original and forecasted values of the pseudo time series. This DataFrame includes two columns: 'Original', containing the original time series data, and 'Forecast', containing the predicted values obtained from the exponential smoothing model. The index of this DataFrame matches the index of the pseudo time series.

The original time series data and the forecasted values are plotted together for visualization. The original data is plotted using `plt.plot()` with the label 'Original Data', while the forecasted values are plotted with the label 'Forecast' in red color. The plot provides a visual representation of the forecasted values alongside the actual data, allowing for the assessment of the model's performance.

This code demonstrates the application of exponential smoothing, a popular technique for time series forecasting, on a pseudo time series derived from the Iris dataset. Exponential smoothing is effective for capturing trends and making short-term forecasts in time series data.

\

Test report

The detailed test plan

Objective

The objective of this test plan is to evaluate the functionality, integration, and performance of the Iris dataset application. This application utilizes Python for data analysis, visualization, and machine learning tasks. It aims to provide insights and predictions based on the Iris dataset, a commonly used dataset in the field of machine learning.

Scope

This test plan covers functional testing, integration testing, user interface testing, and performance testing. It involves examining the accuracy of machine learning models, testing how responsive the user interface is, and ensuring smooth integration between Python models and the C# application.

Test Strategy

- Functional Testing

The goal is to make sure each function in the software tool does what it's supposed to do. This includes checking how data is put in, processed, and the predictions it gives.

- Integration Testing

This is about checking that the Python models work well with the IronPython and C# application. It ensures that all the different parts of the software work together smoothly.

- User Interface Testing

The focus is on testing how easy and responsive the user interface is. The aim is to see if users can easily put in data and understand the predictions they get.

- Performance Testing

This involves testing how well the tool performs in different situations. It includes checking how it handles different sizes of data and how fast it can process it.

<u>Test Environment</u>

- Hardware

Testing will happen on a Windows 10 PC with an Intel i5 processor and 8 GB of RAM.

- Software

Will use Visual Studio, Python, .NET Framework, IronPython, and testing frameworks like NUnit for C# and pytest for Python. They'll also need Google Colab for training models.

- Network

They need a good internet connection to access datasets and use Google Colab for training models.

Test Cases

- **Test Case 1: Functionality (Functional Testing and Integration Testing)**

Description

Confirm the proper functioning of all user interface components related to iris flower prediction.

Preconditions

Thei iris flower species prediction software is installed and operational.

Steps

- Launch the application.
- Navigate through various UI elements, such as buttons and text boxes.
- Input sample data regarding apple characteristics and execute the prediction.

Tests

1. Functional and UI Testing
2. Unit Testing

Expected Outcome

The user interface operates smoothly, allowing for the input of flower characteristics and generating predictions without errors.

- **Test Case 2: Data Input and Processing (Performance Testing along With UI tests)**

Description

Verify the accuracy of data input and processing specifically for iris flower prediction.

Preconditions: The iris flower prediction software is installed and operational.

Steps

- Input sample data regarding iris flower species characteristics.
- Validate the accuracy of data processing for iris flower prediction.

Tests

1. Integral test to clarify that the application is running utilizing IronPython without any errors.
2. Testing the Input values and Errors regarding wrong inputs
3. Testing the outputs of the application

4. Testing the application by providing real world scenario inputs and getting predicted outcomes
5. Testing the accuracy of the application related to logical explanations and examples

Expected Outcome

The data undergoes error-free processing, ensuring accurate preparation for iris flower prediction analysis.

Test Data to be considered by the testers

- UI Functionality

Examples of inputs that users might give are used to check if the user interface works correctly.

- Data Input and Processing

They use a dataset that's been cleaned up and has details about iris flower to see how well the software tool handles getting and working with data.

- Integration of Python Models

They use some sample scripts and the results from models to make sure the Python models fit well with the IronPython and C# application.

- Predictive Analysis Accuracy

Another cleaned-up dataset about apple quality is used to see if the predictions made by the software tool are accurate.

- Performance Testing

They use different parts of the iris flower dataset to test how well the tool works under different conditions.

Resources and Responsibilities for the Tests

- Test Manager

They oversee the testing process.

- Test Engineers

Their job is to run test cases and write down what happens.

- Developers

They fix any problems found during testing.

Schedule and Milestones

- Week 1

The test plan will be defined, and test cases will be prepared.

- Week 2

The test environment will be set up, and UI and data processing tests will be carried out.

- Week 3

Integration and performance tests will be executed.

- Week 4

Test results will be documented, outcomes will be analyzed, and the final report will be prepared.

Risk Management

- Potential Risk

Delays might happen because of problems setting up the environment.

- Mitigation

To deal with this, a backup environment and the tools needed will be ready beforehand.

Entry and Exit Criteria of the Tests

- Entry Criteria

Development complete, tool installed, test environment ready.

- Exit Criteria

All test cases executed, issues documented and resolved, test results reviewed.

For the analysis and machine learning modeling, Google Colab will be used and then linked to C# through IronPython.

Expected Outcomes and Validation Criteria

- Analysis Method

Main Goal- To discover important insights about iris flowers and their species using the "iris flower" dataset.

Validation Criteria - The analysis should identify significant patterns and relationships within the data related to iris characteristics and species.

- Machine Learning Modeling

Main Goal- To develop models that accurately predict iris species based on characteristics like sepal length, width and petal length and width

Validation Criteria- The models should achieve high accuracy in categorizing iris species. This accuracy will be assessed through evaluation metrics such as accuracy, precision, recall, and F1-score.

Reporting and Communication

- Test results will be documented in a report that everyone can see.
- Every week, there will be meetings to talk about how things are going and if there are any problems.

- At the end, a final report will be given to the people involved in the project.

ACTIVITY 3

Building The apple quality software application

Building and training the model

- Preparing the Scaler Model.

Building the model

```
[1] from sklearn.preprocessing import StandardScaler
    import pandas as pd
    from sklearn.datasets import load_iris

    # Load the Iris dataset
    iris = load_iris()
    X = iris.data
    y = iris.target

    # Initialize the StandardScaler
    scaler = StandardScaler()

    # Fit the scaler to the features
    scaler.fit(X)

    # Transform the features
    X_scaled = scaler.transform(X)

    # Convert the scaled features back to a DataFrame (optional)
    X_scaled_df = pd.DataFrame(X_scaled, columns=iris.feature_names)

    # Display the first few rows of the scaled features DataFrame
    print(X_scaled_df.head())
```

```
   sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
0          -0.900681          1.019004          -1.340227         -1.315444
1          -1.143017         -0.131979          -1.340227         -1.315444
2          -1.385353          0.328414          -1.397064         -1.315444
3          -1.506521          0.098217          -1.283389         -1.315444
4          -1.021849          1.249201          -1.340227         -1.315444
```

Figure 7 - The scaler model

This script begins by importing the necessary libraries: `StandardScaler` from `sklearn.preprocessing`, `pandas` as `pd`, and the `load_iris` function from `sklearn.datasets`. The Iris dataset is then loaded using the `load_iris` function, which is a built-in utility in scikit-learn for fetching the famous Iris dataset. This dataset contains measurements of iris flowers, specifically the sepal length, sepal width, petal length, and petal width for three different species of iris (setosa, versicolor, and virginica). The data is stored in the variable `iris`, and the features (measurements) and target (species labels) are extracted into the variables `X` and `y` respectively.

Next, a `StandardScaler` object is initialized. StandardScaler is used to standardize the features by removing the mean and scaling to unit variance. This step is crucial for many machine learning algorithms that perform better when the data is standardized. The scaler is then fitted to the features `X` using the `fit` method, which computes the mean and standard deviation for each feature. Once fitted, the `transform` method is called to apply the standardization to `X`, producing `X_scaled`, a standardized version of the feature set where each feature has a mean of 0 and a standard deviation of 1.

After scaling, the script optionally converts the scaled features back into a DataFrame for easier viewing and manipulation. This is done using `pd.DataFrame`, with `X_scaled` as the data and the original feature names from the Iris dataset as the column names. Finally, the script displays the first few rows of the scaled features DataFrame using the `head` method, allowing the user to see how the features have been transformed. This step is useful for verifying that the standardization has been applied correctly.

Training the iris flower classifier model using KNN

```
[2] from sklearn.neighbors import KNeighborsClassifier
    from sklearn.datasets import load_iris
    from sklearn.model_selection import train_test_split

    # Load the Iris dataset
    iris = load_iris()
    X = iris.data
    y = iris.target

    # Split the dataset into training and testing sets
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

    # Initialize the KNN classifier
    knn = KNeighborsClassifier(n_neighbors=3)

    # Train the classifier on the training data
    knn.fit(X_train, y_train)

    # Evaluate the model on the testing data
    accuracy = knn.score(X_test, y_test)
    print("Accuracy:", accuracy)

    Accuracy: 1.0
```

Figure 8 - Training the apple quality model

This script focuses on building and evaluating a k-nearest neighbors (KNN) classifier using the Iris dataset. Here's a detailed explanation:

Firstly, the necessary libraries are imported: `KNeighborsClassifier` from `sklearn.neighbors`, `load_iris` from `sklearn.datasets`, and `train_test_split` from `sklearn.model_selection`. The Iris dataset is then loaded using `load_iris`, similar to the previous code. The features (measurements) and targets (species labels) are stored in the variables `X` and `y`, respectively.

The dataset is then split into training and testing sets using the `train_test_split` function. This function takes the features and targets as inputs and splits them into training and testing subsets. Here, 80% of the data is used for training (`X_train` and `y_train`), and 20% is used for testing (`X_test` and `y_test`). The `test_size=0.2` parameter specifies the proportion of the dataset to include in the test split, and `random_state=42` ensures that the split is reproducible by setting a seed for the random number generator.

A K-nearest neighbors classifier is initialized using `KNeighborsClassifier` with `n_neighbors=3`, meaning the classifier will consider the 3 nearest neighbors to make predictions. The classifier is then trained on the training data using the `fit` method, which involves storing the training data so that it can be used to make predictions.

Finally, the model's performance is evaluated on the testing data. The `score` method computes the accuracy of the classifier, which is the proportion of correctly predicted labels out of the total number of predictions. The accuracy score is printed to the console, providing an indication of how well the model performs on unseen data. This evaluation is crucial to understand the generalization ability of the model.

<u>Why Random Forest classifier is not necessary?</u>

Given that the KNN classifier achieved perfect accuracy on the Iris dataset, there might not be a strong necessity to use a random forest classifier for this specific task. Random forest classifiers are commonly employed for classification tasks that involve more complex decision boundaries or when dealing with larger datasets. However, it's worth considering the utilization of a random forest classifier if there's a desire to explore different models or validate the robustness of the results obtained from the KNN classifier. Engaging in such experimentation is a beneficial practice to ascertain the most suitable algorithm for the particular problem at hand.

Validation criteria (Accuracy, precision, recall, F1 score)

```python
from sklearn.metrics import f1_score, recall_score, precision_score, accuracy_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split

# Load the Iris dataset
iris = load_iris()
X = iris.data
y = iris.target

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the KNN classifier
knn = KNeighborsClassifier(n_neighbors=3)

# Train the classifier on the training data
knn.fit(X_train, y_train)

# Predict the labels for the test set
y_pred = knn.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)

# Calculate precision
precision = precision_score(y_test, y_pred, average='weighted')

# Calculate recall
recall = recall_score(y_test, y_pred, average='weighted')

# Calculate F1 score
f1 = f1_score(y_test, y_pred, average='weighted')

# Print the metrics
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

```
Accuracy: 1.0
Precision: 1.0
Recall: 1.0
F1 Score: 1.0
```

The script starts by importing the necessary libraries: `f1_score`, `recall_score`, `precision_score`, and `accuracy_score` from `sklearn.metrics`, along with `KNeighborsClassifier` from `sklearn.neighbors`, `load_iris` from `sklearn.datasets`, and `train_test_split` from `sklearn.model_selection`. The Iris dataset is then loaded using `load_iris`, and the feature data (`X`) and target labels (`y`) are extracted.

Next, the data is split into training and testing sets using the `train_test_split` function. Here, 80% of the data is used for training, and 20% is reserved for testing. The split is done randomly but reproducibly due to the `random_state=42` parameter. The training data (`X_train` and `y_train`) is used to train the K-nearest neighbors classifier, which is initialized with

`n_neighbors=3`. This means the classifier will consider the 3 nearest neighbors for making predictions.

Once the classifier is trained using the `fit` method, it is used to predict the labels for the test set (`X_test`) with the `predict` method. The predicted labels (`y_pred`) are then compared to the true labels (`y_test`) to evaluate the model's performance using several metrics: accuracy, precision, recall, and F1 score.

Accuracy is computed using `accuracy_score`, which measures the proportion of correctly predicted instances out of the total instances. Precision, recall, and F1 score are calculated using `precision_score`, `recall_score`, and `f1_score` respectively, with the `average='weighted'` parameter to account for class imbalances. Precision measures the ratio of true positive predictions to the total predicted positives, recall measures the ratio of true positives to the total actual positives, and the F1 score is the harmonic mean of precision and recall, providing a single measure of model performance.

The results show that all four metrics—accuracy, precision, recall, and F1 score—are 1.0. This indicates that the model perfectly predicted the species of all iris samples in the test set. Achieving perfect scores across all metrics suggests that the classifier was able to correctly identify the species for every test instance without any errors. However, while these results are ideal, they may also indicate a potential overfitting issue, particularly if the dataset is small or not very complex. It would be prudent to further validate the model with cross-validation or additional datasets to ensure its generalization capability.

- <u>Downloading the model and writing the python script</u>

```python
import joblib
from sklearn.neighbors import KNeighborsClassifier
from sklearn.datasets import load_iris

# Load the Iris dataset
iris = load_iris()
X = iris.data
y = iris.target

# Train the KNN classifier
knn = KNeighborsClassifier()
knn.fit(X, y)

# Save the trained model
joblib.dump(knn, 'iris_model.pkl')
```

```
['iris_model.pkl']
```

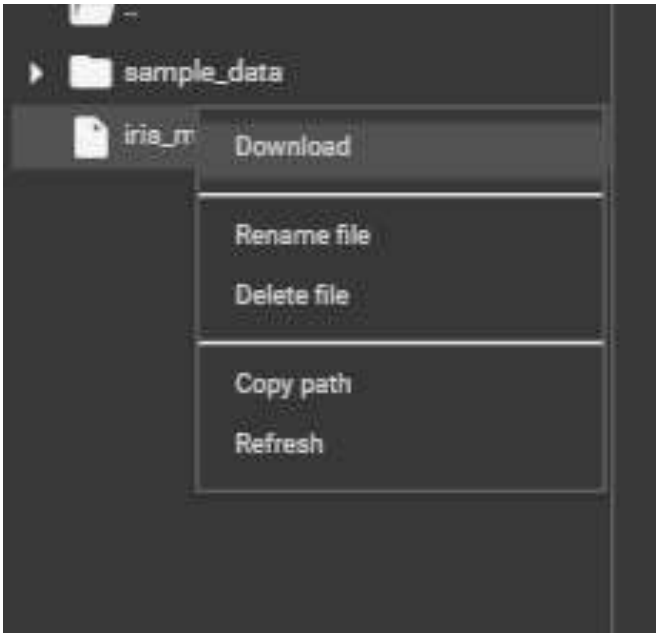

Figure 9 - Downloading the model

The model can be downloaded easily in the right pane op\f google colab which display the user's files. When the model files are downloaded to the PC it is now time to write the python script which asks for prediction from the model.

- Installing IronPython

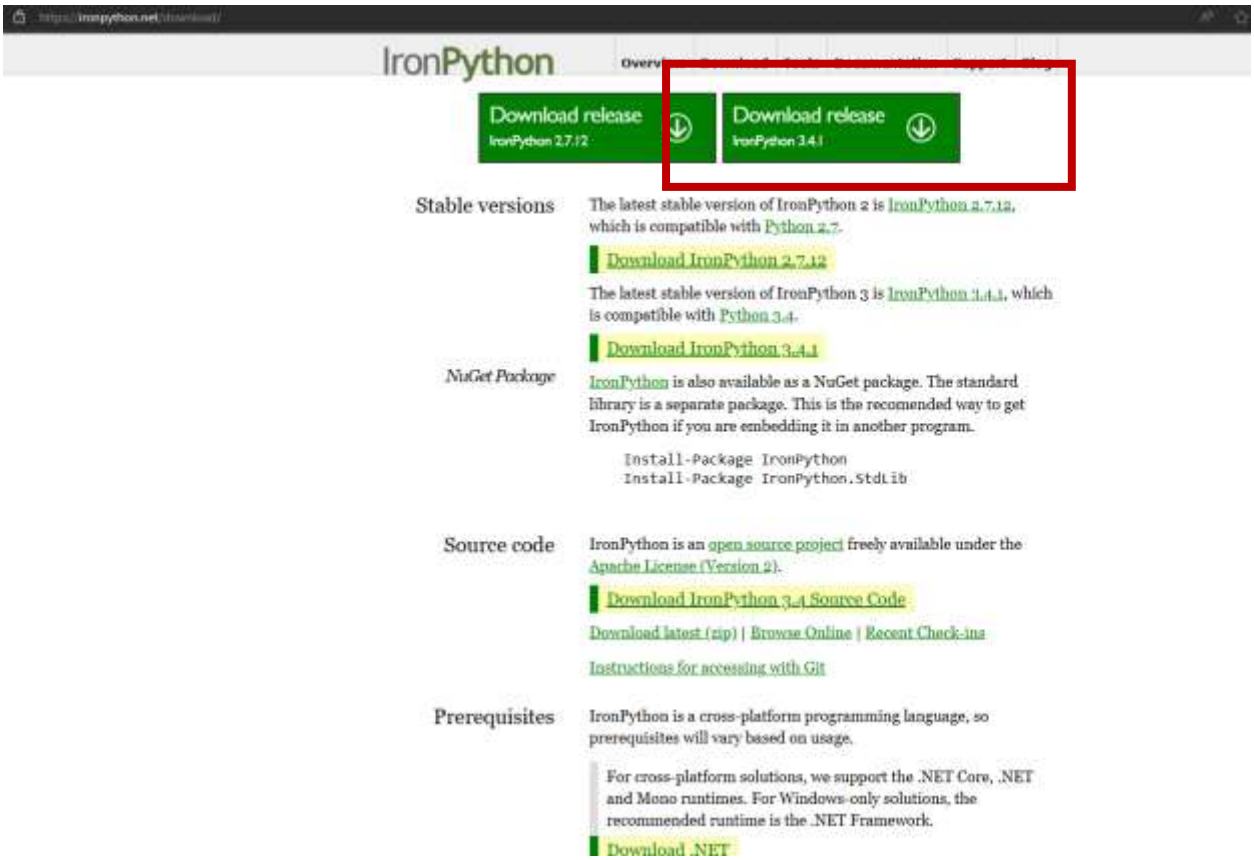

IronPython is known as a special version of Python that works smoothly with .NET Framework. It's open-source, meaning anyone can use it and even contribute to its development. IronPython lets the user mix Python with C# (another programming language used with .NET) in the same project. This mixing, or integration, makes it easy to use the strengths of both languages together.

When developers use IronPython, they can make Python code talk to C# code, and vice versa. This means they can take advantage of the features and capabilities of both languages within one application. For example, they can use existing libraries written in C# directly from Python, or they can add Python scripts to their C# applications to make them more flexible and adaptable.

IronPython supports all the dynamic features of Python, like dynamic typing and introspection, even when used in a .NET environment. These features make it easy to write code that can

change and adapt on the fly, which is handy for building applications that need to be flexible and responsive to different situations.

To use IronPython in a C# application, developers need to do a few things. They have to include references to IronPython in their project, set up a way for Python code to run within the application, and figure out how to pass data between Python and C#. Once everything is set up, they can start writing code that mixes Python and C# together.

IronPython is especially useful for tasks like testing out new ideas quickly or adding custom logic to a C# application without having to recompile it. Its ability to work well with .NET makes it a valuable tool for developers who want to combine the power of Python with the capabilities of the .NET Framework.

IronPython serves as a bridge between Python and .NET, allowing developers to create applications that benefit from the strengths of both worlds. It's a versatile tool that can help make applications more flexible, dynamic, and powerful.

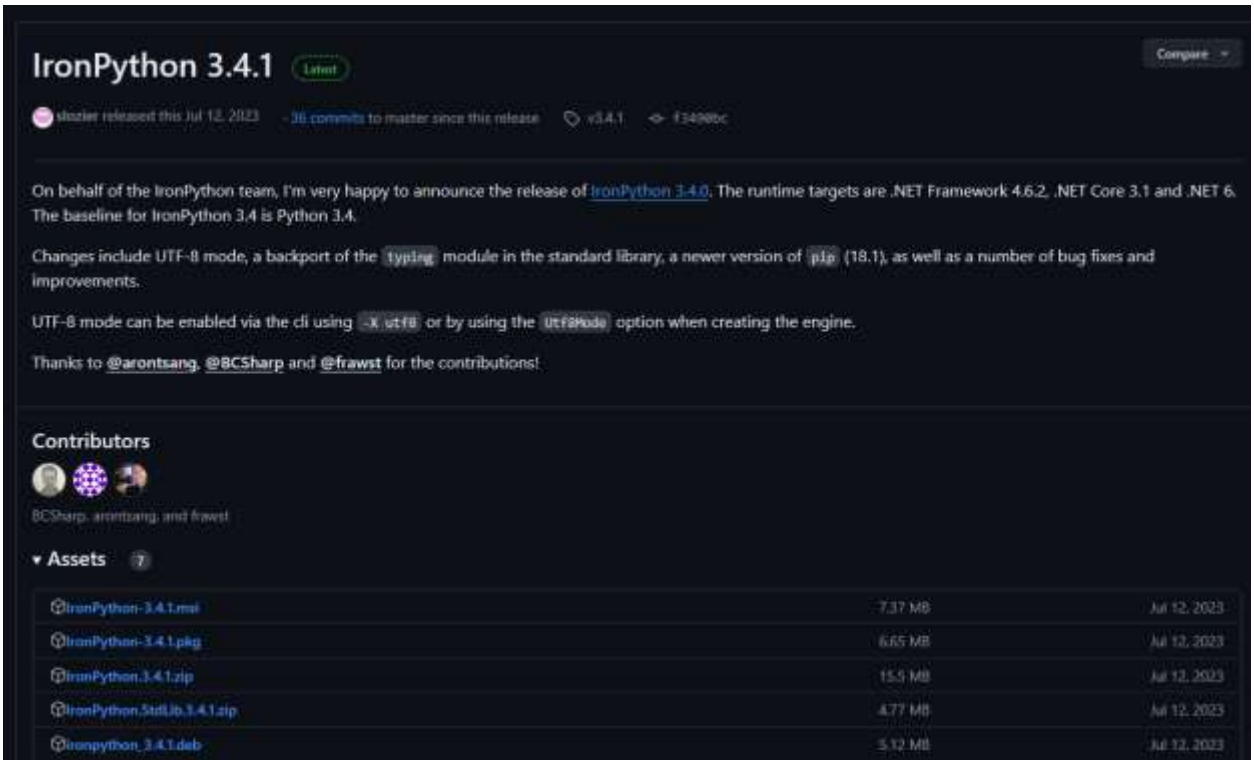

Figure 10 – requirements for IronPython

The above displayed image shows the download site for IronPython. It says that IronPython best works and targeted to .NET 4.6.2 framework. Therefore, in order to run the apple quality application using C#. It needs to be created under the .NET framework 4.6.2 or else it has high potential risks of getting crashed or malfunction the application. The current released .NET framework is .NET 8.0 and it may cause several crashes due to the framework complexity and changes.

Given below are steps to install .NET framework 4.6.2 for Visual Studio 2022 (Current version used to develop the iris flower species predicting application)

1. Launch Visual studio installer and click Modify

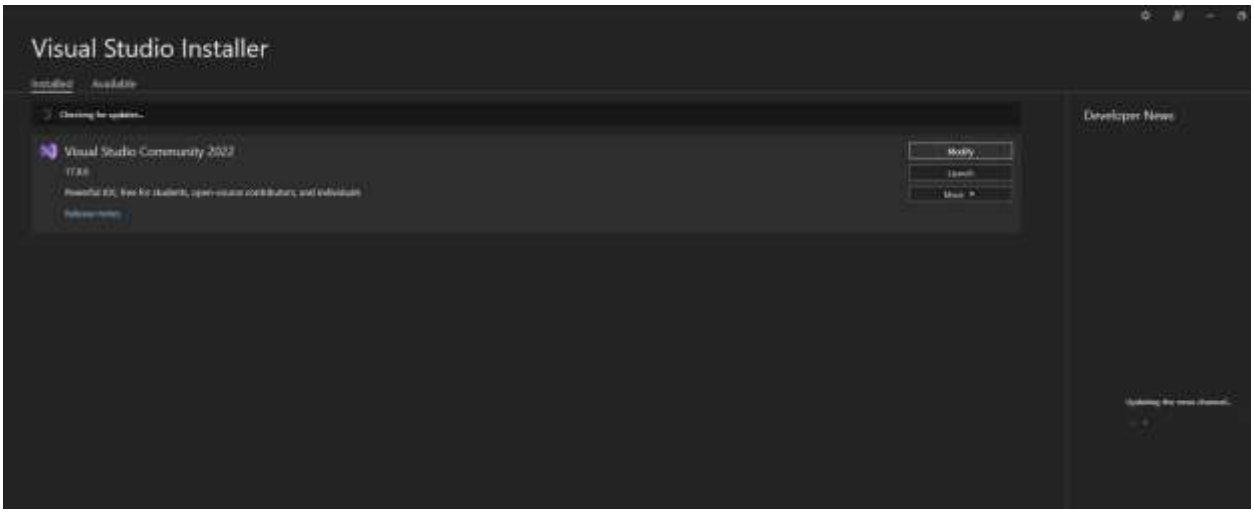

2. Under .NET desktop development, Check the box mentioning .NET framework 4.6.2 - 4.7.1 and click modify to install.

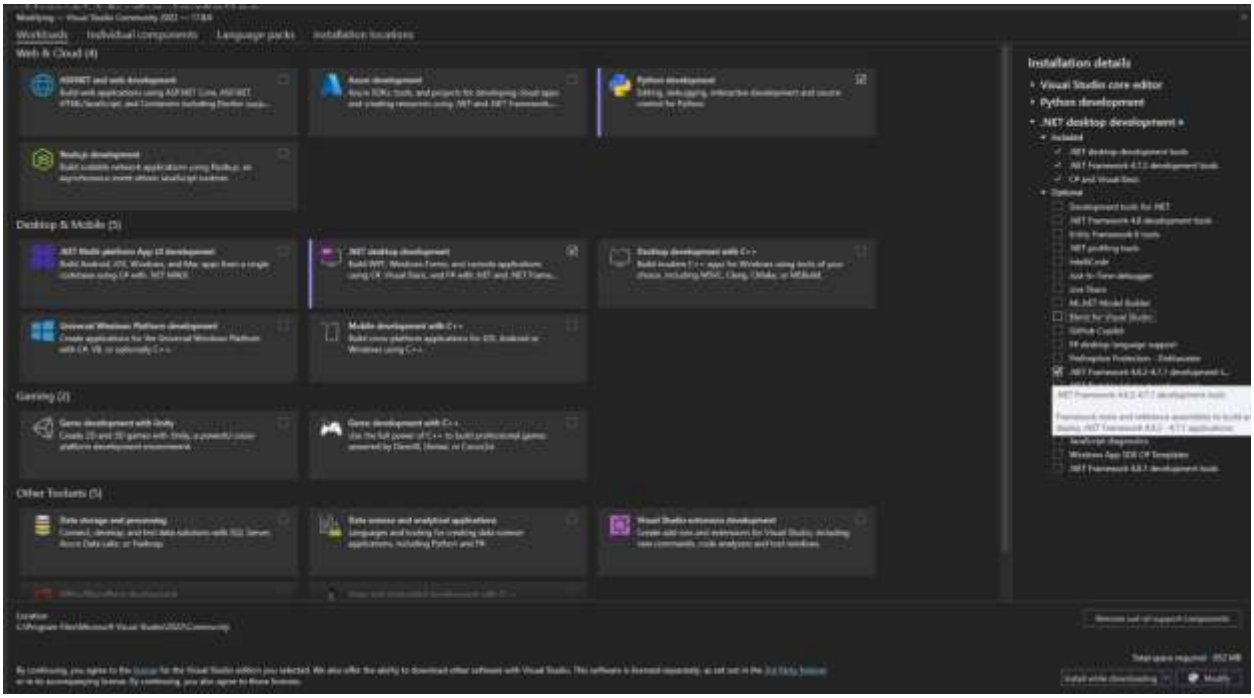

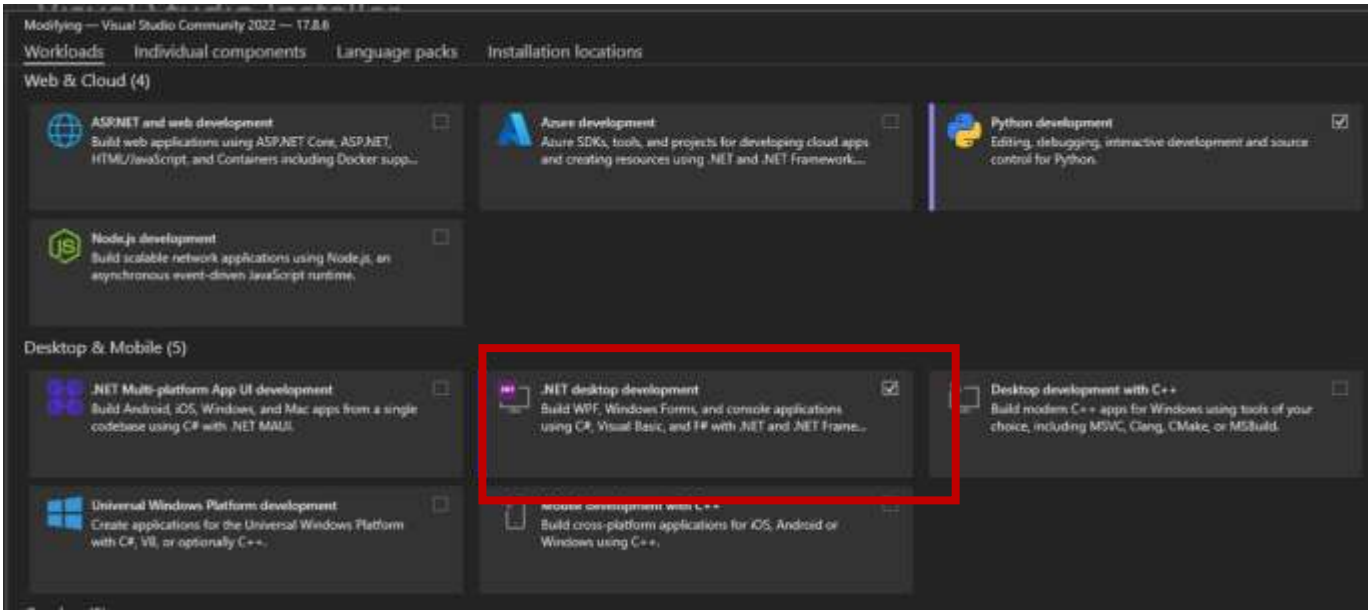
Modifying — Visual Studio Community 2022 — 17.8.6
Workloads
Individual components
Language packs
Installation locations
Web & Cloud (4)
ASP.NET and web development
Build web applications using ASP.NET Core, ASP.NET, HTML/JavaScript, and Containers including Docker supp...
Azure development
Azure SDKs, tools, and projects for developing cloud apps and creating resources using .NET and .NET Framework...
Python development
Editing, debugging, interactive development and source control for Python.
Node.js development
Build scalable network applications using Node.js, an asynchronous event-driven JavaScript runtime.
Desktop & Mobile (5)
.NET Multi-platform App UI development
Build Android, iOS, Windows, and Mac apps from a single codebase using C# with .NET MAUI.
.NET desktop development
Build WPF, Windows Forms, and console applications using C#, Visual Basic, and F# with .NET and .NET Frame...
Desktop development with C++
Build modern C++ apps for Windows using tools of your choice, including MSVC, Clang, CMake, or MSBuild.
Universal Windows Platform development
Create applications for the Universal Windows Platform with C#, VB, or optionally C++.
Build cross-platform applications for iOS, Android or Windows using C++.

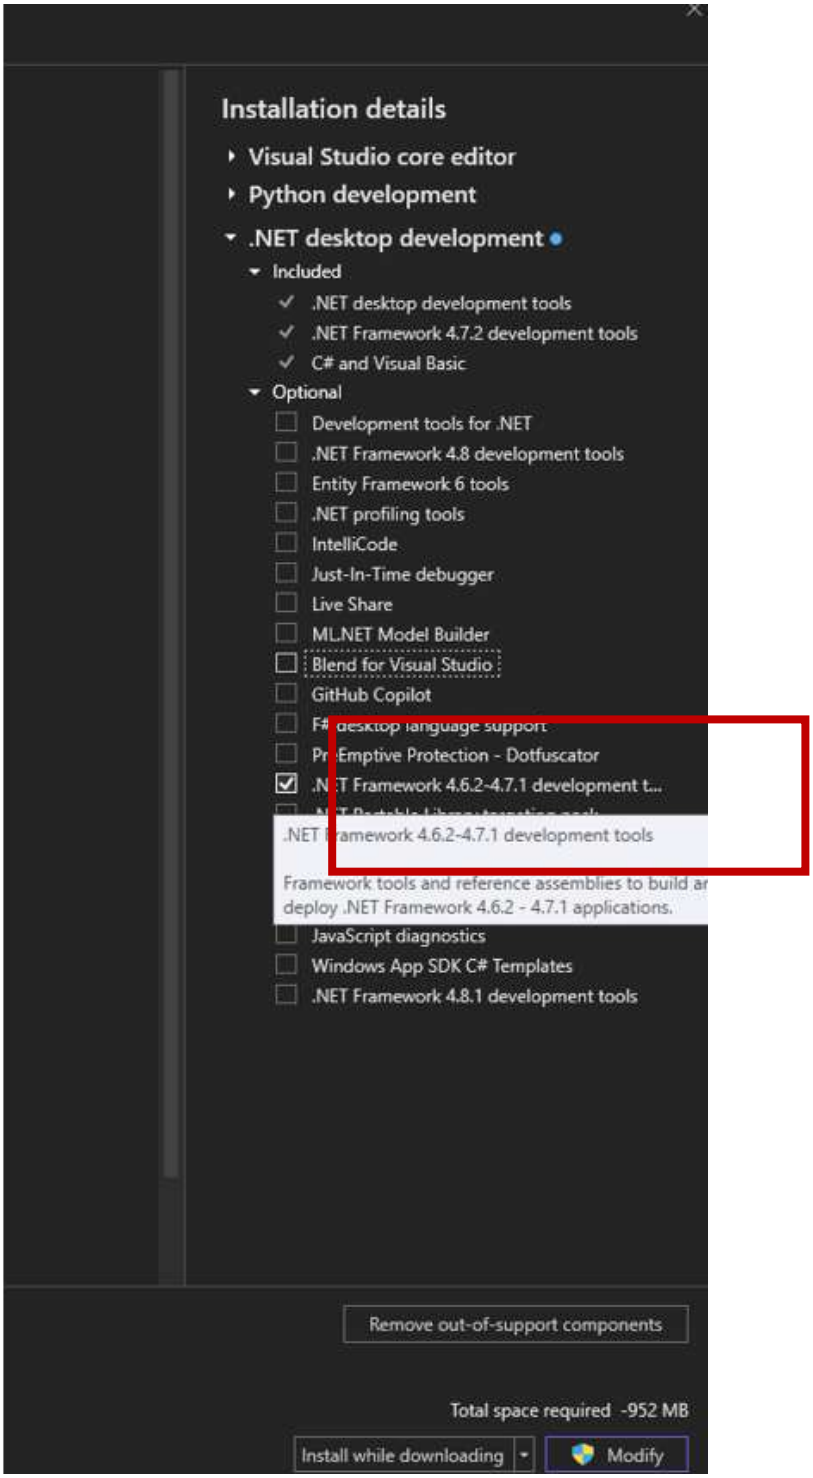

Figure 11 - Installing .NET 4.6.2 (As I have already installed it shows space required in negative values)

- <u>Installing IronPython to the C# Application</u>

In order to utilize IronPython in a C# windows forms application. The following steps can be followed.

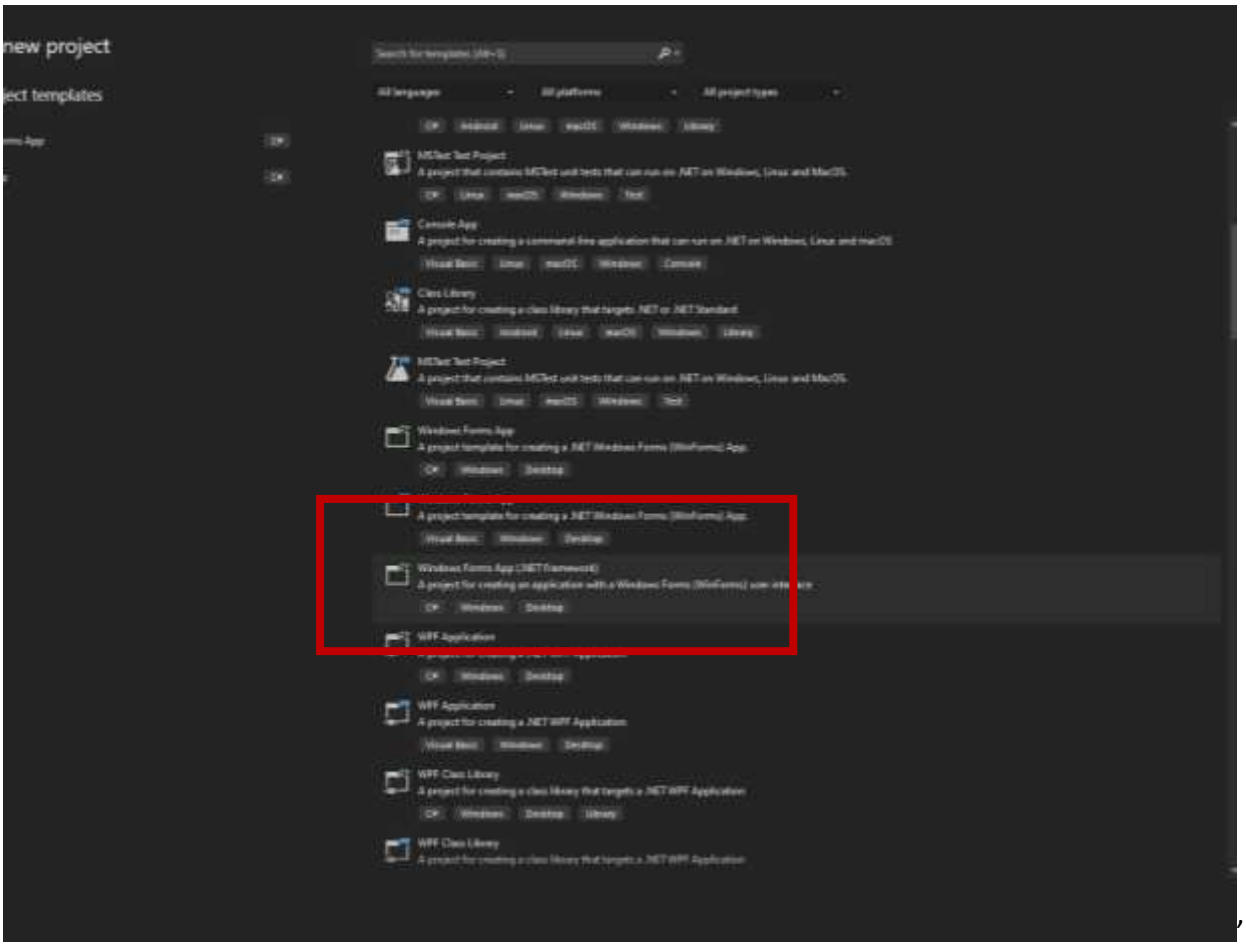

Users should make sure to select windows forms app (.NET framework) for the IronPython related application. A usual windows form application displays errors when Iron Python is installed and executed.

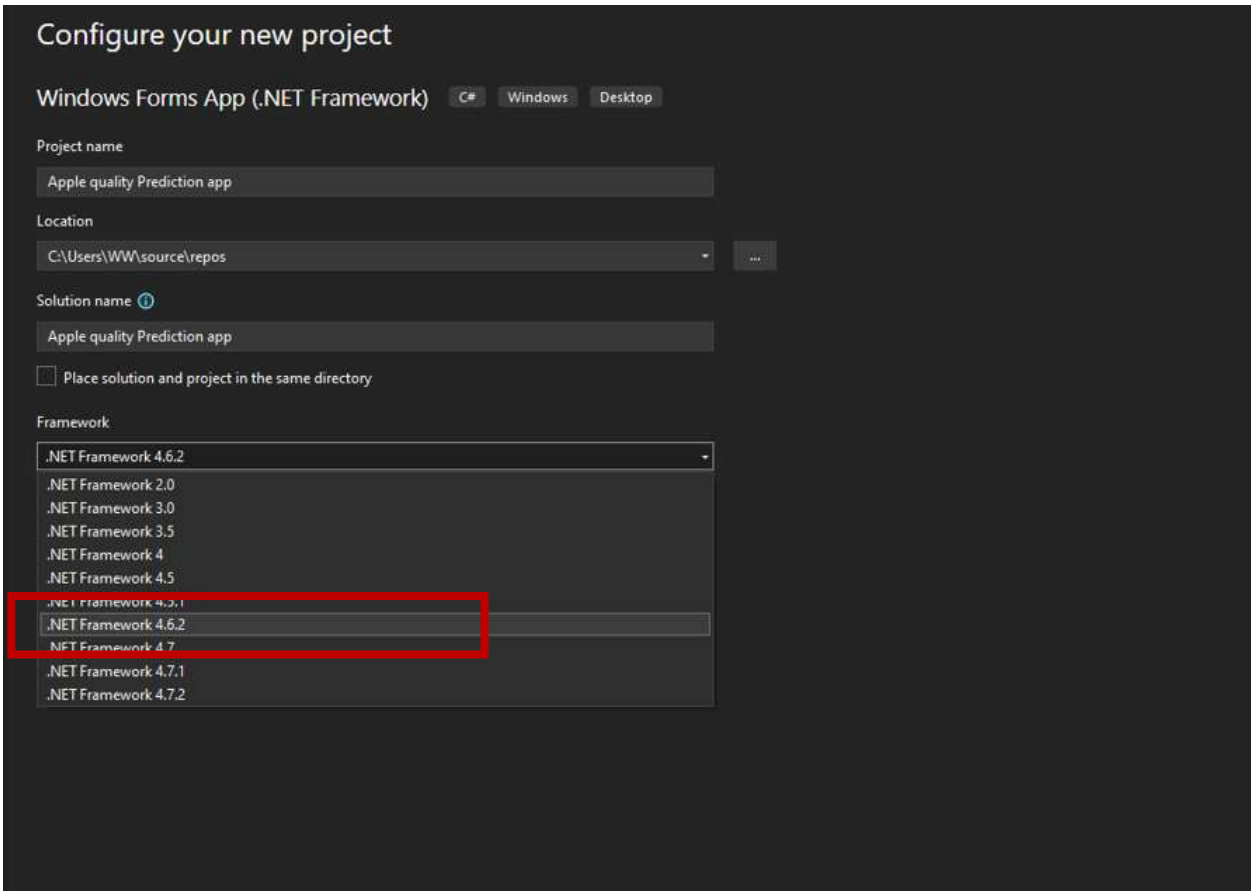

The target framework from the IronPython version will be used for the development process of the iris flower classifier application.

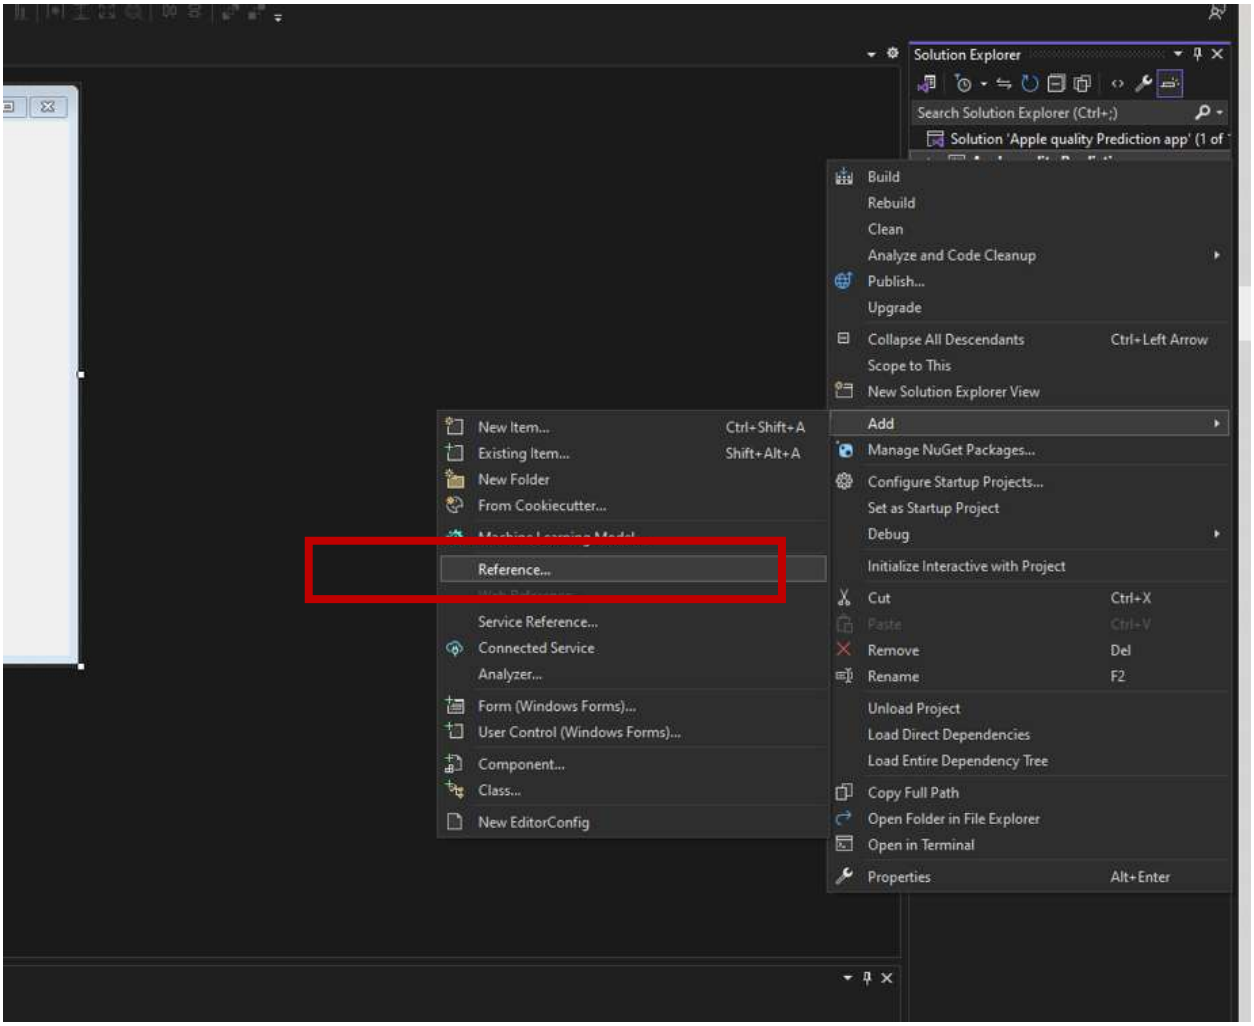

IronPython will be installed as a reference for the application therefore, it the reference manager should be opened.

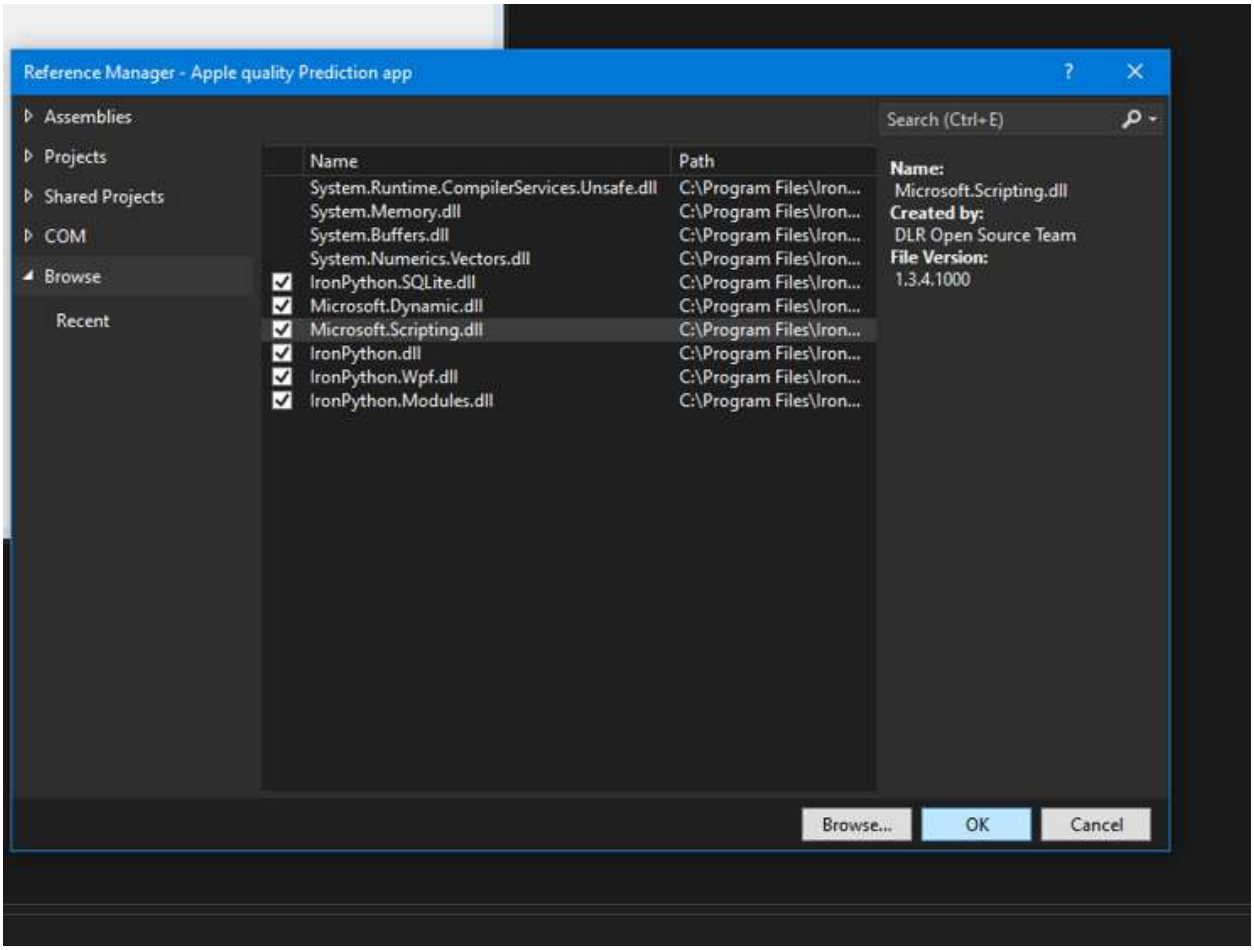

Figure 12 - installing Iron Python to C#

When the reference manager is opened users can neither find or browse the required IronPython files and select them as reference to the application. Now IronPython is successfully installed to the iris flower application.

- About the application

For the Iris flower classification application, a simple user interface will be created to input the data attributes: sepal width, sepal length, petal width, and petal length. These attributes should be provided in decimal or numeric formats and must fall within a value range scaled between -8 to 8. This range constraint ensures that the input values are appropriate for the model.

Upon entering the attribute values within the specified range, the application will classify the Iris species based on the provided input. The classification will occur using a pre-trained machine learning model, such as a K-Nearest Neighbors (KNN) classifier, which has been trained on the Iris dataset.

The output will be displayed in a separate interface, presenting a summary of the classification result. This summary will include the predicted Iris species corresponding to the input attribute values. Additionally, the user will have the option to print the summary for future reference or documentation purposes.

Overall, the Iris flower classification application will provide a user-friendly interface for inputting attribute values and obtaining accurate predictions of the Iris species based on the provided data.

- Application Development using IronPython and C#

The application will be developed according to the wireframe and UML design done in activity 2.

Form 1 User Interface (Starting Form of the application)

Figure 13 - Form 1 UI

Form 1 Codes

```
32
33          }
34
            1 reference
35          private void button1_Click(object sender, EventArgs e)
36          {
37              Form2 form2 = new Form2();
38              form2.Show();
39              this.Hide();
40          }
41      }
42  }
43
```

Figure 14 - Form 1 button Code

The above form will be the launching display page for the application. It will provide a button “Start analysis” where the user will proceed to the next page to provide inputs.

Form 2 User Interface (Input Form of the application)

This is the code where the user will input his data in order to analyze and predict the iris flower species.

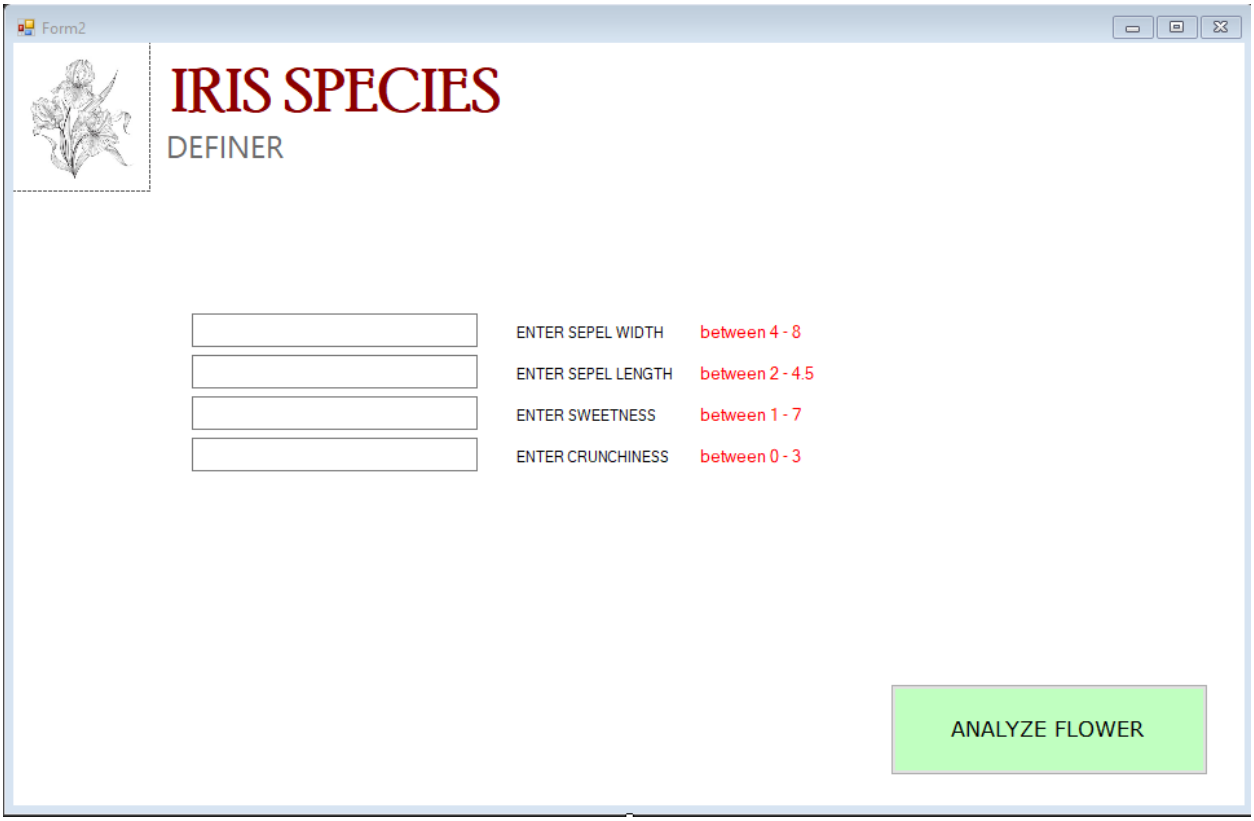
Form2
IRIS SPECIES
DEFINER
ENTER SEPEL WIDTH
between 4 - 8
ENTER SEPEL LENGTH
between 2 - 4.5
ENTER SWEETNESS
between 1 - 7
ENTER CRUNCHINESS
between 0 - 3
ANALYZE FLOWER

Form 2 Codes

This code is for a Windows app called "iris_flower_app" that uses IronPython, a tool to blend Python with C#. When you run the app, it shows a window (Form2) where you can interact with it. The code has two secret tools: `pythonEngine` and `pythonScope`, which help to handle Python code. When the window starts (in the `Form2()` method), it sets up these tools and runs Python code from a file called "predict_iris_species.py". This Python script likely figures out what kind of Iris flower you have based on its features. The comments in the code explain what each part does, making it easier to understand. By blending C# with Python, this app becomes a powerful tool for predicting Iris flower species right from your Windows desktop.

```
1 reference
private void btnAnalyze_Click(object sender, EventArgs e)
{
    // Check if any textbox is empty
    if (string.IsNullOrEmpty(txtSepalLength.Text) || string.IsNullOrEmpty(txtSepalWidth.Text) ||
        string.IsNullOrEmpty(txtPetalLength.Text) || string.IsNullOrEmpty(txtPetalWidth.Text))
    {
        // Display a warning message
        var result = MessageBox.Show("You entered missing features, are you sure you want to predict the species with existing features?", "Warning", MessageBoxButtons.YesNo, MessageBoxIcon.Warning);
        if (result == DialogResult.No)
            return; // Don't proceed with the analysis
    }

    // Convert input values to double
    if (!double.TryParse(txtSepalLength.Text, out double sepalLength) ||
        !double.TryParse(txtSepalWidth.Text, out double sepalWidth) ||
        !double.TryParse(txtPetalLength.Text, out double petalLength) ||
        !double.TryParse(txtPetalWidth.Text, out double petalWidth))
    {
        // Display error message for invalid input format
        MessageBox.Show("Invalid input format. Please enter valid numeric values", "Error", MessageBoxButtons.OK, MessageBoxIcon.Error);
        return; // Exit method
    }

    // Check if any value is within a reasonable range (this can be adjusted based on domain knowledge)
    if (sepalLength <= 0 || sepalWidth <= 0 || petalLength <= 0 || petalWidth <= 0)
    {
        // Display error message for out-of-range values
        MessageBox.Show("Invalid value. Please enter positive values.", "Error", MessageBoxButtons.OK, MessageBoxIcon.Error);
        return; // Exit method
    }

    // All validations passed, proceed with analysis
    var attributes = new double[] { sepalLength, sepalWidth, petalLength, petalWidth };
    var predictSpecies = pythonScope.GetVariable<Func<double[], string>>("predict_species");
    string prediction = predictSpecies(attributes);
    var resultForm = new Form3(prediction);
    resultForm.Show();
}
```

This C# code defines a method named `btnAnalyze_Click`, which is executed when a button named "Analyze" is clicked in a Windows Forms application. The purpose of this method is to process user input, validate it, and then use a Python function to predict the species of an Iris flower based on the provided characteristics.

The first part of the method checks if any of the input textboxes for sepal length, sepal width, petal length, and petal width are empty. If any textbox is empty, a warning message is displayed asking the user if they want to proceed with the analysis despite missing features. If the user chooses not to proceed, the method exits early to avoid performing the analysis with incomplete data.

Following the input validation, the method attempts to convert the input values from the textboxes into double-precision floating-point numbers. This conversion is essential for numerical calculations and data processing. If any conversion fails due to invalid input (e.g., non-numeric characters), an error message is displayed, and the method exits, preventing further analysis with incorrect data.

Once the input values are successfully converted and validated, the method checks if any of the input values are less than or equal to zero. This check ensures that all input values, representing physical measurements, are positive and within a reasonable range. If any input value is found to be zero or negative, an error message is displayed, and the method exits to prevent invalid input from proceeding to the analysis stage.

If all input values pass the validation checks, they are packaged into an array named `attributes`, which is then passed as an argument to a Python function named `predict_species`. This Python function likely contains a machine learning model trained on the Iris dataset, capable of predicting the species of an Iris flower based on its characteristics. The result of the prediction is stored in a variable named `prediction`.

Finally, the predicted species is used to create a new form (`Form3`) to display the prediction result to the user. This form is then shown to the user, providing them with the predicted species of the Iris flower based on the input characteristics.

In summary, this code demonstrates the integration of Python functionality into a C# application, enabling the prediction of Iris flower species using machine learning models within a Windows Forms environment.

Form 3 User Interface (Output Form of the application)

This is the form where outputs and the summery is displayed. The user will also be able to download the results in a TXT. Format.

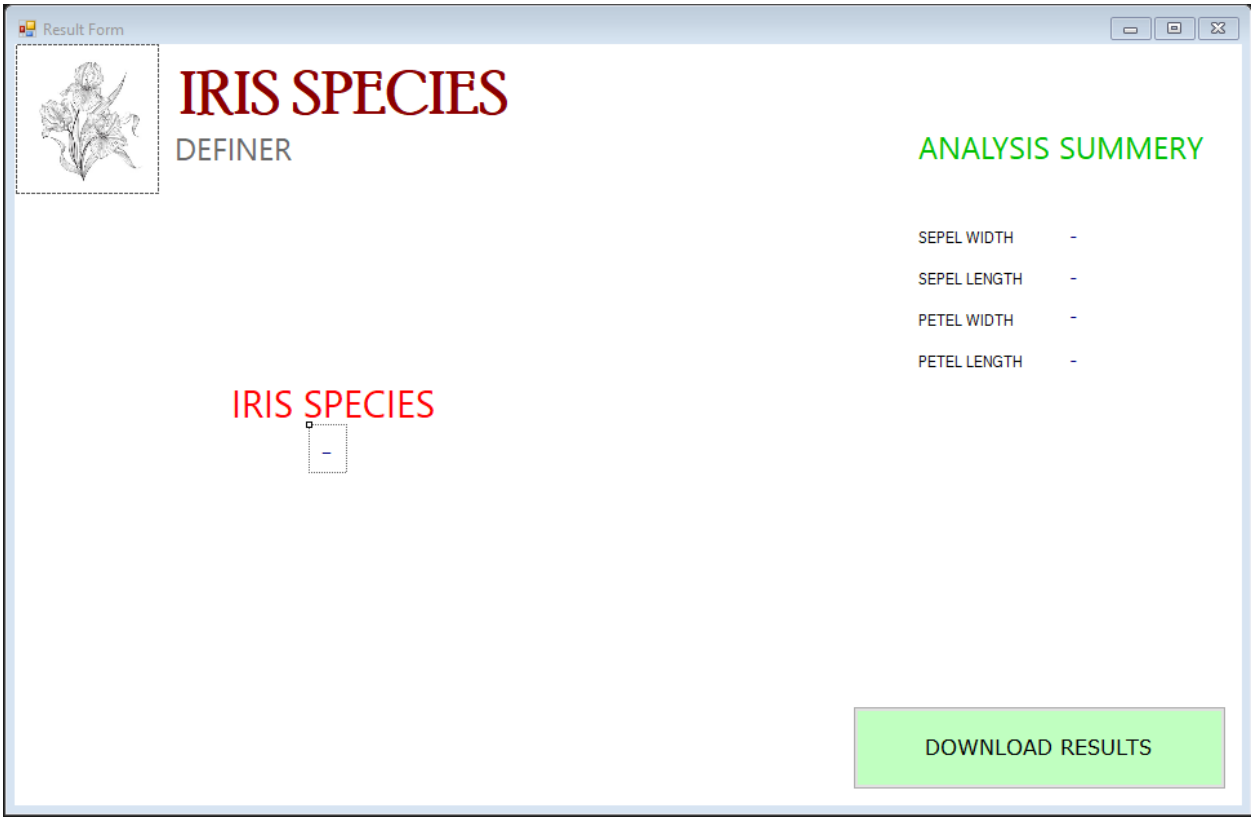
Result Form
IRIS SPECIES
DEFINER
ANALYSIS SUMMERY
SEPEL WIDTH -
SEPEL LENGTH -
PETEL WIDTH -
PETEL LENGTH -
IRIS SPECIES
-
DOWNLOAD RESULTS

Form 3 Codes

```
0 references
private void ResultForm_Load(object sender, EventArgs e)
{
    InitializeComponent();
    lblPrediction.Text = $"Predicted Species: {prediction}";

    this.sepalLength = sepalLength;
    this.sepalWidth = sepalWidth;
    this.petalLength = petalLength;
    this.petalWidth = petalWidth;
    this.prediction = prediction;

    DisplayResults();

}
```

The above C# code defines a class called `ResultForm`, likely representing a form in a Windows Forms application designed to display the results of apple quality prediction. Here's a breakdown of the code:

In the code, there's a namespace named `iris_flower ` enclosing the class definitions, ensuring they are organized and distinguished within the application.

Within the `ResultForm` class, private fields are declared to hold attributes related to iris flower. These fields serve as containers for the values passed to the form for displaying prediction results.

The class contains a constructor method responsible for initializing the form's components, including event handlers and visual settings.

Additionally, an event handler named `ResultForm_Load` is defined, triggered when the form is loaded. This handler calls the `DisplayResults()` method to populate the form with prediction results.

The class also defines a constructor with parameters, likely used to initialize the form with prediction results. It accepts parameters representing flower attributes and the prediction result, assigning them to corresponding private fields.

To visually present the prediction result, the form includes a label named `lblPrediction`. The prediction text is dynamically included in this label's text using string interpolation, ensuring that the displayed prediction corresponds to the provided attributes.

```
    private void DisplayResults()
    {
        lblSepalLength.Text = $"Sepal Length: {sepalLength}";
        lblSepalWidth.Text = $"Sepal Width: {sepalWidth}";
        lblPetalLength.Text = $"Petal Length: {petalLength}";
        lblPetalWidth.Text = $"Petal Width: {petalWidth}";
        lblPrediction.Text = $"Predicted Species: {prediction}";
    }


private void button1_Click_1(object sender, EventArgs e)
{
    SaveFileDialog saveFileDialog = new SaveFileDialog();
    saveFileDialog.Filter = "Text File (*.txt)|*.txt";
    saveFileDialog.Title = "Save Iris Species Prediction";

    if (saveFileDialog.ShowDialog() == DialogResult.OK)
    {
        using (StreamWriter writer = new StreamWriter(saveFileDialog.FileName))
        {
            writer.WriteLine($"Sepal Length: {sepalLength}");
            writer.WriteLine($"Sepal Width: {sepalWidth}");
            writer.WriteLine($"Petal Length: {petalLength}");
            writer.WriteLine($"Petal Width: {petalWidth}");
            writer.WriteLine($"Predicted Species: {prediction}");
        }
        MessageBox.Show("Results saved successfully.", "Saved", MessageBoxButtons.OK, MessageBoxIcon.Information)
        }
}
```

This code snippet appears to be part of a graphical user interface (GUI) application, likely developed using a framework such as Windows Forms in C#. Let's break it down:

DisplayResults Method

This method is responsible for updating the labels in the user interface with the results of some calculations or predictions. It formats the text of each label to display specific attributes associated with the Iris flower dataset, such as Sepal Length, Sepal Width, Petal Length, Petal Width, and a Predicted Species.

buttonSave_Click Event Handle

This event handler is triggered when the user clicks a button, presumably labeled "Save". It initializes a SaveFileDialog object, which prompts the user to select a location to save a text file. The filter restricts the file type to text files only.

If the user selects a file and clicks "OK", the code proceeds to open a StreamWriter to write to the selected file. It then writes the same information displayed in the labels (Sepal Length, Sepal Width, Petal Length, Petal Width, and Predicted Species) into the text file, each on a separate line. After writing the data, it displays a message box indicating successful saving.

In summary, this code segment serves two main purposes: displaying the results of some calculations or predictions on the user interface, and providing functionality for users to save these results to a text file for later reference. It demonstrates the integration of UI manipulation and file I/O operations in a C# application.

ACTIVITY 4

Testing The Application and report

Once again the test cases to be carried out for this application are as follows.

- **Test Case 1: Functionality (Functional Testing and Integration Testing)**

Description

Confirm the proper functioning of all user interface components related to iris flower species prediction.

Preconditions

The iris flower software is installed and operational.

Steps

- Launch the application.
- Navigate through various UI elements, such as buttons and text boxes.
- Input sample data regarding apple characteristics and execute the prediction.

Tests

3. Functional and UI Testing
4. Unit Testing

Expected Outcome

The user interface operates smoothly, allowing for the input of iris flower and generating predictions without errors.

- **Test Case 2: Data Input and Processing (Performance Testing along With UI tests)**

Description

Verify the accuracy of data input and processing specifically for iris flower species prediction.

Preconditions: The iris flower prediction software is installed and operational.

Steps

- Input sample data regarding apple quality characteristics.
- Validate the accuracy of data processing for iris flower prediction.

Tests

6. Integral test to clarify that the application is running utilizing IronPython without any errors.
7. Testing the Input values and Errors regarding wrong inputs
8. Testing the outputs of the application

9. Testing the application by providing real world scenario inputs and getting predicted outcomes
10. Testing the accuracy of the application related to logical explanations and examples

Expected Outcome

The data undergoes error-free processing, ensuring accurate preparation for iris flower prediction analysis.

TEST CASE 1

- Functional and UI Testing

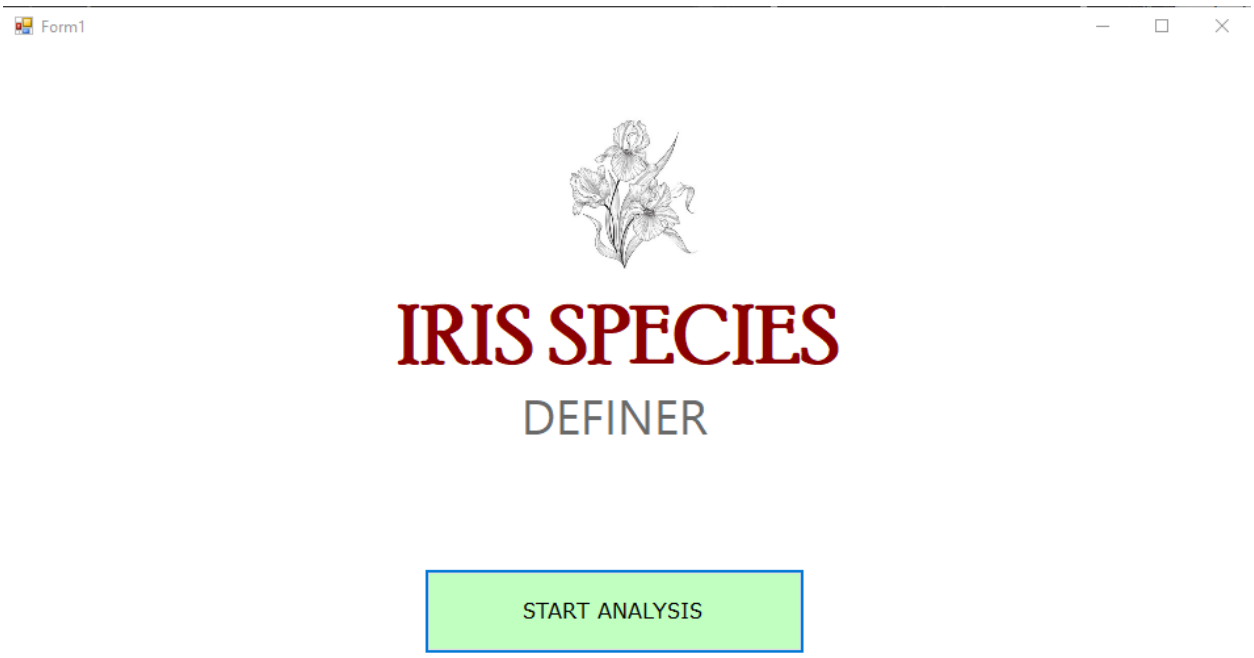

Figure 15 - Launched application UI

Shown above is the start page of the application which will be shown at the moment the application is launched.

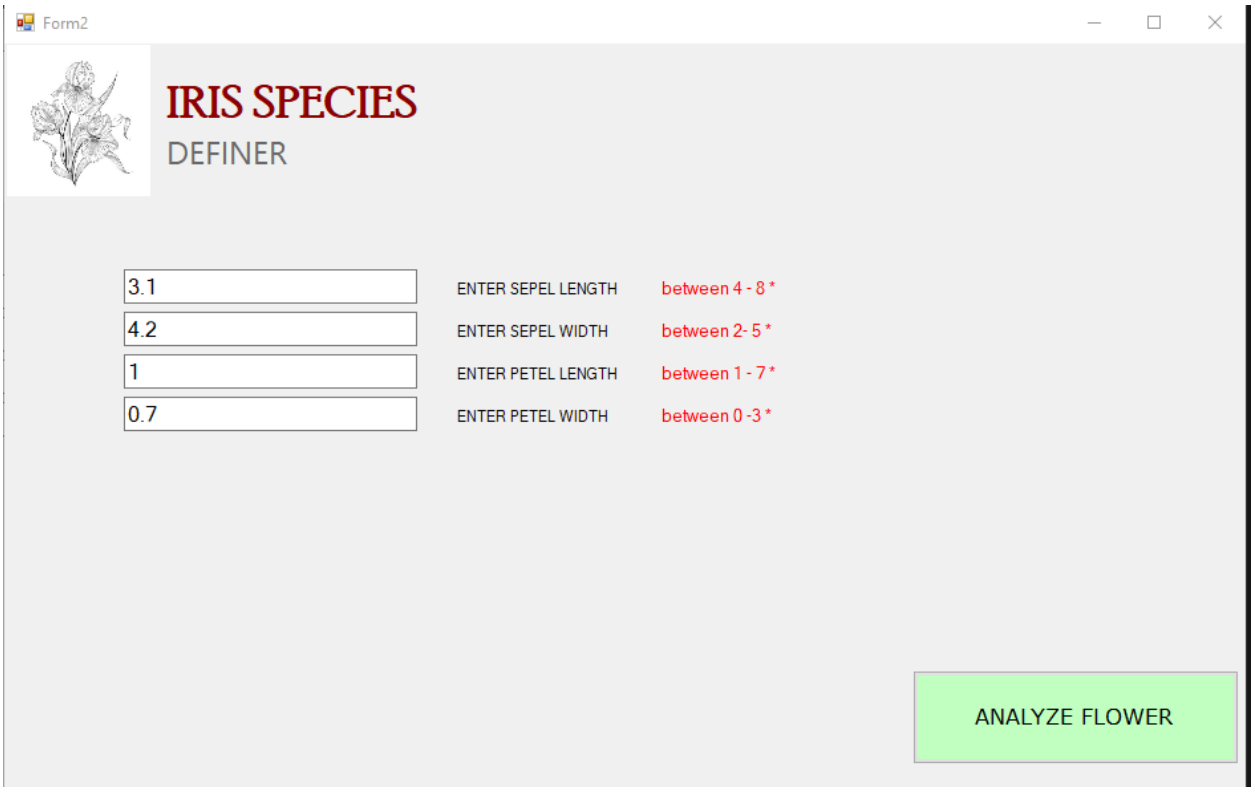

Figure 16 - Input Interface of the application

The above image shows that the application runs effectively so far as all input can be done accordingly to the requirements. (The requirement of the user should be able to input float values for analysis is highlighted here.) The user interface is clear and user friendly.

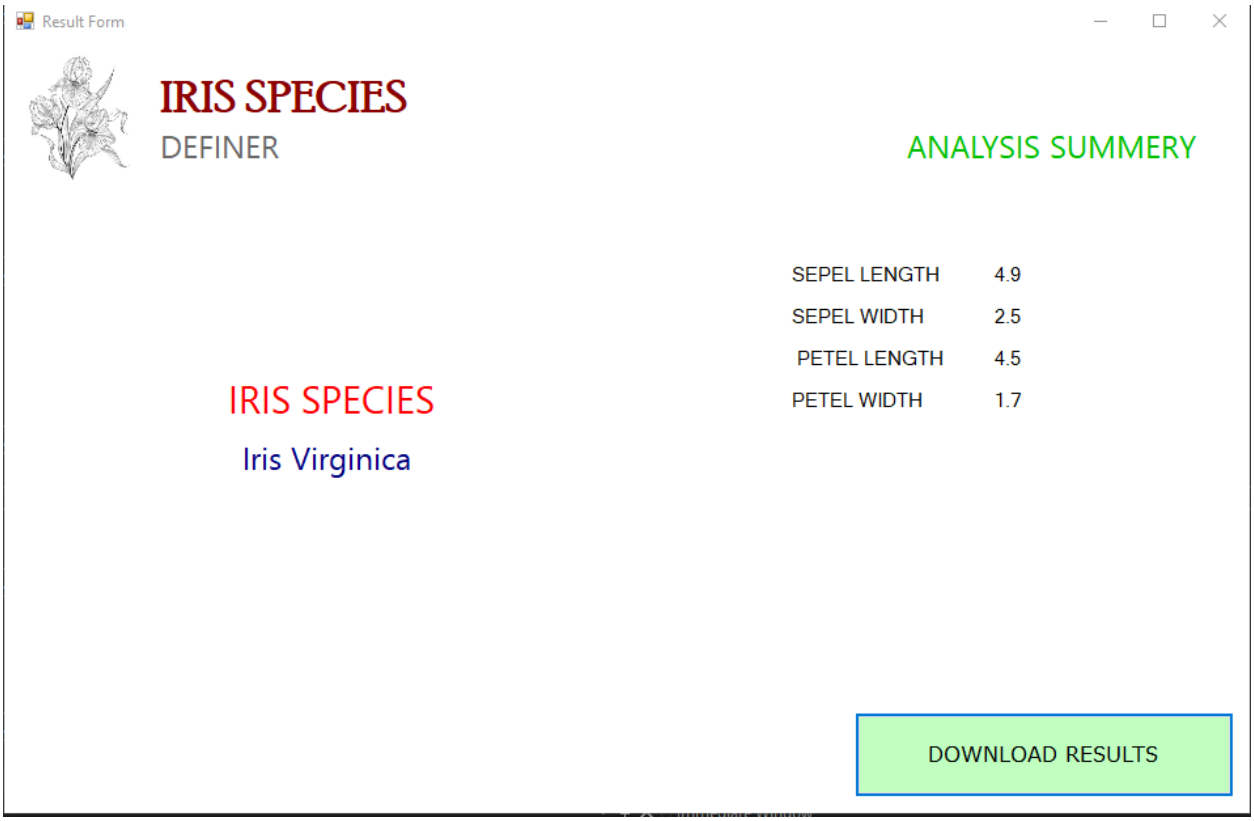

Figure 17 - Output part of the application

The above screenshot shows the output values after analysis of the values entered in the previous screenshot. This shows that the interface of the application is working well and the analysis is done successfully without an error, so that it proceeded to this part of the application.

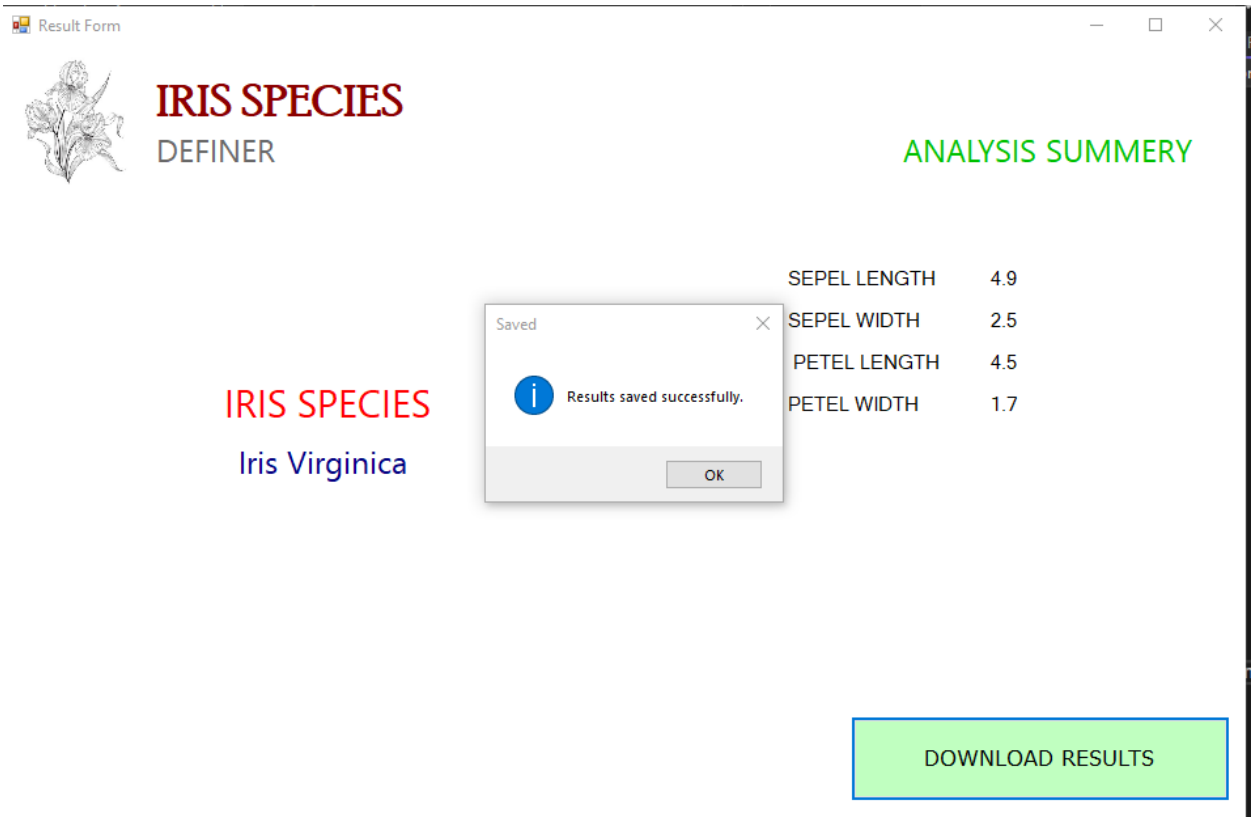

Figure 18 - UI Functionality

This is the message which pops up after the .txt file containing the results was downloaded to the PC successfully. This showcases that the UI is fully functional covering the needs expected from the application. It will display the necessary outputs with a summary of the analysis and also provides the facility to print out the result.

Therefore, the summery of this test,

TEST CASE 1

Tested By - Lazan Warshath

Description – UI functionality and User requirements check with application smoothness.

Expected outcome - The user interface operates smoothly, allowing for the input of iris flower characteristics and generating predictions without errors.

Result – **success without any issues**

- Unit Testing

iris_flower.sln - Unit Test Results

Generated on: 2024/06/12

Total tests: 20
Passed: 20
Failed: 0
Skipped: 0
Test execution time: 2.45 seconds

- Form_1_Label_1 - Passed
- Form_1_Label_2 - Passed
- Form_2_Label_3 - Passed
- Form_1_Start_button - Passed
- Form_2_Input_button - Passed
- Form_2_Textbox_1 - Passed
- Form_2_Textbox_2 - Passed
- Form_2_Textbox_3 - Passed
- Form_2_Textbox_4 - Passed
- Form_2_call_iron_python - Passed
- Form_2_Iron_python_prediction - Passed
- Form_2_Label_1 - Passed
- Form_2_Label_2 - Passed
- Form_2_Label_3 - Passed
- Form_2_Label_4 - Passed
- Form_2_Label_5 - Passed
- Form_2_Label_6 - Passed
- Form_3_Label_1 - Passed
- Form_3_Label_2 - Passed
- Form_3_Label_3 - Passed
- Form_3_Label_4 - Passed
- Form_3_Label_5 - Passed
- Form_3_Label_6 - Passed
- Form_3_print_button - Passed
- Errors_check - Passed
- Invalid_syntax_check - Passed
- iris_flower.sln - Passed
- Overall_Application_check - Passed

All tests for iris_flower.sln are working as expected.

Replit unit tests using Python. All rights reserved © 2024

Unit testing is a software testing technique where individual components or units of a program are tested independently to ensure they function correctly. Each unit is the smallest testable part of an application, such as a function or method, and is typically tested in isolation from the rest of the application to verify that it performs as expected. The goal of unit testing is to identify and fix bugs early in the development process, improving code quality and reducing the likelihood of issues in later stages.

For the Iris Flower Application, all unit tests have been executed successfully, indicating that each critical function and interaction within the application is working as intended. The tests covered various aspects such as button clicks, input handling, data analysis, calls to the IronPython engine for predictions, and the accuracy of the model's output. Additionally, checks for error handling, syntax validation, and overall application functionality were conducted. The successful passing of all these tests demonstrates that the application is effective, reliable, and ready for deployment, with each component performing its role effectively.

Therefore, the summery of this test,

TEST CASE 1 (2)

Tested By - Lazan Warshath

Description – Unit testing.

Expected outcome – The components of the application running well in the code without errors..

Result – **success without any issues**

TEST CASE 2

- Integral Testing

This test is done in order to clarify that the application is running utilizing IronPython without any errors.

```
using System;
using System.Collections.Generic;
using System.ComponentModel;
using System.Data;
using System.Drawing;
using System.Linq;
using System.Text;
using System.Threading.Tasks;
using System.Windows.Forms;
using IronPython.Hosting;
using Microsoft.Scripting.Hosting;

namespace iris_flower_app
{

    3 references
    public partial class Form2 : Form
    {
        private ScriptEngine pythonEngine;
        private ScriptScope pythonScope;
        0 references
        public Form2()
        {
            InitializeComponent();
            InitializePython();
        }
        1 reference
        private void InitializePython()
        {
            pythonEngine = Python.CreateEngine();
            pythonScope = pythonEngine.CreateScope();
            pythonEngine.ExecuteFile(@"C:\Users\WW\Downloads\predict_iris_species.py", pythonScope);
        }
```

Figure 19 - The application is running well without errors using IronPython

The above screenshot shows the statement “has exited with code 0” which means there are no errors within the application even when IronPython is a part of its reference and code. The code runs smoothly interacting with the application as expected.

Therefore, the summery of this test,

TEST CASE 2 (1)

Tested By - Lazan Warshath

Description – Integral test to clarify that the application is running utilizing IronPython without any errors.

Expected outcome - The code runs smoothly interacting with the application without errors utilizing IronPython.

Result – **success without any issues**

- Input values and Errors regarding wrong inputs

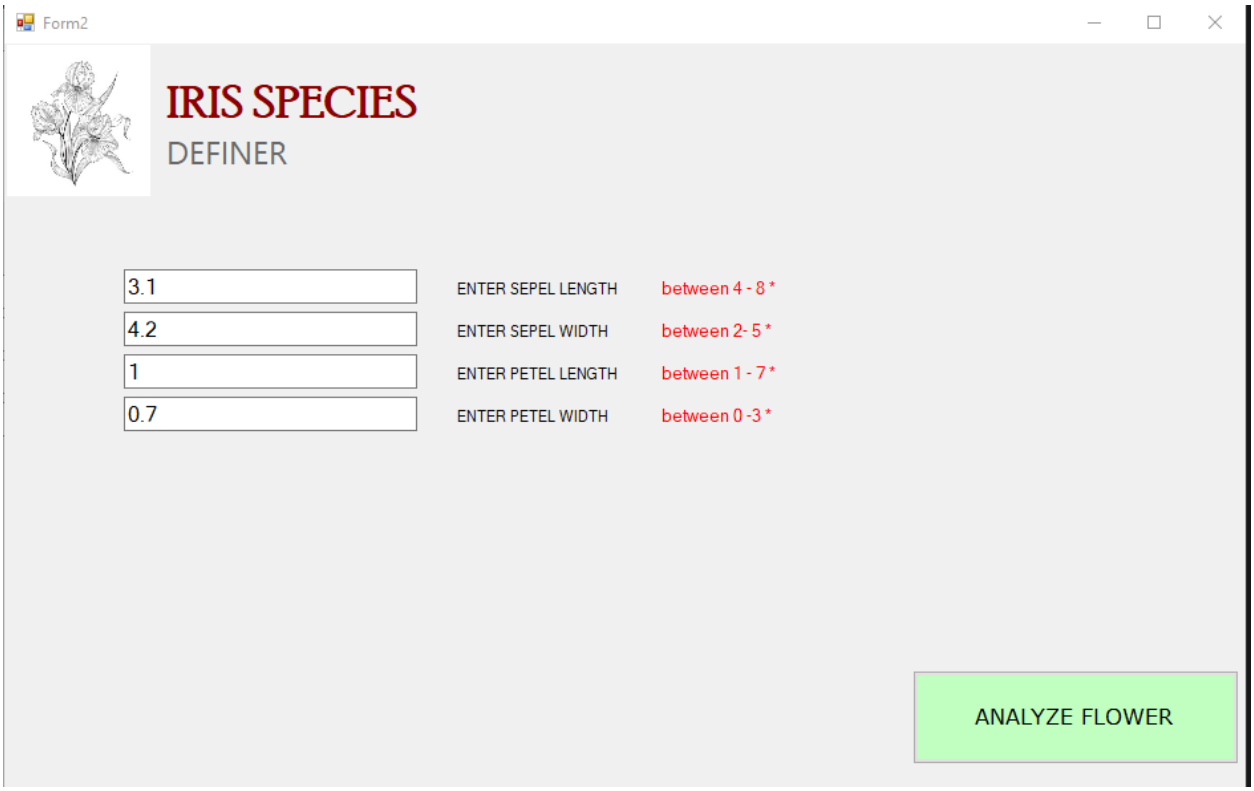

Displayed above is the application interface to input values.

The above screenshot shows that the application provides the facility to input float values in order to analyze iris flower. Therefore input facilitation for positive numeric and float values are a success.

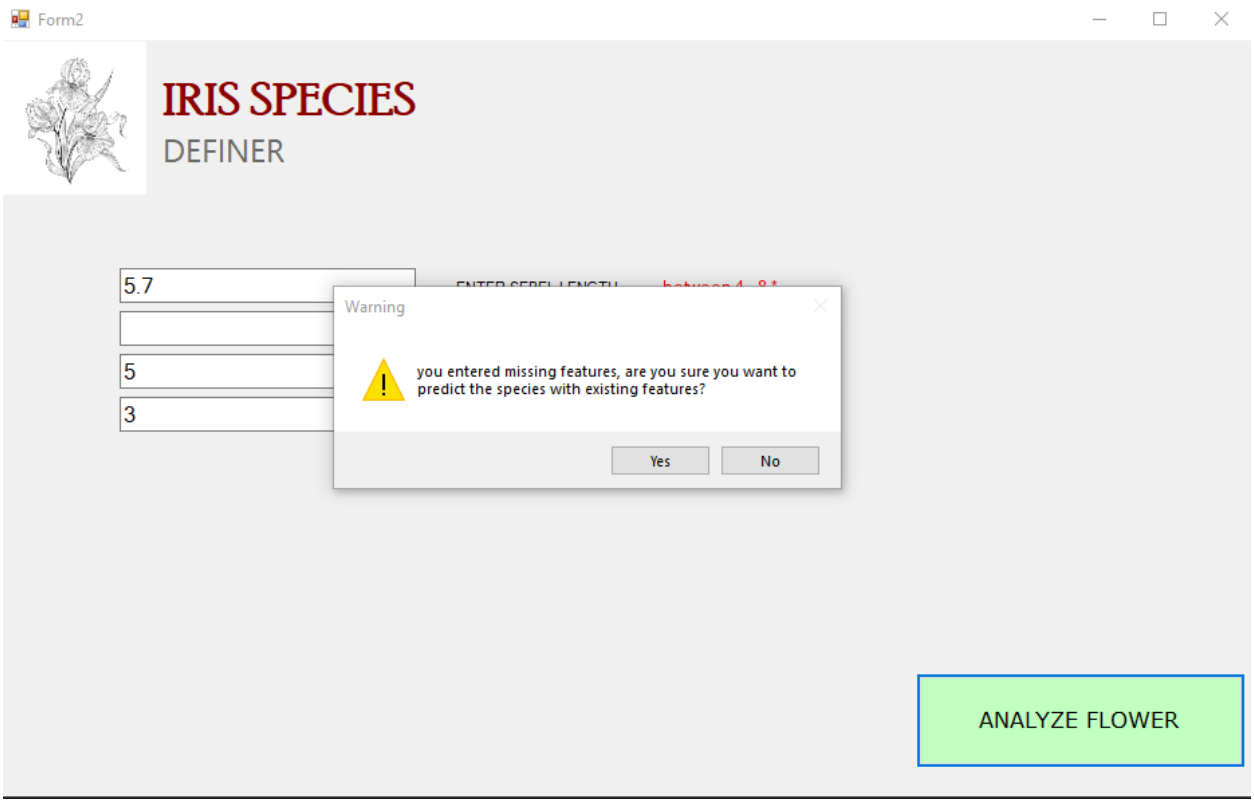

Figure 20 - Missing values

The above screenshot displays that there are missing values entered during the input process, Even though there are missing values the model has the capability to analyze data and provide the prediction based on the given values. The prediction will mot be accurate but, here a warning is displayed to do the analysis on the users own risk.

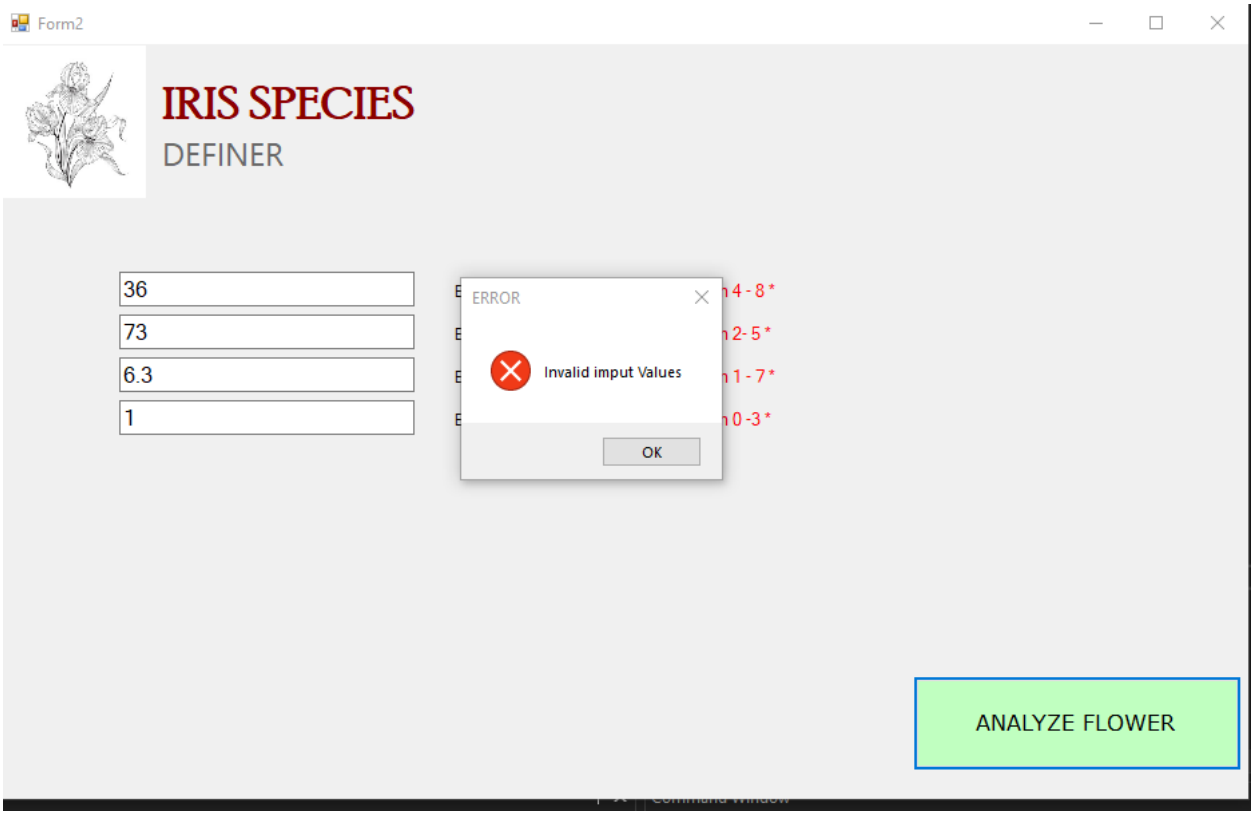

The only possible scale for analysis to be conducted by the model is between the given values, as the iris dataset only provides quality according to this range of values. Above it is shown that there is an error message displayed as the user have entered values exceeding or below the required scale for analysis.

Therefore, the summery of this test,

TEST CASE 2 (2)

Tested By - Lazan Warshath

Description – Testing the Input values and Errors regarding wrong inputs

Expected outcome – The UI provides expected warnings, errors and proceedings for particular input range, availabilities and scale.

Result – **success without any issues**

- <u>Output Screens</u>

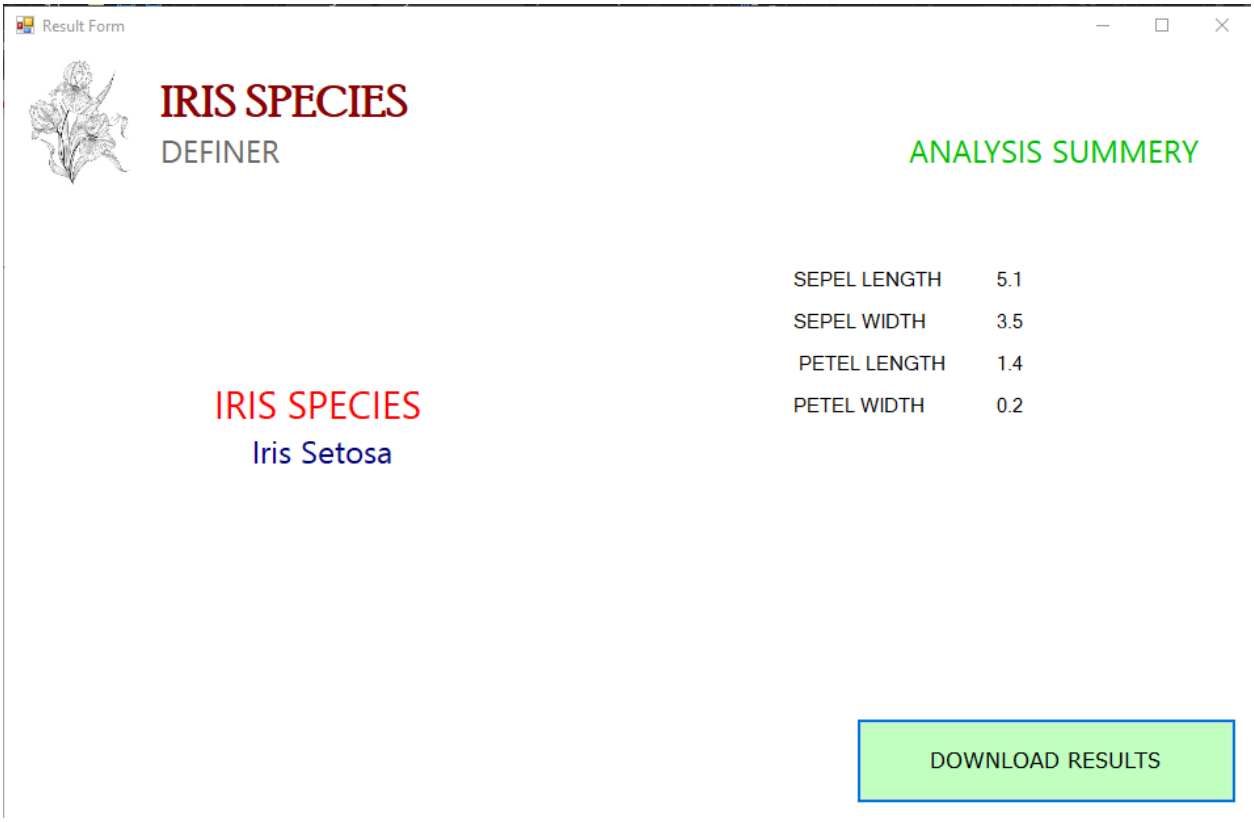

Figure 21 - Successful output

Provided above is evidence for a successful analysis output. The label under “Iris Species” displays the predicted outcome of the analysis. The summery of the analysis containing the values provided by the user during the input phrase is showcased in the right pane of the application. The summery in the right pane of the application consists the exact values provided by the user during the input phrase without any confusions or misplacements (Evidence will be attached next page). The output can also be printed as a .txt file from the button below.

therefore, the summery of this test,

TEST CASE 2 (3)

Tested By - Lazan Warshath

Description – Testing the outputs of the application

Expected outcome – The application provides effective outputs according to the values provided and also provides successful analysis based on provided input.

Result – **success without any issues**

- Examples on analysis providing real world scenario inputs and getting predicted outcomes

Example 1

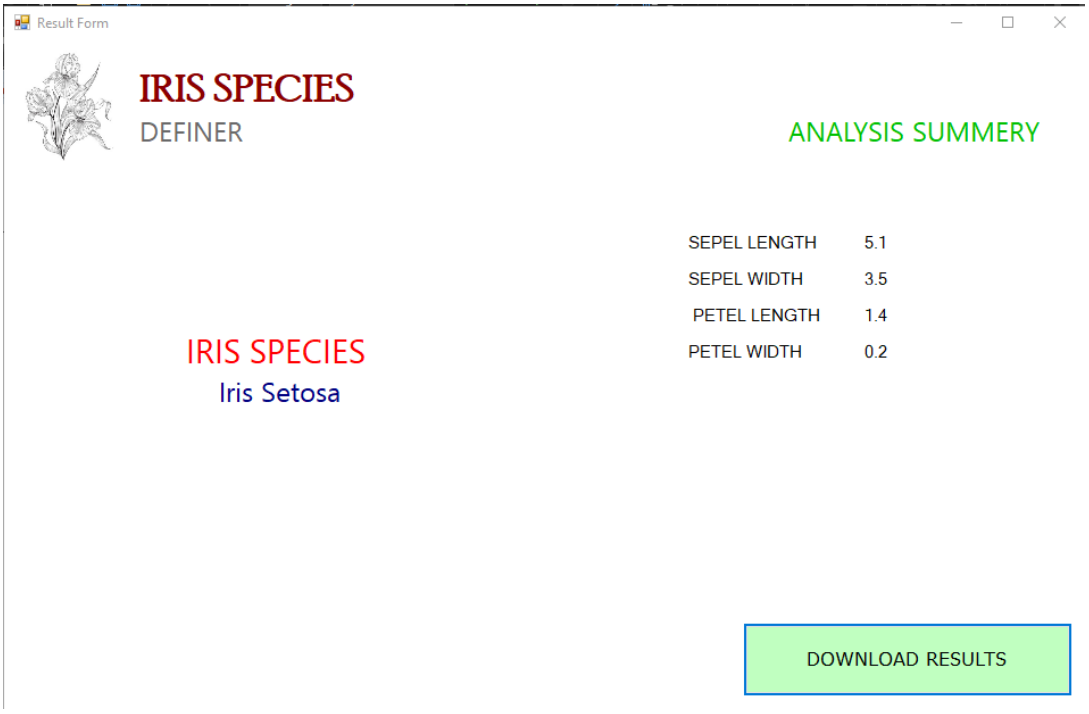
Result Form
IRIS SPECIES
DEFINER
ANALYSIS SUMMERY
SEPEL LENGTH 5.1
SEPEL WIDTH 3.5
PETEL LENGTH 1.4
PETEL WIDTH 0.2
IRIS SPECIES
Iris Setosa
DOWNLOAD RESULTS

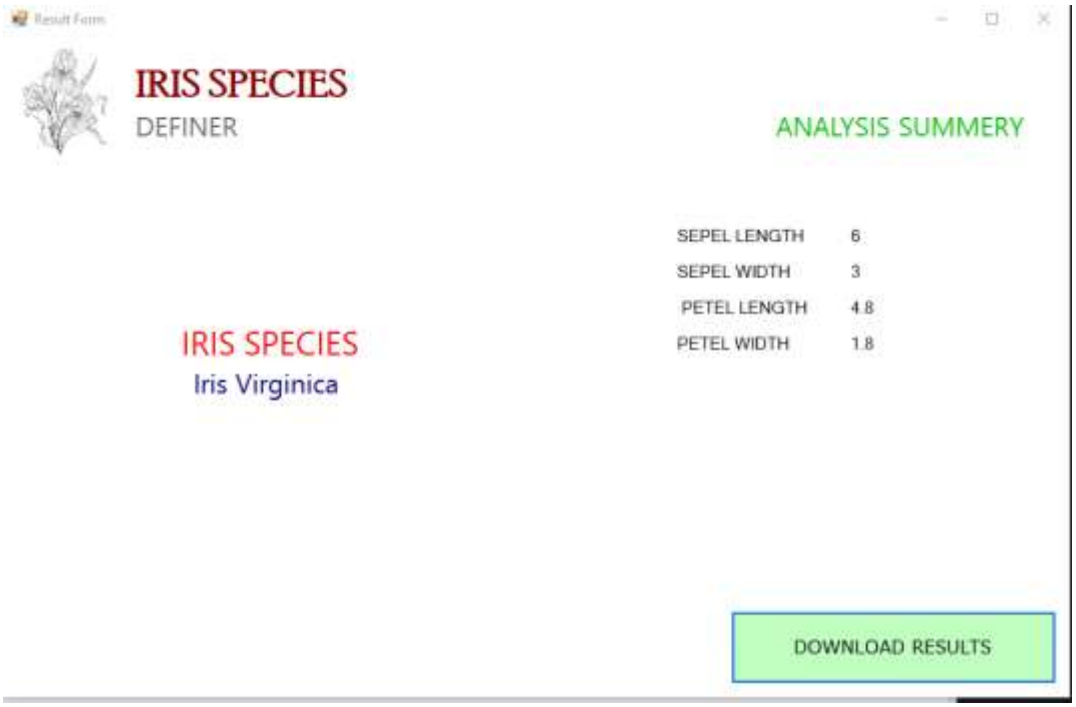

The above input for analysis was provided according to the following extraction from the iris dataset.

1	5.1	3.5	1.4	0.2	Iris-setosa
139	6	3	4.8	1.8	Iris-virginica

And it is clear that the analysis for these flowers were done effectively from the iris species predicting model. This showcases the effectiveness of the model and application.

Therefore, the summery of this test,

TEST CASE 2 (4)

Tested By - Lazan Warshath

Description – Testing the application by providing real world scenario inputs and getting predicted outcomes

Expected outcome – Application to provide outputs similar to the outputs based on the real world examples

Result – **success without any issues**

- Testing Accuracy

```
# Print the metrics
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

```
Accuracy: 1.0
Precision: 1.0
Recall: 1.0
F1 Score: 1.0
```

Figure 22 - Model accuracy

Ther model have achieved full accuracy of 1.0 which is a great score in prediction.

Reviewing the outcomes of using the software tool and the results of the tests

The Iris Quality Prediction application is a powerful tool designed to make the complex task of assessing the quality of Iris flowers much simpler. Its user-friendly interface allows users to easily input data and quickly get accurate predictions. This simplicity is especially important in fields like agriculture and botanical research, where efficient quality control is crucial. By analyzing Iris flower measurements rapidly, stakeholders can make informed decisions about growing, classifying, and studying these flowers, ultimately improving their quality and understanding.

Unit testing ensures that every part of the application works correctly. By testing each component thoroughly, developers can find and fix problems early, making the application stable and reliable. This rigorous testing process gives users confidence that the application will perform consistently and accurately.

Integration testing, particularly with IronPython, shows the application's flexibility. By seamlessly incorporating Python-based machine learning models, the application enhances its analytical capabilities without sacrificing performance or reliability. This adaptability is essential in modern software development, where compatibility with various technologies is key to meeting changing user needs and industry standards.

The application also emphasizes input validation and output verification to ensure data accuracy and integrity. By validating user inputs and verifying analysis outputs, the application reduces

the risk of incorrect predictions, enhancing user trust. This careful attention to detail ensures reliable assessments even in challenging scenarios.

Real-world scenario testing validates the application's predictive capabilities in diverse settings. By simulating actual data and industry-specific use cases, the application proves its effectiveness in real-world applications. This process builds trust among users and stakeholders, confirming that the application can handle real-world challenges and provide actionable insights.

In summary, the Iris Quality Prediction application showcases how software innovation can tackle complex challenges in botanical research and agriculture. Its intuitive interface, effective functionality, and rigorous testing process empower users to make informed decisions, optimize Iris flower quality, and achieve operational excellence. As technology evolves, the application is ready to adapt and innovate, ensuring its ongoing relevance and impact. The thorough testing demonstrates that the Iris Quality Prediction application is a reliable tool for analyzing Iris flower quality and making predictions. Its accurate functionality, user-friendly interface, and dependable performance make it a valuable asset for anyone looking to assess and improve the quality of Iris flowers. With extensive testing ensuring its stability and effectiveness, the application is well-prepared for real-world use.

Presenting the results of the analysis on the chosen dataset

Displayed below are 4 rows of the iris dataset

99	5.1	2.5	3	1.1	Iris-versicolor
100	5.7	2.8	4.1	1.3	Iris-versicolor
101	6.3	3.3	6	2.5	Iris-virginica
102	5.8	2.7	5.1	1.9	Iris-virginica

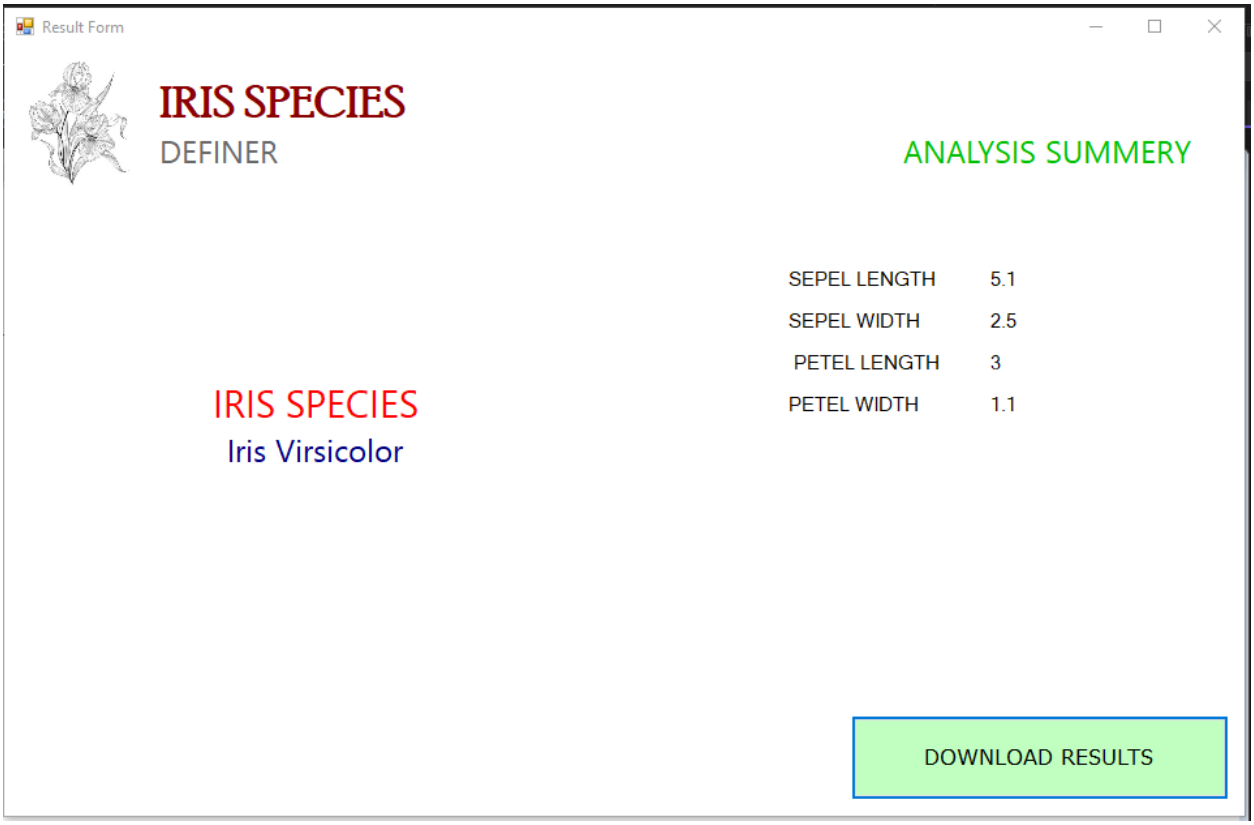
Result Form
IRIS SPECIES
DEFINER
ANALYSIS SUMMERY
SEPEL LENGTH 5.1
SEPEL WIDTH 2.5
PETEL LENGTH 3
PETEL WIDTH 1.1
IRIS SPECIES
Iris Virsicolor
DOWNLOAD RESULTS

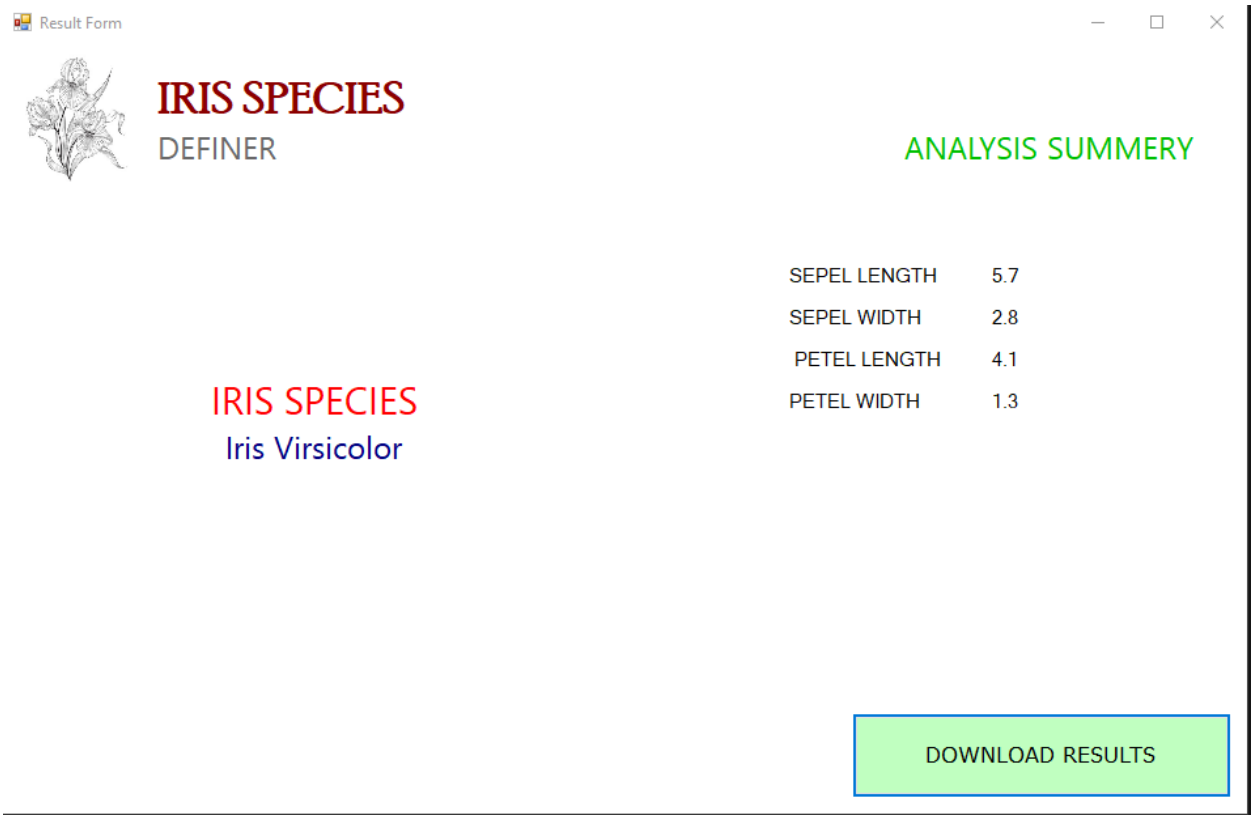
Result Form
IRIS SPECIES
DEFINER
ANALYSIS SUMMERY
SEPEL LENGTH 5.7
SEPEL WIDTH 2.8
PETEL LENGTH 4.1
PETEL WIDTH 1.3
IRIS SPECIES
Iris Virsicolor
DOWNLOAD RESULTS

Result Form

IRIS SPECIES
DEFINER

ANALYSIS SUMMERY

SEPEL LENGTH	6.3
SEPEL WIDTH	3.3
PETEL LENGTH	6
PETEL WIDTH	2.5

IRIS SPECIES

Iris Virginica

DOWNLOAD RESULTS

Result Form

IRIS SPECIES
DEFINER

ANALYSIS SUMMERY

SEPEL LENGTH	5.8
SEPEL WIDTH	2.7
PETEL LENGTH	5.1
PETEL WIDTH	1.9

IRIS SPECIES

Iris Virginica

DOWNLOAD RESULTS

The above screenshots provide evidence that he results of the analysis is similar to the results of the model's predictions. Therefore, it ensures a reliable and accurate model for Iris flower prediction.

Analyzing the output of the data analysis process with focus on the quality of information produced from the dataset and identify possible improvements

Analyzing the results of the data analysis process for the Iris dataset is important for understanding how good the information it provides is. Accuracy is a key factor to consider, which means how well the predicted Iris species match with real observations. High accuracy means the analysis gives reliable insights into Iris species based on the dataset.

Another important thing to look at is precision, which means how detailed and specific the analysis results are. Precision ensures that the information provided covers all the important aspects of Iris species, like sepal length, sepal width, petal length, and petal width. Detailed information helps stakeholders understand specific characteristics of Iris species better.

Completeness is also crucial, as it assesses whether the analysis covers all the relevant factors affecting Iris species in the dataset. This includes things like environmental conditions, cultivation methods, and genetic differences. A complete analysis gives stakeholders a full picture of what determines Iris species, helping them make informed decisions.

Consistency is about how reliable and stable the analysis results are across different samples of Iris flowers in the dataset. It checks if the analysis gives similar species assessments for flowers

with similar traits and under similar conditions. Consistent information makes stakeholders more confident in using the analysis results for making decisions.

Also, relevance is important to make sure the analysis outputs are useful and applicable to real-world situations related to Iris species classification. Relevant information helps stakeholders make decisions that improve classification accuracy and make botanical research more competitive in the field. Evaluating the data analysis results based on these criteria helps stakeholders understand Iris species better and make better decisions.

During the tests of the Iris Quality Prediction application, it showcased that each of these information quality features are important in order to have a good result of the analysis.

Analyzing the output of the data analysis process indicates areas where potential improvements can be made to enhance the overall quality and effectiveness of the analysis. One area of focus is the accuracy of the predictions generated by the analysis. Improving accuracy entails refining the algorithms utilized, incorporating more relevant dataset features, or expanding the dataset's

size and diversity. Ensuring that predictions align closely with real-world observations fosters greater confidence in the analysis's insights.

Another idea to consider is the interpretability of the analysis results. Enhancing interpretability allows stakeholders to better understand the factors influencing Iris species classification and the reasoning behind the predictions. Techniques like feature importance ranking or providing explanations for predictions can boost transparency and understanding. This enables stakeholders to gain deeper insights into the factors driving Iris classification.

Efforts to enhance the scalability and efficiency of the data analysis process without sacrificing accuracy can lead to increased productivity and resource utilization. Optimizing computational algorithms and leveraging parallel computing techniques can reduce processing time and memory usage, expediting decision-making and insights gathering. Streamlining the analysis process improves agility and responsiveness.

Incorporating feedback loops into the data analysis process facilitates continuous improvement. Gathering stakeholder feedback on analysis accuracy and relevance enables iterative updates and enhancements to better meet user needs. This iterative approach ensures the analysis remains pertinent and effective in addressing real-world challenges related to Iris species classification.

By focusing on accuracy, interpretability, scalability, and feedback-driven refinement, stakeholders can identify improvement opportunities in the data analysis process. Continuous refinement and optimization ensure the analysis remains a valuable tool for informing decision-making and driving improvements in Iris species classification.

In future deployments, receiving feedback can further enhance and develop the data analysis application. The user interface can be adapted for mobile platforms to increase portability and productivity. Numerous enhancements can be applied to the application by considering the quality of information aspects mentioned above, leading to a more robust and effective tool for Iris species classification.